



Intel® OSPRay

An Open, Scalable, Parallel, Ray Tracing
Based Rendering Engine for High-Fidelity
Visualization

Version 3.3.0
October 9, 2024

Contents

| | |
|---|-----------|
| 1 OSPRay Overview | 4 |
| 1.1 OSPRay Support and Contact | 4 |
| 2 Building and Finding OSPRay | 6 |
| 2.1 Prerequisites | 6 |
| 2.2 CMake Superbuild | 7 |
| 2.3 Standard CMake Build | 9 |
| 2.4 Finding an OSPRay Install with CMake | 10 |
| 3 OSPRay API | 11 |
| 3.1 Initialization and Shutdown | 11 |
| 3.1.1 Command Line Arguments | 11 |
| 3.1.2 Manual Device Instantiation | 11 |
| 3.1.3 Environment Variables | 14 |
| 3.1.4 Error Handling and Status Messages | 14 |
| 3.1.5 Loading OSPRay Extensions at Runtime | 15 |
| 3.1.6 Shutting Down OSPRay | 15 |
| 3.2 Objects | 15 |
| 3.2.1 Parameters | 16 |
| 3.2.2 Data | 17 |
| 3.3 Volumes | 19 |
| 3.3.1 Structured Regular Volume | 19 |
| 3.3.2 Structured Spherical Volume | 20 |
| 3.3.3 Adaptive Mesh Refinement (AMR) Volume | 21 |
| 3.3.4 Unstructured Volume | 22 |
| 3.3.5 VDB Volume | 23 |
| 3.3.6 Particle Volume | 25 |
| 3.3.7 Transfer Function | 26 |
| 3.3.8 VolumetricModels | 26 |
| 3.4 Geometries | 27 |
| 3.4.1 Mesh | 27 |
| 3.4.2 Subdivision | 28 |
| 3.4.3 Spheres | 28 |
| 3.4.4 Curves | 29 |
| 3.4.5 Boxes | 30 |
| 3.4.6 Planes | 30 |
| 3.4.7 Isosurfaces | 31 |
| 3.4.8 GeometricModels | 31 |
| 3.5 Lights | 32 |
| 3.5.1 Photometric Lights | 32 |
| 3.5.2 Directional Light / Distant Light | 33 |
| 3.5.3 Point Light / Sphere Light | 34 |
| 3.5.4 Spotlight / Ring Light | 34 |
| 3.5.5 Quad Light | 35 |

| | | |
|----------|--|-----------|
| 3.5.6 | Cylinder Light | 35 |
| 3.5.7 | HDRI Light | 36 |
| 3.5.8 | Ambient Light | 36 |
| 3.5.9 | Sun-Sky Light | 37 |
| 3.5.10 | Emissive Objects | 37 |
| 3.6 | Materials | 37 |
| 3.6.1 | OBJ Material | 37 |
| 3.6.2 | Principled | 39 |
| 3.6.3 | CarPaint | 39 |
| 3.6.4 | Metal | 40 |
| 3.6.5 | Alloy | 42 |
| 3.6.6 | Glass | 42 |
| 3.6.7 | ThinGlass | 44 |
| 3.6.8 | MetallicPaint | 44 |
| 3.6.9 | Luminous | 45 |
| 3.7 | Texture | 46 |
| 3.7.1 | Texture2D | 46 |
| 3.7.2 | Volume Texture | 47 |
| 3.7.3 | Texture Transformations | 48 |
| 3.8 | Cameras | 48 |
| 3.8.1 | Perspective Camera | 49 |
| 3.8.2 | Orthographic Camera | 50 |
| 3.8.3 | Panoramic Camera | 50 |
| 3.9 | Scene Hierarchy | 53 |
| 3.9.1 | Groups | 53 |
| 3.9.2 | Instances | 54 |
| 3.9.3 | World | 54 |
| 3.10 | Renderers | 55 |
| 3.10.1 | SciVis Renderer | 57 |
| 3.10.2 | Ambient Occlusion Renderer | 57 |
| 3.10.3 | Path Tracer | 57 |
| 3.11 | Framebuffer | 59 |
| 3.11.1 | Image Operation | 61 |
| 3.12 | Rendering | 62 |
| 3.12.1 | Asynchronous Rendering | 62 |
| 3.12.2 | Synchronous Rendering | 63 |
| 3.12.3 | Rendering and ospCommit | 63 |
| 3.12.4 | Picking | 64 |
| 4 | Modules and Devices | 65 |
| 4.1 | CPU | 65 |
| 4.2 | GPU (Beta) | 65 |
| | Known Issues | 66 |
| 4.3 | Distributed Rendering with MPI | 66 |
| 4.3.1 | MPI Offload Rendering | 66 |
| 4.3.2 | MPI Distributed Rendering | 68 |
| 4.3.3 | Interaction with User Modules | 69 |
| 4.4 | MultiDevice | 69 |
| 5 | Tutorials | 71 |
| 5.1 | ospTutorial | 71 |
| 5.2 | ospExamples | 72 |
| 5.3 | ospMPIDistribTutorial | 73 |
| 5.4 | ospMPIDistribTutorialSpheres and ospMPIDistribTutorialVolume | 74 |
| 5.5 | ospMPIDistribTutorialPartialRepl | 75 |
| 5.6 | ospMPIDistribTutorialReplicated | 75 |

Chapter 1

OSPRay Overview

Intel® OSPRay is an **open source**, **scalable**, and **portable** **ray** tracing engine for high-performance, high-fidelity visualization on Intel Architecture CPUs, Intel Xe GPUs, and Aarch64/ARM64 CPUs. OSPRay is part of the [Intel Rendering Toolkit \(Render Kit\)](#) and is released under the permissive [Apache 2.0 license](#).

The purpose of OSPRay is to provide an open, powerful, and easy-to-use rendering library that allows one to easily build applications that use ray tracing based rendering for interactive applications (including both surface- and volume-based visualizations). OSPRay runs on anything from laptops, to workstations, to compute nodes in HPC systems.

OSPRay internally builds on top of Intel [Embree](#), Intel [Open VKL](#), and Intel [Open Image Denoise](#). The CPU implementation is based on Intel [ISPC \(Implicit SPMD Program Compiler\)](#) and fully exploits modern instruction sets like Intel SSE4, AVX, AVX2, AVX-512 and NEON to achieve high rendering performance. Hence, a CPU with support for at least SSE4.1 is required to run OSPRay on x86_64 architectures, or a CPU with support for NEON is required to run OSPRay on ARM64 architectures.

OSPRay's GPU implementation (beta status) is based on the [SYCL](#) cross-platform programming language implemented by [Intel oneAPI Data Parallel C++ \(DPC++\)](#) and currently supports Intel Arc™ GPUs on Linux and Windows, and Intel Data Center GPU Flex and Max Series on Linux, exploiting ray tracing hardware support.

1.1 OSPRay Support and Contact

OSPRay is under active development, and though we do our best to guarantee stable release versions a certain number of bugs, as-yet-missing features, inconsistencies, or any other issues are still possible. For any such requests or findings please use [OSPRay's GitHub Issue Tracker](#) (or, if you should happen to have a fix for it, you can also send us a pull request).

To receive release announcements simply “Watch” the OSPRay repository on GitHub.

Changes in v3.3.0:

- Add more framebuffer channels and options:
 - Parameter `bool projectedDepth` to switch from euclidean to projected distance for channel `OSP_FB_DEPTH` (often the distance to the image plane, or simply depth)

- Channel `OSP_FB_FIRST_NORMAL` holds the world-space normal of the *first* hit (as opposed to `OSP_FB_NORMAL`, which holds the normal of the first *non-specular* hit for denoising)
 - Channel `OSP_FB_POSITION` holds the world-space position of the first hit
- Improved sampling of layered materials
 - New parameter `specularMetallic` for the Principled material to optionally disable the influence of `specular` to metallicness, improving compatibility with glTF `KHR_materials_specular`
 - Improvements to and documentation of the pathtracer's Shadow Catcher feature (enabled via parameter `shadowCatcherPlane`)

Changes in v3.2.0:

- Sampling improvements:
 - Better performance (lower rendering time and faster convergence)
 - More pleasing blue noise enabled when the total number of frames to be accumulated is known in advance and set as the `targetFrames` parameter at the framebuffer
 - Note a maximum of 64k samples is supported
- Improved `denoiser` image operation:
 - User-controlled quality levels via parameter `quality`
 - Optionally denoise alpha channel as well, enabled via parameter `denoiseAlpha`
- Support half-precision (16 bit float) texture formats `OSP_TEXTURE_[RGBA16F|RGB16F|RA16F|R16F]` and two-channel 32 bit float textures `OSP_TEXTURE_RA32F`
- New parameter `limitIndirectLightSamples` for the `pathtracer` which limits the number of light samples after the first non-specular (i.e., diffuse and glossy) bounce to at most one
- Implement MIP Mapping for better texture filtering. If the additional memory per texture needed cannot be spared, applications can disable the generation of MIP maps with device parameter `disableMipMapGeneration`
- The backplate (background texture) is now always sampled at the pixel center and thus not blurred by the pixel filter anymore
- Avoid color bleeding across eye-subimages when stereo rendering
- Superbuild uses binary packages of Open VKL
- Removed Intel ISPCRT dependency (ISPC compiler is still needed):
 - oneAPI Level Zero Loader is no longer necessary
 - `zeContext` and `zeDeviceDevice` parameters are no longer supported
 - `ispcrtContext` and `ispcrtDeviceDevice` parameters are no longer supported
- Clarify the size of `OSP_BOOL` to be 1 byte
- Fix artifacts occasionally appearing with `gpu` device
- The new minimum versions of dependencies:
 - Embree v4.3.3 (better error reporting)
 - Open Image Denoise v2.3 (better image quality with `HIGH` quality mode, added `FAST` quality mode)
 - rkcommon v1.14.0

For the complete history of changes have a look at the [CHANGELOG](#).

Chapter 2

Building and Finding OSPRay

The latest OSPRay sources are always available at the [OSPRay GitHub repository](#). The default `master` branch should always point to the latest bugfix release.

2.1 Prerequisites

OSPRay currently supports Linux, Mac OS X, and Windows. In addition, before you can build OSPRay you need the following prerequisites:

- You can clone the latest OSPRay sources via:

```
git clone https://github.com/RenderKit/ospray.git
```

- To build OSPRay you need [CMake](#), any form of C++11 compiler (we recommend using GCC, but also support Clang, MSVC, and [Intel® C++ Compiler \(icc\)](#)), and standard Linux development tools.
- Additionally you require a copy of the [Intel® Implicit SPMD Program Compiler \(ISPC\)](#), version 1.23.0 or later. Please obtain a release of ISPC from the [ISPC downloads page](#). If ISPC is not found by CMake its location can be hinted with the variable `ISPC_EXECUTABLE`.
- OSPRay builds on top of the [Intel Rendering Toolkit \(Render Kit\) common library \(rkcommon\)](#). The library provides abstractions for tasking, aligned memory allocation, vector math types, among others. For users who also need to build rkcommon, we recommend the default the Intel [Threading Building Blocks \(TBB\)](#) as tasking system for performance and flexibility reasons. TBB must be built from source when targeting ARM CPUs, or can be built from source as part of the [superbuild](#). Alternatively you can set CMake variable `RKCOMMON_TASKING_SYSTEM` to `OpenMP` or `Internal`.
- OSPRay also heavily uses Intel [Embree](#), installing version 4.3.3 or newer is required. If Embree is not found by CMake its location can be hinted with the variable `embree_DIR`.
- OSPRay supports volume rendering (enabled by default via `OSPRAY_ENABLE_VOLUMES`), which heavily uses Intel [Open VKL](#), version 2.0.1 or newer is required. If Open VKL is not found by CMake its location can be hinted with the variable `openvkl_DIR`, or disable `OSPRAY_ENABLE_VOLUMES`.
- OSPRay also provides an optional module implementing the `denoiser` image operation, which is enabled by `OSPRAY_MODULE_DENOISER`. This module requires Intel [Open Image Denoise](#) in version 2.3.0 or newer. You

may need to hint the location of the library with the CMake variable `OpenImageDenoise_DIR`.

- For the optional MPI modules (enabled by `OSPRAY_MODULE_MPI`), which provide the `mpiOffload` and `mpiDistributed` devices, you need an MPI library and [Google Snappy](#).
- The optional example application, the test suit and benchmarks need some version of OpenGL and GLFW as well as [GoogleTest](#) and [Google Benchmark](#)

Depending on your Linux distribution you can install these dependencies using `yum` or `apt-get`. Some of these packages might already be installed or might have slightly different names.

Type the following to install the dependencies using `yum`:

```
sudo yum install cmake.x86_64
sudo yum install tbb.x86_64 tbb-devel.x86_64
```

Type the following to install the dependencies using `apt-get`:

```
sudo apt-get install cmake-curses-gui
sudo apt-get install libtbb-dev
```

Under Mac OS X these dependencies can be installed using [MacPorts](#):

```
sudo port install cmake tbb
```

Under Windows please directly use the appropriate installers for [CMake](#), [TBB](#), [ISPC](#) (for your Visual Studio version) and [Embree](#).

Additional Prerequisites for GPU Build

To build OSPRay's GPU module you need

- a SYCL compiler, either the open source [oneAPI DPC++ Compiler 2023-10-26](#) or the latest [Intel oneAPI DPC++/C++ Compiler](#)
- a recent [CMake](#), version 3.25.3 or higher

2.2 CMake Superbuild

For convenience, OSPRay provides a CMake Superbuild script which will pull down OSPRay's dependencies and build OSPRay itself. By default, the result is an install directory, with each dependency in its own directory.

Run with:

```
mkdir build
cd build
cmake [<OSPRAY_SOURCE_DIR>/scripts/superbuild]
cmake --build .
```

On Windows make sure to select a 64 bit generator, e.g.

```
cmake -G "Visual Studio 17 2022" [<OSPRAY_SOURCE_DIR>/scripts/superbuild]
```

The resulting `install` directory (or the one set with `CMAKE_INSTALL_PREFIX`) will have everything in it, with one subdirectory per dependency.

CMake options to note (all have sensible defaults):

CMAKE_INSTALL_PREFIX will be the root directory where everything gets installed.

BUILD_JOBS sets the number given to `make -j` for parallel builds.

INSTALL_IN_SEPARATE_DIRECTORIES toggles installation of all libraries in separate or the same directory.

BUILD_OPENVKL whether to enable volume rendering via Open VKL

BUILD_EMBREE_FROM_SOURCE set to OFF will download a pre-built version of Embree.

BUILD_OIDN_FROM_SOURCE set to OFF will download a pre-built version of Open Image Denoise.

OIDN_VERSION determines which version of Open Image Denoise to pull down.

BUILD_OSPRAY_MODULE_MPI set to ON to build OSPRay's MPI module for data-replicated and distributed parallel rendering on multiple nodes.

BUILD_GPU_SUPPORT enables beta GPU support, fetching the SYCL variants of the dependencies and builds **OSPRAY_MODULE_GPU**

BUILD_TBB_FROM_SOURCE set to ON to build TBB from source (required for ARM support). The default setting is OFF.

For the full set of options, run:

```
cmake [<OSPRAY_SOURCE_DIR>/scripts/superbuild]
```

or

```
cmake-gui [<OSPRAY_SOURCE_DIR>/scripts/superbuild]
```

2.2.1 Cross-Compilation with the Superbuild

The superbuild can be passed a [CMake Toolchain file](#) to configure for cross-compilation. This is done by passing the toolchain file when running `cmake`. When cross compiling it is also likely that you'll want to build TBB and Embree from source to ensure they're built for the correct target, rather than the target the Github binaries are built for. It may also be necessary to disable specific ISAs for the target by passing `BUILD_ISA_<ISA_NAME>=OFF` as well.

```
mkdir build
cd build
cmake --toolchain [toolchain_file.cmake] [path/to/this/directory]
-DBUILD_TBB_FROM_SOURCE=ON \
-DBUILD_EMBREE_FROM_SOURCE=ON \
<other arguments>
```

While OSPRay supports ARM natively, it may be desirable to cross-compile it for `x86_64` to run in Rosetta depending on the application integrating OSPRay. This can be done using the toolchain file `toolchains/macos-rosetta.cmake`, and by disabling all non-SSE ISAs when building. This can also be done by launching an `x86_64` bash shell and then compiling as usual in this environment, which will cause the compilation chain to target `x86_64`. The `BUILD_ISA_<ISA NAME>=OFF` flags should be passed to disable all ISAs besides SSE4 for Rosetta:

```
arch -x86_64 bash
mkdir build
cd build
cmake [path/to/this/directory]
-DBUILD_TBB_FROM_SOURCE=ON \
-DBUILD_EMBREE_FROM_SOURCE=ON \
```

```
-DBUILD_ISA_AVX=OFF \
-DBUILD_ISA_AVX2=OFF \
-DBUILD_ISA_AVX512=OFF \
<other arguments>
```

2.3 Standard CMake Build

2.3.1 Compiling OSPRay on Linux and Mac OS X

Assuming the above requisites are all fulfilled, building OSPRay through CMake is easy:

- Create a build directory, and go into it

```
mkdir ospray/build
cd ospray/build
```

(We do recommend having separate build directories for different configurations such as release, debug, etc.).

- The compiler CMake will use will default to whatever the CC and CXX environment variables point to. Should you want to specify a different compiler, run cmake manually while specifying the desired compiler. The default compiler on most linux machines is `gcc`, but it can be pointed to `clang` instead by executing the following:

```
cmake -DCMAKE_CXX_COMPILER=clang++ -DCMAKE_C_COMPILER=clang ..
```

CMake will now use Clang instead of GCC. If you are OK with using the default compiler on your system, then simply skip this step. Note that the compiler variables cannot be changed after the first `cmake` or `ccmake` run.

- Open the CMake configuration dialog

```
ccmake ..
```

- Make sure to properly set build mode and enable the components you need, etc.; then type '`c`'onfigure and '`g`'enerate. When back on the command prompt, build it using

```
make
```

- You should now have `libospray.[so,dylib]` as well as a set of [example applications](#).

2.3.2 Compiling OSPRay on Windows

On Windows using the CMake GUI (`cmake-gui.exe`) is the most convenient way to configure OSPRay and to create the Visual Studio solution files:

- Browse to the OSPRay sources and specify a build directory (if it does not exist yet CMake will create it).
- Click “Configure” and select as generator the Visual Studio version you have; OSPRay needs “Visual Studio 15 2017 Win64” or newer, 32 bit builds are not supported, e.g., “Visual Studio 17 2022”.

- If the configuration fails because some dependencies could not be found then follow the instructions given in the error message, e.g., set the variable `embree_DIR` to the folder where Embree was installed and `openvkl_DIR` to where Open VKL was installed.
- Optionally change the default build options, and then click “Generate” to create the solution and project files in the build directory.
- Open the generated `OSPRay.sln` in Visual Studio, select the build configuration and compile the project.

Alternatively, OSPRay can also be built without any GUI, entirely on the console. In the Visual Studio command prompt type:

```
cd path\to\ospray
mkdir build
cd build
cmake -G "Visual Studio 17 2022" [-D VARIABLE=value] ..
cmake --build . --config Release
```

Use `-D` to set variables for CMake, e.g., the path to Embree with `"-D embree_DIR=\path\to\embree"`.

You can also build only some projects with the `--target` switch. Additional parameters after `--` will be passed to `msbuild`. For example, to build in parallel only the OSPRay library without the example applications use

```
cmake --build . --config Release --target ospray -- /m
```

2.4 Finding an OSPRay Install with CMake

Client applications using OSPRay can find it with CMake’s `find_package()` command. For example,

```
find_package(OSPRay 3.0.0 REQUIRED)
```

finds OSPRay via OSPRay’s configuration file `osprayConfig.cmake`¹. Once found, the following is all that is required to use OSPRay:

```
target_link_libraries(${client_target} ospray::ospray)
```

This will automatically propagate all required include paths, linked libraries, and compiler definitions to the client CMake target (either an executable or library).

Advanced users may want to link to additional targets which are exported in OSPRay’s CMake config, which includes all installed modules. All targets built with OSPRay are exported in the `ospray::` namespace, therefore all targets locally used in the OSPRay source tree can be accessed from an install. For example, `ospray_module_cpu` can be consumed directly via the `ospray::ospray_module_cpu` target. All targets have their libraries, includes, and definitions attached to them for public consumption (please [report bugs](#) if this is broken!).

¹This file is usually in `${install_location}/[lib|lib64]/cmake/ospray-$version/`. If CMake does not find it automatically, then specify its location in variable `ospray_DIR` (either an environment variable or CMake variable).

Chapter 3

OSPRay API

To access the OSPRay API you first need to include the OSPRay header

```
#include "ospray/ospray.h"
```

where the API is compatible with C99 and C++.

3.1 Initialization and Shutdown

To use the API, OSPRay must be initialized with a “device”. A device is the object which implements the API. Creating and initializing a device can be done in either of two ways: command line arguments using `ospInit` or manually instantiating a device and setting parameters on it.

3.1.1 Command Line Arguments

The first is to do so by giving OSPRay the command line from `main()` by calling

```
OSPError ospInit(int *argc, const char **argv);
```

OSPRay parses (and removes) its known command line parameters from your application’s `main` function. For an example see the [tutorial](#). For possible error codes see section [Error Handling and Status Messages](#). It is important to note that the arguments passed to `ospInit` are processed in order they are listed. The following parameters (which are prefixed by convention with “`--osp:`”) are understood:

3.1.2 Manual Device Instantiation

The second method of initialization is to explicitly create the device and possibly set parameters. This method looks almost identical to how other [objects](#) are created and used by OSPRay (described in later sections). The first step is to create the device with

```
OSPDevice ospNewDevice(const char *type);
```

where the `type` string maps to a specific device implementation. OSPRay always provides the “cpu” device, which maps to a fast, local CPU implementation. Other devices can also be added through additional modules, such as distributed MPI device implementations. See next Chapter for details.

Once a device is created, you can call

```
void ospDeviceSetParam(OSPObject, const char *id, OSPDataType type, const void *mem);
```

Table 3.1 – Command line parameters accepted by OSPRay’s `ospInit`.

| Parameter | Description |
|---|--|
| <code>--osp:debug</code> | enables various extra checks and debug output, and disables multi-threading |
| <code>--osp:num-threads=<n></code> | use <code>n</code> threads instead of per default using all detected hardware threads |
| <code>--osp:log-level=<str></code> | set logging level; valid values (in order of severity) are <code>none</code> , <code>error</code> , <code>warning</code> , <code>info</code> , and <code>debug</code> |
| <code>--osp:warn-as-error</code> | send <code>warning</code> and <code>error</code> messages through the <code>error</code> callback, otherwise send <code>warning</code> messages through the <code>message</code> callback; must have sufficient <code>logLevel</code> to enable warnings |
| <code>--osp:verbose</code> | shortcut for <code>--osp:log-level=info</code> and enable debug output on <code>cout</code> , error output on <code>cerr</code> |
| <code>--osp:vv</code> | shortcut for <code>--osp:log-level=debug</code> and enable debug output on <code>cout</code> , error output on <code>cerr</code> |
| <code>--osp:load-modules=<name>[, ...]</code> | load one or more modules during initialization; equivalent to calling <code>ospLoadModule(name)</code> |
| <code>--osp:log-output=<dst></code> | convenience for setting where status messages go; valid values for <code>dst</code> are <code>cerr</code> and <code>cout</code> |
| <code>--osp:error-output=<dst></code> | convenience for setting where error messages go; valid values for <code>dst</code> are <code>cerr</code> and <code>cout</code> |
| <code>--osp:device=<name></code> | use <code>name</code> as the type of device for OSPRay to create; e.g., <code>--osp:device=cpu</code> gives you the default <code>cpu</code> device; Note if the device to be used is defined in a module, remember to pass <code>--osp:load-modules=<name></code> first |
| <code>--osp:set-affinity=<n></code> | if 1, bind software threads to hardware threads; 0 disables binding; default is 0 |
| <code>--osp:device-params=<param>:<value>[, ...]</code> | set one or more other device parameters; equivalent to calling <code>ospDeviceSet*(param, value)</code> |

to set parameters on the device. The semantics of setting parameters is exactly the same as `ospSetParam`, which is documented below in the [parameters](#) section. The following parameters can be set on all devices:

Once parameters are set on the created device, the device must be committed with

```
void ospDeviceCommit(OSPDevice);
```

To use the newly committed device, you must call

```
void ospSetCurrentDevice(OSPDevice);
```

This then sets the given device as the object which will respond to all other OSPRay API calls.

Device handle lifetimes are managed with two calls, the first which increments the internal reference count to the given `OSPDevice`

```
void ospDeviceRetain(OSPDevice)
```

and the second which decrements the reference count

```
void ospDeviceRelease(OSPDevice)
```

Table 3.2 – Parameters shared by all devices.

| Type | Name | Description |
|--------|-------------------------|---|
| int | numThreads | number of threads which OSPRay should use |
| bool | disableMipMapGeneration | disable the default generation of MIP maps for textures (e.g., to save the additional memory needed) |
| uint | logLevel | logging level; valid values (in order of severity) are OSP_LOG_NONE, OSP_LOG_ERROR, OSP_LOG_WARNING, OSP_LOG_INFO, and OSP_LOG_DEBUG |
| string | logOutput | convenience for setting where status messages go; valid values are <code>cerr</code> and <code>cout</code> |
| string | errorOutput | convenience for setting where error messages go; valid values are <code>cerr</code> and <code>cout</code> |
| bool | debug | set debug mode; equivalent to <code>logLevel=debug</code> and <code>numThreads=1</code> |
| bool | warnAsError | send warning and error messages through the error callback, otherwise send warning messages through the message callback; must have sufficient <code>logLevel</code> to enable warnings |
| bool | setAffinity | bind software threads to hardware threads if set to 1; 0 disables binding omitting the parameter will let OSPRay choose |

Users can change parameters on the device after initialization (from either method above), by calling

```
OSPDevice ospGetCurrentDevice();
```

This function returns the handle to the device currently used to respond to OSPRay API calls, where users can set/change parameters and recommit the device. If changes are made to the device that is already set as the current device, it does not need to be set as current again. Note this API call will increment the ref count of the returned device handle, so applications must use `ospDeviceRelease` when finished using the handle to avoid leaking the underlying device object. If there is no current device set, this will return an invalid NULL handle.

When a device is created, its reference count is initially 1. When a device is set as the current device, it internally has its reference count incremented. Note that `ospDeviceRetain` and `ospDeviceRelease` should only be used with reference counts that the application tracks: removing reference held by the current set device should be handled by `ospShutdown`. Thus, `ospDeviceRelease` should only decrement the reference counts that come from `ospNewDevice`, `ospGetCurrentDevice`, and the number of explicit calls to `ospDeviceRetain`.

OSPRay allows applications to query runtime properties of a device in order to do enhanced validation of what device was loaded at runtime. The following function can be used to get these device-specific properties (attributes about the device, not parameter values)

```
int64_t ospDeviceGetProperty(OSPDevice, OSPDeviceProperty);
```

It returns an integer value of the queried property and the following properties can be provided as parameter:

```
OSP_DEVICE_VERSION
OSP_DEVICE_VERSION_MAJOR
OSP_DEVICE_VERSION_MINOR
OSP_DEVICE_VERSION_PATCH
OSP_DEVICE_SO_VERSION
```

3.1.3 Environment Variables

OSPRay's generic device parameters can be overridden via environment variables for easy changes to OSPRay's behavior without needing to change the application (variables are prefixed by convention with "OSPRAY_"):

Table 3.3 – Environment variables interpreted by OSPRay.

| Variable | Description |
|----------------------|--|
| OSPRAY_NUM_THREADS | equivalent to --osp:num-threads |
| OSPRAY_LOG_LEVEL | equivalent to --osp:log-level |
| OSPRAY_LOG_OUTPUT | equivalent to --osp:log-output |
| OSPRAY_ERROR_OUTPUT | equivalent to --osp:error-output |
| OSPRAY_DEBUG | equivalent to --osp:debug |
| OSPRAY_WARN_AS_ERROR | equivalent to --osp:warn-as-error |
| OSPRAY_SET_AFFINITY | equivalent to --osp:set-affinity |
| OSPRAY_LOAD_MODULES | equivalent to --osp:load-modules, can be a comma separated list of modules which will be loaded in order |
| OSPRAY_DEVICE | equivalent to --osp:device: |

Note that these environment variables take precedence over values specified through `ospInit` or manually set device parameters.

3.1.4 Error Handling and Status Messages

The following errors are currently used by OSPRay:

| Name | Description |
|-----------------------|--|
| OSP_NO_ERROR | no error occurred |
| OSP_UNKNOWN_ERROR | an unknown error occurred |
| OSP_INVALID_ARGUMENT | an invalid argument was specified |
| OSP_INVALID_OPERATION | the operation is not allowed for the specified object |
| OSP_OUT_OF_MEMORY | there is not enough memory to execute the command |
| OSP_UNSUPPORTED_CPU | the CPU is not supported (minimum ISA is SSE4.1 on x86_64 and NEON on ARM64) |
| OSP_VERSION_MISMATCH | a module could not be loaded due to mismatching version |

Table 3.4 – Possible error codes, i.e., valid named constants of type `OSPError`.

These error codes are either directly return by some API functions, or are recorded to be later queried by the application via

```
OSPError ospDeviceGetLastErrorCode(OSPDevice);
```

A more descriptive error message can be queried by calling

```
const char* ospDeviceGetLastErrorMsg(OSPDevice);
```

Alternatively, the application can also register a callback function of type

```
typedef void (*OSPErrorCallback)(void *userData, OSPError, const char* errorDetails);
```

via

```
void ospDeviceSetErrorCallback(OSPDevice, OSPErrorCallback, void *userData);
```

to get notified when errors occur.

Applications may be interested in messages which OSPRay emits, whether for debugging or logging events. Applications can call

```
void ospDeviceSetStatusCallback(OSPDevice, OSPStatusCallback, void *userData);
```

in order to register a callback function of type

```
typedef void (*OSPStatusCallback)(void *userData, const char* messageText);
```

which OSPRay will use to emit status messages. By default, OSPRay uses a callback which does nothing, so any output desired by an application will require that a callback is provided. Note that callbacks for C++ `std::cout` and `std::cerr` can be alternatively set through `ospInit` or the `OSPRAY_LOG_OUTPUT` environment variable.

Applications can clear either callback by passing `NULL` instead of an actual function pointer.

3.1.5 Loading OSPRay Extensions at Runtime

OSPRay's functionality can be extended via plugins (which we call "modules"), which are implemented in shared libraries. To load module `name` from `libospray_module_<name>.so` (on Linux and Mac OS X) or `ospray_module_<name>.dll` (on Windows) use

```
OSPError ospLoadModule(const char *name);
```

Modules are searched in OS-dependent paths. `ospLoadModule` returns `OSP_NO_ERROR` if the plugin could be successfully loaded.

3.1.6 Shutting Down OSPRay

When the application is finished using OSPRay (typically on application exit), the OSPRay API should be finalized with

```
void ospShutdown();
```

This API call ensures that the current device is cleaned up appropriately. Due to static object allocation having non-deterministic ordering, it is recommended that applications call `ospShutdown` before the calling application process terminates.

3.2 Objects

All entities of OSPRay (the `renderer`, `volumes`, `geometries`, `lights`, `cameras`, ...) are a logical specialization of `OSPObject` and share common mechanism to deal with parameters and lifetime.

An important aspect of object parameters is that parameters do not get passed to objects immediately. Instead, parameters are not visible at all to objects until they get explicitly committed to a given object via a call to

```
void ospCommit(OSPObject);
```

at which time all previously additions or changes to parameters are visible at the same time. If a user wants to change the state of an existing object (e.g., to change the origin of an already existing camera) it is perfectly valid to do so, as long as the changed parameters are recommitted.

The commit semantic allow for batching up multiple small changes, and specifies exactly when changes to objects will occur. This can impact performance and consistency for devices crossing a PCI bus or across a network.

Note that OSPRay uses reference counting to manage the lifetime of all objects, so one cannot explicitly “delete” any object. Instead, to indicate that the application does not need and does not access the given object anymore, call

```
void ospRelease(OSPObject);
```

This decreases its reference count and if the count reaches 0 the object will automatically get deleted. Passing NULL is not an error. Note that every handle returned via the API needs to be released when the object is no longer needed, to avoid memory leaks.

Sometimes applications may want to have more than one reference to an object, where it is desirable for the application to increment the reference count of an object. This is done with

```
void ospRetain(OSPObject);
```

It is important to note that this is only necessary if the application wants to call `ospRelease` on an object more than once: objects which contain other objects as parameters internally increment/decrement ref counts and should not be explicitly done by the application.

3.2.1 Parameters

Parameters allow to configure the behavior of and to pass data to objects. However, objects do *not* have an explicit interface for reasons of high flexibility and a more stable compile-time API. Instead, parameters are passed separately to objects in an arbitrary order, and unknown parameters will simply be ignored (though a warning message will be posted). The following function allows adding various types of parameters with name `id` to a given object:

```
void ospSetParam(OSPObject, const char *id, OSPDataType type, const void *mem);
```

The valid parameter names for all `OSPObject`s and what types are valid are discussed in future sections.

Note that `mem` must always be a pointer to the object, otherwise accidental type casting can occur. This is especially true for pointer types (`OSP_VOID_PTR` and `OSPObject` handles), as they will implicitly cast to `void\ *`, but be incorrectly interpreted. To help with some of these issues, there also exist variants of `ospSetParam` for specific types, such as `ospSetInt` and `ospSetVec3f` in the OSPRay utility library (found in `ospray_util.h`). Note that half precision float parameters `OSP_HALF`, `OSP_VEC[234]H` are not supported.

Users can also remove parameters that have been explicitly set from `ospSetParam`. Any parameters which have been removed will go back to their default value during the next commit unless a new parameter was set after the parameter was removed. To remove a parameter, use

```
void ospRemoveParam(OSPObject, const char *id);
```

3.2.2 Data

OSPRay consumes data arrays from the application using a specific object type, `OSPData`. There are several components to describing a data array: element type, 1/2/3 dimensional striding, and whether the array is shared with the application or copied into opaque, OSPRay-owned memory.

Shared data arrays require that the application's array memory outlives the lifetime of the created `OSPData`, as OSPRay is referring to application memory. Where this is not preferable, applications use opaque arrays to allow the `OSPData` to own the lifetime of the array memory. However, opaque arrays dictate the cost of copying data into it, which should be kept in mind.

Thus, the most efficient way to specify a data array from the application is to create a shared data array, which is done with

```
OSPData ospNewSharedData(const void *sharedData,
    OSPDataType,
    uint64_t numItems1,
    int64_t byteStride1 = 0,
    uint64_t numItems2 = 1,
    int64_t byteStride2 = 0,
    uint64_t numItems3 = 1,
    int64_t byteStride3 = 0,
    OSPDeleterCallback = NULL,
    void *userData = NULL);
```

The call returns an `OSPData` handle to the created array. The calling program guarantees that the `sharedData` pointer will remain valid for the duration that this data array is being used. The number of elements `numItems` must be positive (there cannot be an empty data object). The data is arranged in three dimensions, with specializations to two or one dimension (if some `numItems` are 1). The distance between consecutive elements (per dimension) is given in bytes with `byteStride` and can also be negative. If `byteStride` is zero it will be determined automatically (e.g., as `sizeof(type)`). Strides do not need to be ordered, i.e., `byteStride2` can be smaller than `byteStride1`, which is equivalent to a transpose. However, if the stride should be calculated, then an ordering in dimensions is assumed to disambiguate, i.e., `byteStride1 < byteStride2 < byteStride3`.

An application can pass ownership of shared data to OSPRay (for example, when it temporarily created a modified version of its data only to make it compatible with OSPRay) by providing a deleter function that OSPRay will call whenever the time comes to deallocate the shared buffer. The deleter function has the following signature:

```
typedef void (*OSPDeleterCallback)(const void *userData, const void *sharedData);
```

where `sharedData` will receive the address of the buffer and `userData` will receive whatever additional state the function needs to perform the deletion (both provided to `ospNewSharedData` when sharing the data with OSPRay).

The enum type `OSPDataType` describes the different element types that can be represented in OSPRay; valid constants are listed in the table below.

If the elements of the array are handles to objects, then their reference counter is incremented.

An opaque `OSPData` with memory allocated by OSPRay is created with

```
OSPData ospNewData(OSPDataType,
    uint64_t numItems1,
    uint64_t numItems2 = 1,
    uint64_t numItems3 = 1);
```

Table 3.5 – Valid named constants for OSPDataType.

| Type / Name | Description |
|----------------------------|---|
| OSP_DEVICE | API device object reference |
| OSP_DATA | data reference |
| OSP_OBJECT | generic object reference |
| OSP_CAMERA | camera object reference |
| OSP_FRAMEBUFFER | framebuffer object reference |
| OSP_FUTURE | future object reference |
| OSP_LIGHT | light object reference |
| OSP_MATERIAL | material object reference |
| OSP_TEXTURE | texture object reference |
| OSP_RENDERER | renderer object reference |
| OSP_WORLD | world object reference |
| OSP_GROUP | group object reference |
| OSP_INSTANCE | instance object reference |
| OSP_GEOMETRY | geometry object reference |
| OSP_GEOMETRIC_MODEL | geometric model object reference |
| OSP_VOLUME | volume object reference |
| OSP_VOLUMETRIC_MODEL | volumetric model object reference |
| OSP_TRANSFER_FUNCTION | transfer function object reference |
| OSP_IMAGE_OPERATION | image operation object reference |
| OSP_STRING | C-style zero-terminated character string |
| OSP_BOOL | 8 bit boolean |
| OSP_CHAR, OSP_VEC[234]C | 8 bit signed character scalar and [234]-element vector |
| OSP_UCHAR, OSP_VEC[234]UC | 8 bit unsigned character scalar and [234]-element vector |
| OSP_SHORT, OSP_VEC[234]S | 16 bit unsigned integer scalar and [234]-element vector |
| OSP USHORT, OSP_VEC[234]US | 16 bit unsigned integer scalar and [234]-element vector |
| OSP_INT, OSP_VEC[234]I | 32 bit signed integer scalar and [234]-element vector |
| OSP_UINT, OSP_VEC[234]UI | 32 bit unsigned integer scalar and [234]-element vector |
| OSP_LONG, OSP_VEC[234]L | 64 bit signed integer scalar and [234]-element vector |
| OSP ULONG, OSP_VEC[234]UL | 64 bit unsigned integer scalar and [234]-element vector |
| OSP_HALF, OSP_VEC[234]H | 16 bit half precision floating-point scalar and [234]-element vector (IEEE 754 binary16) |
| OSP_FLOAT, OSP_VEC[234]F | 32 bit single precision floating-point scalar and [234]-element vector |
| OSP_DOUBLE, OSP_VEC[234]D | 64 bit double precision floating-point scalar and [234]-element vector |
| OSP_BOX[1234]I | 32 bit integer box (lower + upper bounds) |
| OSP_BOX[1234]F | 32 bit single precision floating-point box (lower + upper bounds) |
| OSP_LINEAR[23]F | 32 bit single precision floating-point linear transform ([23] vectors) |
| OSP_AFFINE[23]F | 32 bit single precision floating-point affine transform (linear transform plus translation) |
| OSP_QUATF | 32 bit single precision floating-point quaternion, in (i, j, k, w) layout |
| OSP_VOID_PTR | raw memory address (only found in module extensions) |

To allow for (partial) copies or updates of data arrays use

```
void ospCopyData(const OSPData source,
                 OSPData destination,
                 uint64_t destinationIndex1 = 0,
                 uint64_t destinationIndex2 = 0,
                 uint64_t destinationIndex3 = 0);
```

which will copy the whole¹ content of the `source` array into `destination` at the given location `destinationIndex`. The `OSPDataTypes` of the data objects must match. The region to be copied must be valid inside the destination, i.e., in all dimensions, `destinationIndex + sourceSize <= destinationSize`. The affected region `[destinationIndex, destinationIndex + sourceSize)` is marked as dirty, which may be used by OSPRay to only process or update that sub-region (e.g., updating an acceleration structure). If the destination array is shared with `OSPData` by the application (created with `ospNewSharedData`), then

- the source array must be shared as well (thus `ospCopyData` cannot be used to read opaque data)
- if source and destination memory overlaps (aliasing), then behavior is undefined
- except if source and destination regions are identical (including matching strides), which can be used by application to mark that region as dirty (instead of the whole `OSPData`)

To add a data array as parameter named `id` to another object call also use

```
void ospSetObject(OSPObject, const char *id, OSPData);
```

3.3 Volumes

Volumes are volumetric data sets with discretely sampled values in 3D space, typically a 3D scalar field. To create a new volume object of given type `type` use

```
OSPVolume ospNewVolume(const char *type);
```

Note that OSPRay's implementation forwards `type` directly to Open VKL, allowing new Open VKL volume types to be usable within OSPRay without the need to change (or even recompile) OSPRay.

3.3.1 Structured Regular Volume

Structured volumes only need to store the values of the samples, because their addresses in memory can be easily computed from a 3D position. A common type of structured volumes are regular grids.

Structured regular volumes are created by passing the type string “`structuredRegular`” to `ospNewVolume`. Structured volumes are represented through an `OSPData` 3D array `data` (which may or may not be shared with the application). The voxel data must be laid out in xyz-order² and can be compact (best for performance) or can have a stride between voxels, specified through the `byteStride1` parameter when creating the `OSPData`. Only 1D strides are supported, additional strides between scanlines (2D, `byteStride2`) and slices (3D, `byteStride3`) are not.

The parameters understood by structured volumes are summarized in the table below.

¹ The number of items to be copied is defined by the size of the source array.

² For consecutive memory addresses the x-index of the corresponding voxel changes the quickest.


Table 3.6 – Configuration parameters for structured regular volumes.

| Type | Name | Default | Description |
|---------|----------------|-----------------------------|---|
| vec3f | gridOrigin | (0, 0, 0) | origin of the grid in object-space |
| vec3f | gridSpacing | (1, 1, 1) | size of the grid cells in object-space |
| OSPData | data | | the actual voxel 3D <code>data</code> |
| bool | cellCentered | false | whether the data is provided per cell (as opposed to per vertex) |
| uint | filter | OSP_VOLUME_FILTER_LINEAR | filter used for reconstructing the field, also allowed is OSP_VOLUME_FILTER_NEAREST and OSP_VOLUME_FILTER_CUBIC |
| uint | gradientFilter | same as <code>filter</code> | filter used during gradient computations |
| float | background | NaN | value that is used when sampling an undefined region outside the volume domain |

The size of the volume is inferred from the size of the 3D array `data`, as is the type of the voxel values (currently supported are: OSP_UCHAR, OSP_SHORT, OSP USHORT, OSP_HALF, OSP_FLOAT, and OSP_DOUBLE). Data can be provided either per cell or per vertex (the default), selectable via the `cellCentered` parameter (which will also affect the computed bounding box).

3.3.2 Structured Spherical Volume

Structured spherical volumes are also supported, which are created by passing a type string of “structuredSpherical” to `ospNewVolume`. The grid dimensions and parameters are defined in terms of radial distance r , inclination angle θ , and azimuthal angle ϕ , conforming with the ISO convention for spherical coordinate systems. The coordinate system and parameters understood by structured spherical volumes are summarized below.

**Figure 3.1** – Coordinate system of structured spherical volumes.

The dimensions (r, θ, ϕ) of the volume are inferred from the size of the 3D array `data`, as is the type of the voxel values (currently supported are: OSP_UCHAR, OSP_SHORT, OSP USHORT, OSP_HALF, OSP_FLOAT, and OSP_DOUBLE).

These grid parameters support flexible specification of spheres, hemispheres, spherical shells, spherical wedges, and so forth. The grid extents (computed as `[gridOrigin, gridOrigin + (dimensions - 1) * gridSpacing]`) however must be constrained such that:

Table 3.7 – Configuration parameters for structured spherical volumes.

| Type | Name | Default | Description |
|---------|----------------|---------------------------------|--|
| vec3f | gridOrigin | (0, 0, 0) | origin of the grid in units of (r, θ, ϕ) ; angles in degrees |
| vec3f | gridSpacing | (1, $180/dim.y$, $360/dim.z$) | size of the grid cells in units of (r, θ, ϕ) , per default covering the full sphere; angles in degrees |
| OSPData | data | | the actual voxel 3D data |
| uint | filter | OSP_VOLUME_FILTER_LINEAR | filter used for reconstructing the field, also allowed is OSP_VOLUME_FILTER_NEAREST |
| uint | gradientFilter | same as <code>filter</code> | filter used during gradient computations |
| float | background | NaN | value that is used when sampling an undefined region outside the volume domain |

- $r \geq 0$
- $0 \leq \theta \leq 180$
- $0 \leq \phi \leq 360$

3.3.3 Adaptive Mesh Refinement (AMR) Volume

OSPRay currently supports block-structured (Berger-Colella) AMR volumes. Volumes are specified as a list of blocks, which exist at levels of refinement in potentially overlapping regions. Blocks exist in a tree structure, with coarser refinement level blocks containing finer blocks. The cell width is equal for all blocks at the same refinement level, though blocks at a coarser level have a larger cell width than finer levels.

There can be any number of refinement levels and any number of blocks at any level of refinement. An AMR volume type is created by passing the type string “`amr`” to `ospNewVolume`.

Blocks are defined by three parameters: their bounds, the refinement level in which they reside, and the scalar data contained within each block.

Note that cell widths are defined *per refinement level*, not per block.

Table 3.8 – Configuration parameters for AMR volumes.

| Type | Name | Default | Description |
|-----------|--------------|-----------------|---|
| uint | method | OSP_AMR_CURRENT | OSPMRMethod sampling method. Supported methods are: OSP_AMR_CURRENT OSP_AMR_FINEST OSP_AMR_OCTANT |
| float[] | cellWidth | NULL | array of each level's cell width |
| box3i[] | block.bounds | NULL | data array of grid sizes (in voxels) for each AMR block |
| int[] | block.level | NULL | array of each block's refinement level |
| OSPData[] | block.data | NULL | data array of OSPData containing the actual scalar voxel data, only OSP_FLOAT is supported as OSPDataType |
| vec3f | gridOrigin | (0, 0, 0) | origin of the grid |
| vec3f | gridSpacing | (1, 1, 1) | size of the grid cells |
| float | background | NaN | value that is used when sampling an undefined region outside the volume domain |

Lastly, note that the `gridOrigin` and `gridSpacing` parameters act just like the structured volume equivalent, but they only modify the root (coarsest level) of refinement.

In particular, OSPRay's / Open VKL's AMR implementation was designed to cover Berger-Colella [1] and Chombo [2] AMR data. The `method` parameter above determines the interpolation method used when sampling the volume.

`OSP_AMR_CURRENT` finds the finest refinement level at that cell and interpolates through this “current” level

`OSP_AMR_FINEST` will interpolate at the closest existing cell in the volume-wide finest refinement level regardless of the sample cell's level

`OSP_AMR_OCTANT` interpolates through all available refinement levels at that cell. This method avoids discontinuities at refinement level boundaries at the cost of performance

Details and more information can be found in the publication for the implementation [3].

1. M.J. Berger and P. Colella, “Local adaptive mesh refinement for shock hydrodynamics.” *Journal of Computational Physics* 82.1 (1989): 64-84. DOI: 10.1016/0021-9991(89)90035-1
2. M. Adams, P. Colella, D.T. Graves, J.N. Johnson, N.D. Keen, T.J. Ligocki, D.F. Martin, P.W. McCorquodale, D. Modiano, P.O. Schwartz, T.D. Sternberg, and B. Van Straalen, “Chombo Software Package for AMR Applications – Design Document”, Lawrence Berkeley National Laboratory Technical Report LBNL-6616E.
3. I. Wald, C. Brownlee, W. Usher, and A. Knoll, “CPU volume rendering of adaptive mesh refinement data”. *SIGGRAPH Asia 2017 Symposium on Visualization – SA '17*, 18(8), 1–8. DOI: 10.1145/3139295.3139305

3.3.4 Unstructured Volume

Unstructured volumes can have their topology and geometry freely defined. Geometry can be composed of tetrahedral, hexahedral, wedge or pyramid cell types. The data format used is compatible with VTK and consists of multiple arrays: vertex positions and values, vertex indices, cell start indices, cell types, and cell values. An unstructured volume type is created by passing the type string “unstructured” to `ospNewVolume`.

Sampled cell values can be specified either per-vertex (`vertex.data`) or per-cell (`cell.data`). If both arrays are set, `cell.data` takes precedence.

Similar to a mesh, each cell is formed by a group of indices into the vertices. For each vertex, the corresponding (by array index) data value will be used for sampling when rendering, if specified. The index order for a tetrahedron is the same as VTK_TETRA: bottom triangle counterclockwise, then the top vertex.

For hexahedral cells, each hexahedron is formed by a group of eight indices into the vertices and data values. Vertex ordering is the same as VTK_HEXAHEDRON: four bottom vertices counterclockwise, then top four counterclockwise.

For wedge cells, each wedge is formed by a group of six indices into the vertices and data values. Vertex ordering is the same as VTK_WEDGE: three bottom vertices counterclockwise, then top three counterclockwise.

For pyramid cells, each cell is formed by a group of five indices into the vertices and data values. Vertex ordering is the same as VTK_PYRAMID: four bottom vertices counterclockwise, then the top vertex.

To maintain VTK data compatibility, the `index` array may be specified with cell sizes interleaved with vertex indices in the following format: $n, id_1, \dots, id_n, m, id_1, \dots, id_m$. This alternative `index` array layout can be enabled through the `indexPrefixed` flag (in which case, the `cell.type` parameter must be omitted).

Table 3.9 – Configuration parameters for unstructured volumes.

| Type | Name | Default | Description |
|---------------------|--------------------|---------|--|
| vec3f[] | vertex.position | | <code>data</code> array of vertex positions |
| float[] | vertex.data | | <code>data</code> array of vertex data values to be sampled |
| uint32[] / uint64[] | index | | <code>data</code> array of indices (into the vertex array(s)) that form cells |
| bool | indexPrefixed | false | indicates that the <code>index</code> array is compatible to VTK, where the indices of each cell are prefixed with the number of vertices |
| uint32[] / uint64[] | cell.index | | <code>data</code> array of locations (into the <code>index</code> array), specifying the first index of each cell |
| float[] | cell.data | | <code>data</code> array of cell data values to be sampled |
| uint8[] | cell.type | | <code>data</code> array of cell types (VTK compatible), only set if <code>indexPrefixed = false</code> . Supported types are: <code>OSP_TETRAHEDRON</code> <code>OSP_HEXAHEDRON</code> <code>OSP_WEDGE</code> <code>OSP_PYRAMID</code> |
| bool | hexIterative | false | hexahedron interpolation method, defaults to fast non-iterative version which could have rendering inaccuracies may appear if hex is not parallelepiped |
| bool | precomputedNormals | false | whether to accelerate by precomputing, at a cost of 12 bytes/face |
| float | background | NaN | value that is used when sampling an undefined region outside the volume domain |

3.3.5 VDB Volume

VDB volumes implement a data structure that is very similar to the data structure outlined in Museth [1], they are created by passing the type string “vdb” to `ospNewVolume`.

The data structure is a hierarchical regular grid at its core: Nodes are regular grids, and each grid cell may either store a constant value (this is called a tile), or child pointers. Nodes in VDB trees are wide: Nodes on the first level have a resolution of 32^3 voxels, on the next level 16^3 , and on the leaf level 8^3 voxels. All nodes on a given level have the same resolution. This makes it easy to find the node containing a coordinate using shift operations (see [1]). VDB leaf nodes are implicit in OSPRay / Open VKL: they are stored as pointers to user-provided data.

VDB volumes interpret input data as constant cells (which are then potentially filtered). This is in contrast to `structuredRegular` volumes, which have a vertex-centered interpretation.

The VDB implementation in OSPRay / Open VKL follows the following goals:

- Efficient data structure traversal on vector architectures.
- Enable the use of industry-standard .vdb files created through the OpenVDB library.
- Compatibility with OpenVDB on a leaf data level, so that .vdb file may be loaded with minimal overhead.

VDB volumes have the following parameters:

The `nodesPackedDense` and `nodesPackedTile` together with `node.format` parameters may be provided instead of `node.data`; this packed data




Figure 3.2 – Topology of VDB volumes.

Table 3.10 – Configuration parameters for VDB volumes.

| Type | Name | Description |
|-----------|------------------|---|
| int | maxSamplingDepth | do not descend further than to this depth during sampling, the maximum value and the default is 3 |
| uint32[] | node.level | level on which each input node exists, may be 1, 2 or 3 (levels are counted from the root level = 0 down) |
| vec3i[] | node.origin | the node origin index (per input node) |
| OSPData[] | node.data | data arrays with the node data (per input node). Nodes that are tiles are expected to have single-item arrays. Leaf-nodes with grid data expected to have compact 3D arrays in zyx layout (z changes most quickly) with the correct number of voxels for the level. Only OSP_FLOAT is supported as field OSPDataType. |
| OSPData | nodesPackedDense | optionally provided instead of node.data, a single array of all dense node data in a contiguous zyx layout, provided in the same order as the corresponding node.* parameters |
| OSPData | nodesPackedTile | optionally provided instead of node.data, a single array of all tile node data in a contiguous layout, provided in the same order as the corresponding node.* parameters |
| uint32[] | node.format | for each input node, whether it is of format OSP_VOLUME_FORMAT_DENSE_ZYX (and thus stored in nodesPackedDense), or OSP_VOLUME_FORMAT_TILE (stored in nodesPackedTile) |
| uint | filter | filter used for reconstructing the field, default is OSP_VOLUME_FILTER_LINEAR, alternatively OSP_VOLUME_FILTER_NEAREST, or OSP_VOLUME_FILTER_CUBIC. |
| uint | gradientFilter | filter used for reconstructing the field during gradient computations, default same as filter |
| float | background | value that is used when sampling an undefined region outside the volume domain, default NaN |

layout may provide better performance.

1. Museth, K. VDB: High-Resolution Sparse Volumes with Dynamic Topology.

ACM Transactions on Graphics 32(3), 2013. DOI: 10.1145/2487228.2487235

3.3.6 Particle Volume

Particle volumes consist of a set of points in space. Each point has a position, a radius, and a weight typically associated with an attribute. Particle volumes are created by passing the type string “`particle`” to `ospNewVolume`.

A radial basis function defines the contribution of that particle. Currently, we use the Gaussian radial basis function

$$\phi(P) = w \exp\left(-\frac{(P - p)^2}{2r^2}\right),$$

where P is the particle position, p is the sample position, r is the radius and w is the weight. At each sample, the scalar field value is then computed as the sum of each radial basis function ϕ , for each particle that overlaps it.

The OSPRay / Open VKL implementation is similar to direct evaluation of samples in Reda et al. [2]. It uses an Embree-built BVH with a custom traversal, similar to the method in [1].

Table 3.11 – Configuration parameters for particle volumes.

| Type | Name | Default | Description |
|----------------------|--------------------------------------|-------------------|---|
| <code>vec3f[]</code> | <code>particle.position</code> | | <code>data</code> array of particle positions |
| <code>float[]</code> | <code>particle.radius</code> | | <code>data</code> array of particle radii |
| <code>float[]</code> | <code>particle.weight</code> | <code>NULL</code> | optional <code>data</code> array of particle weights, specifying the height of the kernel. |
| <code>float</code> | <code>radiusSupportFactor</code> | 3.0 | The multiplier of the particle radius required for support. Larger radii ensure smooth results at the cost of performance. In the Gaussian kernel, the radius is one standard deviation (σ), so a value of 3 corresponds to 3σ . |
| <code>float</code> | <code>clampMaxCumulativeValue</code> | 0 | The maximum cumulative value possible, set by user. All cumulative values will be clamped to this, and further traversal (RBF summation) of particle contributions will halt when this value is reached. A value of zero or less turns this off. |
| <code>bool</code> | <code>estimateValueRanges</code> | <code>true</code> | Enable heuristic estimation of value ranges which are used in internal acceleration structures as well as for determining the volume’s overall value range. When set to <code>false</code> , the user <i>must</i> specify <code>clampMaxCumulativeValue</code> , and all value ranges will be assumed $[0 - \text{clampMaxCumulativeValue}]$. Disabling this switch may improve volume commit time, but will make volume rendering less efficient. |

1. A. Knoll, I. Wald, P. Navratil, A. Bowen, K. Reda, M.E., Papka, and K. Gaither, “RBF Volume Ray Casting on Multicore and Manycore CPUs”, 2014, Computer Graphics Forum, 33: 71–80. doi:10.1111/cgf.12363
2. K. Reda, A. Knoll, K. Nomura, M. E. Papka, A. E. Johnson and J. Leigh, “Visualizing large-scale atomistic simulations in ultra-resolution immersive environments”, 2013 IEEE Symposium on Large-Scale Data Analysis and Visualization (LDAV), Atlanta, GA, 2013, pp. 59–65.

3.3.7 Transfer Function

Transfer functions map the scalar values of volumes to color and opacity and thus they can be used to visually emphasize certain features of the volume. To create a new transfer function of given type `type` use

```
OSPTransferFunction ospNewTransferFunction(const char *type);
```

The returned handle can be assigned to a volumetric model (described below) as parameter “`transferFunction`” using `ospSetObject`.

One type of transfer function that is supported by OSPRay is the linear transfer function, which interpolates between given equidistant colors and opacities. It is created by passing the string “`piecewiseLinear`” to `ospNewTransferFunction` and it is controlled by these parameters:

| Type | Name | Description |
|---------|---------|---|
| vec3f[] | color | data array of colors (linear RGB) |
| float[] | opacity | data array of opacities |
| box1f | value | domain (scalar range) this function maps from |

Table 3.12 – Parameters accepted by the linear transfer function.

The arrays `color` and `opacity` can be of different length.

3.3.8 VolumetricModels

Volumes in OSPRay are given volume rendering appearance information through `VolumetricModels`. This decouples the physical representation of the volume (and possible acceleration structures it contains) to rendering-specific parameters (where more than one set may exist concurrently). To create a volume instance, call

```
OSPVolumeModel ospNewVolumetricModel(OSPVolume);
```

The passed volume can be `NULL` as long as the volume to be used is passed as a parameter. If both a volume is specified on object creation and as a parameter, the parameter value is used. If the parameter value is later removed, the volume object passed on object creation is again used.

Table 3.13 – Parameters understood by `VolumetricModel`.

| Type | Name | Default | Description |
|---------------------|------------------|---------|--|
| OSPVolume | volume | | optional <code>volume</code> object this model references |
| OSPTransferFunction | transferFunction | | <code>transfer function</code> to use |
| float | densityScale | 1.0 | makes volumes uniformly thinner or thicker |
| float | anisotropy | 0.0 | anisotropy of the (Henyey-Greenstein) phase function in [-1-1] (<code>path tracer</code> only), default to isotropic scattering |
| uint32 | id | -1u | optional user ID, for <code>framebuffer</code> channel OSP_FB_ID_OBJECT |

3.4 Geometries

Geometries in OSPRay are objects that describe intersectable surfaces. To create a new geometry object of given type `type` use

```
OSPGeometry ospNewGeometry(const char *type);
```

Note that in the current implementation geometries are limited to a maximum of 2^{32} primitives.

3.4.1 Mesh

A mesh consisting of either triangles or quads is created by calling `ospNewGeometry` with type string “mesh”. Once created, a mesh recognizes the following parameters:

Table 3.14 – Parameters defining a mesh geometry.

| Type | Name | Description |
|----------------------------------|-------------------------------------|---|
| <code>vec3f[]</code> | <code>vertex.position</code> | <code>data</code> array of vertex positions, overridden by <code>motion.*</code> arrays |
| <code>vec3f[]</code> | <code>normal</code> | <code>data</code> array of face-varying normals, overridden by <code>motion.*</code> arrays |
| <code>vec3f[]</code> | <code>vertex.normal</code> | <code>data</code> array of vertex-varying normals, overridden by <code>motion.*</code> arrays |
| <code>vec4f[] / vec3f[]</code> | <code>color</code> | <code>data</code> array of face-varying colors (linear RGBA/RGB) |
| <code>vec4f[] / vec3f[]</code> | <code>vertex.color</code> | <code>data</code> array of vertex-varying colors (linear RGBA/RGB) |
| <code>vec2f[]</code> | <code>texcoord</code> | <code>data</code> array of face-varying texture coordinates |
| <code>vec2f[]</code> | <code>vertex.texcoord</code> | <code>data</code> array of vertex-varying texture coordinates |
| <code>vec3ui[] / vec4ui[]</code> | <code>index</code> | <code>data</code> array of (either triangle or quad) indices (into the vertex array(s)) |
| <code>bool</code> | <code>quadSoup</code> | when no explicit <code>index</code> is given, indicates whether to assume a ‘soup’ of quads instead of triangles, default false |
| <code>vec3f[][]</code> | <code>motion.vertex.position</code> | <code>data</code> array of vertex position arrays (uniformly distributed keys for deformation motion blur) |
| <code>vec3f[][]</code> | <code>motion.normal</code> | <code>data</code> array of face-varying normal arrays (uniformly distributed keys for deformation motion blur) |
| <code>vec3f[][]</code> | <code>motion.vertex.normal</code> | <code>data</code> array of vertex-varying normal arrays (uniformly distributed keys for deformation motion blur) |
| <code>box1f</code> | <code>time</code> | time associated with first and last key in <code>motion.*</code> arrays (for deformation motion blur), default [0, 1] |

The data type of index arrays differentiates between the underlying geometry, triangles are used for a index with `vec3ui` type and quads for `vec4ui` type. Quads are internally handled as a pair of two triangles, thus mixing triangles and quads is supported by encoding some triangle as a quad with the last two vertex indices being identical ($w=z$).

The `vertex.position` array is mandatory to create a valid mesh.

The `index` array is optional. If none is provided, a ‘triangle soup’ is assumed, i.e., each three consecutive vertices form one triangle; unless the boolean `quadSoup` is set to true, then a ‘quad soup’ is assumed i.e., each four subsequent vertices form one quad. If the size of the `vertex.position` array is not a multiple of three for triangles or four for quads, the remainder vertices are ignored.

Face-varying attributes (`normal`, `motion.normal`, `color`, `texcoord`) map unique values to each vertex of a primitive/face (triangle or quad), thus attributes

can be different for the same vertex that is shared by multiple primitives. Essentially, face-varying attributes are a ‘attribute soup’ and behave similar to the implicit index, the size of the array must be at least three times the number of triangles or four times the number of quads, respectively. Face-varying attributes take precedence over the respective vertex attributes (`vertex.normal`, `motion.vertex.normal`, `vertex.color`, `vertex.texcoord`) when both arrays of the same attribute are present.

3.4.2 Subdivision

A mesh consisting of subdivision surfaces, created by specifying a geometry of type “`subdivision`”. Once created, a subdivision recognizes the following parameters:

Table 3.15 – Parameters defining a Subdivision geometry.

| Type | Name | Description |
|----------------------|----------------------------------|---|
| <code>vec3f[]</code> | <code>vertex.position</code> | <code>data</code> array of vertex positions |
| <code>vec4f[]</code> | <code>color</code> | optional <code>data</code> array of face-varying colors (linear RGBA) |
| <code>vec4f[]</code> | <code>vertex.color</code> | optional <code>data</code> array of vertex-varying colors (linear RGBA) |
| <code>vec2f[]</code> | <code>texcoord</code> | optional <code>data</code> array of vertex-varying texture coordinates |
| <code>vec2f[]</code> | <code>vertex.texcoord</code> | optional <code>data</code> array of vertex-varying texture coordinates |
| <code>float</code> | <code>level</code> | global level of tessellation, default 5 |
| <code>uint[]</code> | <code>index</code> | <code>data</code> array of indices (into the vertex array(s)) |
| <code>float[]</code> | <code>index.level</code> | optional <code>data</code> array of per-edge levels of tessellation, overrides global level |
| <code>uint[]</code> | <code>face</code> | optional <code>data</code> array holding the number of indices/edges (3 to 15) per face, defaults to 4 (a pure quad mesh) |
| <code>vec2i[]</code> | <code>edgeCrease.index</code> | optional <code>data</code> array of edge crease indices |
| <code>float[]</code> | <code>edgeCrease.weight</code> | optional <code>data</code> array of edge crease weights |
| <code>uint[]</code> | <code>vertexCrease.index</code> | optional <code>data</code> array of vertex crease indices |
| <code>float[]</code> | <code>vertexCrease.weight</code> | optional <code>data</code> array of vertex crease weights |
| <code>uint</code> | <code>mode</code> | <code>OSPSubdivisionMode</code> subdivision edge boundary mode, supported modes are: <code>OSP_SUBDIVISION_NO_BOUNDARY</code> <code>OSP_SUBDIVISION_SMOOTH_BOUNDARY</code> (default) <code>OSP_SUBDIVISION_PIN_CORNERS</code> <code>OSP_SUBDIVISION_PIN_BOUNDARY</code> <code>OSP_SUBDIVISION_PIN_ALL</code> |

The `vertex` and `index` arrays are mandatory to create a valid subdivision surface. If no `face` array is present then a pure quad mesh is assumed (the number of indices must be a multiple of 4). Optionally supported are edge and vertex creases.

3.4.3 Spheres

A geometry consisting of individual spheres, each of which can have an own radius, is created by calling `ospNewGeometry` with type string “`sphere`”. The spheres will not be tessellated but rendered procedurally and are thus perfectly round. To allow a variety of sphere representations in the application this geometry allows a flexible way of specifying the data of center position and radius within a `data` array:

Table 3.16 – Parameters defining a spheres geometry.

| Type | Name | Default | Description |
|---------|-----------------|---------|---|
| vec3f[] | sphere.position | | data array of center positions |
| float[] | sphere.radius | NULL | optional data array of the per-sphere radius |
| vec3f[] | sphere.normal | NULL | optional data array of normals (only for “oriented disc”) |
| vec2f[] | sphere.texcoord | NULL | optional data array of texture coordinates (constant per sphere) |
| float | radius | 0.01 | default radius for all spheres (if <code>sphere.radius</code> is not set) |
| uint | type | | <code>OSPSphereType</code> for rendering the sphere. Supported types are: <code>OSP_SPHERE</code> (default) <code>OSP_DISC</code> <code>OSP_ORIENTED_DISC</code> |

3.4.4 Curves

A geometry consisting of multiple curves is created by calling `ospNewGeometry` with type string “curve”. The parameters defining this geometry are listed in the table below.

Table 3.17 – Parameters defining a curves geometry.

| Type | Name | Description |
|----------|------------------------|---|
| vec4f[] | vertex.position_radius | data array of vertex position and per-vertex radius |
| vec2f[] | vertex.texcoord | data array of per-vertex texture coordinates |
| vec4f[] | vertex.color | data array of corresponding vertex colors (linear RGBA) |
| vec3f[] | vertex.normal | data array of curve normals (only for “ribbon” curves) |
| vec4f[] | vertex.tangent | data array of curve tangents (only for “hermite” curves) |
| uint32[] | index | data array of indices to the first vertex or tangent of a curve segment |
| uint | type | <code>OSPCurveType</code> for rendering the curve. Supported types are: <code>OSP_FLAT</code> <code>OSP_ROUND</code> <code>OSP_RIBBON</code> <code>OSP_DISJOINT</code> |
| uint | basis | <code>OSPCurveBasis</code> for defining the curve. Supported bases are: <code>OSP_LINEAR</code> <code>OSP_BEZIER</code> <code>OSP_BSPLINE</code> <code>OSP_HERMITE</code> <code>OSP_CATMULL_ROM</code> |

Positions in `vertex.position_radius` parameter supports per-vertex varying radii with data type `vec4f[]` and instantiate Embree curves internally for the relevant type/basis mapping.

The following section describes the properties of different curve basis’ and how they use the data provided in data buffers:

OSP_LINEAR The indices point to the first of 2 consecutive control points in the vertex buffer. The first control point is the start and the second control point the end of the line segment. The curve goes through all control points listed in the vertex buffer.

OSP_BEZIER The indices point to the first of 4 consecutive control points in the vertex buffer. The first control point represents the start point of the curve, and the 4th control point the end point of the curve. The Bézier basis is interpolating, thus the curve does go exactly through the first and fourth control vertex.

OSP_BSPLINE The indices point to the first of 4 consecutive control points in the vertex buffer. This basis is not interpolating, thus the curve does in general not go through any of the control points directly. Using this basis, 3 control points can be shared for two continuous neighboring curve segments, e.g., the curves (p_0, p_1, p_2, p_3) and (p_1, p_2, p_3, p_4) are C1 continuous. This feature make this basis a good choice to construct continuous multi-segment curves, as memory consumption can be kept minimal.

OSP_HERMITE It is necessary to have both vertex buffer and tangent buffer for using this basis. The indices point to the first of 2 consecutive points in the vertex buffer, and the first of 2 consecutive tangents in the tangent buffer. This basis is interpolating, thus does exactly go through the first and second control point, and the first order derivative at the begin and end matches exactly the value specified in the tangent buffer. When connecting two segments continuously, the end point and tangent of the previous segment can be shared.

OSP_CATMULL_ROM The indices point to the first of 4 consecutive control points in the vertex buffer. If (p_0, p_1, p_2, p_3) represent the points then this basis goes through p_1 and p_2 , with tangents as $(p_2 - p_0)/2$ and $(p_3 - p_1)/2$.

The following section describes the properties of different curve types' and how they define the geometry of a curve:

OSP_FLAT This type enables faster rendering as the curve is rendered as a connected sequence of ray facing quads.

OSP_ROUND This type enables rendering a real geometric surface for the curve which allows closeup views. This mode renders a sweep surface by sweeping a varying radius circle tangential along the curve.

OSP_RIBBON The type enables normal orientation of the curve and requires a normal buffer be specified along with vertex buffer. The curve is rendered as a flat band whose center approximately follows the provided vertex buffer and whose normal orientation approximately follows the provided normal buffer. Not supported for basis **OSP_LINEAR**.

OSP_DISJOINT Only supported for basis **OSP_LINEAR**; the segments are open and not connected at the joints, i.e., the curve segments are either individual cones or cylinders.

3.4.5 Boxes

OSPRay can directly render axis-aligned bounding boxes without the need to convert them to quads or triangles. To do so create a boxes geometry by calling `ospNewGeometry` with type string "box".

| Type | Name | Description |
|----------------------|------------------|----------------------------------|
| <code>box3f[]</code> | <code>box</code> | <code>data</code> array of boxes |

Table 3.18 – Parameters defining a boxes geometry.

3.4.6 Planes

OSPRay can directly render planes defined by plane equation coefficients in its implicit form $ax + by + cz + d = 0$. By default planes are infinite but their extents

can be limited by defining optional bounding boxes. A planes geometry can be created by calling `ospNewGeometry` with type string “plane”.

| Type | Name | Description |
|----------------------|---------------------------------|--|
| <code>vec4f[]</code> | <code>plane.coefficients</code> | <code>data</code> array of plane coefficients (a, b, c, d) |
| <code>box3f[]</code> | <code>plane.bounds</code> | optional <code>data</code> array of bounding boxes |

Table 3.19 – Parameters defining a planes geometry.

3.4.7 Isosurfaces

OSPRay can directly render multiple isosurfaces of a volume without first tessellating them. To do so create an isosurfaces geometry by calling `ospNewGeometry` with type string “isosurface”. The appearance information of the surfaces is set through the Geometric Model. Per-isosurface colors can be set by passing per-primitive colors to the Geometric Model, in order of the isosurface array.

| Type | Name | Description |
|------------------------|-----------------------|---|
| <code>float</code> | <code>isovalue</code> | single isovalue |
| <code>float[]</code> | <code>isovalue</code> | <code>data</code> array of isovalue |
| <code>OSPVolume</code> | <code>volume</code> | handle of the <code>Volume</code> to be isosurfaced |

Table 3.20 – Parameters defining an isosurfaces geometry.

3.4.8 GeometricModels

Geometries are matched with surface appearance information through GeometricModels. These take a geometry, which defines the surface representation, and applies either full-object or per-primitive color and material information. To create a geometric model, call

```
OSPGeometricModel ospNewGeometricModel(OSPGeometry);
```

The passed geometry can be NULL as long as the geometry to be used is passed as a parameter. If both a geometry is specified on object creation and as a parameter, the parameter value is used. If the parameter value is later removed, the geometry object passed on object creation is again used.

Color and material are fetched with the primitive ID of the hit (clamped to the valid range, thus a single color or material is fine), or mapped first via the `index` array (if present). All parameters are optional, however, some renderers (notably the `path tracer`) require a material to be set. Materials are either handles of `OSPMaterial`, or indices into the `material` array on the `renderer`, which allows to build a `world` which can be used by different types of renderers.

An `invertNormals` flag allows to invert (shading) normal vectors of the rendered geometry. That is particularly useful for clipping. By changing normal vectors orientation one can control whether inside or outside of the clipping geometry is being removed. For example, a clipping geometry with normals oriented outside clips everything what's inside.

Table 3.21 – Parameters understood by GeometricModel.

| Type | Name | Description |
|---|---------------|--|
| OSPGeometry | geometry | optional <code>geometry</code> object this model references |
| OSPMaterial / OSPMaterial[] / uint32 / uint32[] | material | optional (<code>data</code> array of per-primitive) <code>material</code> , may be an index into the <code>material</code> parameter on the renderer (if it exists) |
| vec4f / vec4f[] | color | optional (<code>data</code> array of per-primitive) color assigned to the geometry (linear RGBA) |
| uint8[] | index | optional <code>data</code> array of per-primitive indices into <code>color</code> and <code>material</code> |
| bool | invertNormals | inverts all shading normals (<code>Ns</code>), default false |
| uint32 | id | optional user ID, for <code>framebuffer</code> channel <code>OSP_FB_ID_OBJECT</code> , default -1u |

Table 3.22 – Parameters accepted by all lights.

| Type | Name | Default | Description |
|-------|-------------------|---------|---|
| vec3f | color | white | color of the light (linear RGB) |
| float | intensity | 1 | intensity of the light (a factor) |
| uint | intensityQuantity | | <code>OSPIIntensityQuantity</code> to set the radiometric quantity represented by <code>intensity</code> . The default value depends on the light source. |
| bool | visible | true | whether the light can be directly seen |

3.5 Lights

To create a new light source of given type `type` use

```
OSPLight ospNewLight(const char *type);
```

All light sources accept the following parameters:

In OSPRay the `intensity` parameter of a light source can correspond to different types of radiometric quantities. The type of the value represented by a light's `intensity` parameter is set using `intensityQuantity`, which accepts values from the enum type `OSPIIntensityQuantity`. The supported types of `OSPIIntensityQuantity` differ between the different light sources (see documentation of each specific light source).

3.5.1 Photometric Lights

Measured light sources (IES, EULUMDAT, ...) are supported by the `sphere`, `spot`, and `quad` lights when setting an `intensityDistribution` `data` array to modulate the intensity per direction. The mapping is using the C-γ coordinate system (see also below figure): the values of the first (or only) dimension of `intensityDistribution` are uniformly mapped to γ in [0–π]; the first intensity value to 0, the last value to π, thus at least two values need to be present.

If the array has a second dimension then the intensities are not rotational symmetric around the main direction (where angle γ is zero), but are accordingly mapped to the C-halfplanes in [0–2π]; the first “row” of values to 0 and 2π, the other rows such that they have uniform distance to its neighbors. The orientation of the C0-plane is specified via `c0`.

Table 3.23 – Types of radiometric quantities used to interpret a light’s `intensity` parameter.

| Name | Description |
|--|--|
| <code>OSP_INTENSITY_QUANTITY_POWER</code> | the overall amount of light energy emitted by the light source into the scene, unit is W |
| <code>OSP_INTENSITY_QUANTITY_INTENSITY</code> | the overall amount of light emitted by the light in a given direction, unit is W/sr |
| <code>OSP_INTENSITY_QUANTITY_RADIANCE</code> | the amount of light emitted by a point on the light source in a given direction, unit is W/sr/m ² |
| <code>OSP_INTENSITY_QUANTITY_IRRADIANCE</code> | the amount of light arriving at a surface point, assuming the light is oriented towards to the surface, unit is W/m ² |
| <code>OSP_INTENSITY_QUANTITY_SCALE</code> | a linear scaling factor for light sources with a built-in quantity (e.g., <code>HDRI</code> , or <code>sunSky</code> , or when using <code>intensityDistribution</code>). |




Figure 3.3 – C- γ coordinate system for the mapping of `intensityDistribution` with photometric lights.

Table 3.24 – Special parameters for photometric lights.

| Type | Name | Description |
|---------|------------------------------------|---|
| float[] | <code>intensityDistribution</code> | luminous intensity distribution for photometric lights; can be 2D for asymmetric illumination; values are assumed to be uniformly distributed |
| vec3f | <code>c0</code> | orientation, i.e., direction of the C0-(half)plane (only needed if illumination via <code>intensityDistribution</code> is asymmetric) |

When using an `intensityDistribution` then the default and only valid value for `intensityQuantity` is `OSP_INTENSITY_QUANTITY_SCALE`.

The following light types are supported by most OSPRay renderers.

3.5.2 Directional Light / Distant Light

The distant light (or traditionally the directional light) is thought to be far away (outside of the scene), thus its light arrives (almost) as parallel rays. It is created by passing the type string “`distant`” to `ospNewLight`. The distant light supports `OSP_INTENSITY_QUANTITY_RADIANCE` and `OSP_INTENSITY_QUANTITY_IRRADIANCE` (default) as `intensityQuantity` parameter value. In addition to the [general parameters](#) understood by all lights the distant light supports the following special parameters:

Setting the angular diameter to a value greater than zero will result in soft shadows when the renderer uses stochastic sampling (like the [path tracer](#)). For instance, the apparent size of the sun is about 0.53°.

| Type | Name | Default | Description |
|-------|-----------------|-----------|--|
| vec3f | direction | (0, 0, 1) | main emission direction of the distant light |
| float | angularDiameter | 0 | apparent size (angle in degree) of the light |

Table 3.25 – Special parameters accepted by the distant light.

3.5.3 Point Light / Sphere Light

The sphere light (or the special case point light) is a light emitting uniformly in all directions from the surface toward the outside. It does not emit any light toward the inside of the sphere. It is created by passing the type string “sphere” to `ospNewLight`. The point light supports only `OSP_INTENSITY_QUANTITY_SCALE` when `intensityDistribution` is set, or otherwise `OSP_INTENSITY_QUANTITY_POWER`, `OSP_INTENSITY_QUANTITY_INTENSITY` (then default) and `OSP_INTENSITY_QUANTITY_RADIANCE` as `intensityQuantity` parameter value. In addition to the [general parameters](#) understood by all lights and the [photometric parameters](#) the sphere light supports the following special parameters:

| Type | Name | Default | Description |
|-------|-----------|-----------|--|
| vec3f | position | (0, 0, 0) | the center of the sphere light |
| float | radius | 0 | the size of the sphere light |
| vec3f | direction | (0, 0, 1) | main orientation of <code>intensityDistribution</code> |

Table 3.26 – Special parameters accepted by the sphere light.

Setting the radius to a value greater than zero will result in soft shadows when the renderer uses stochastic sampling (like the [path tracer](#)).

3.5.4 Spotlight / Ring Light

The spotlight is a light emitting into a cone of directions. It is created by passing the type string “spot” to `ospNewLight`. The spotlight supports only `OSP_INTENSITY_QUANTITY_SCALE` when `intensityDistribution` is set, or otherwise `OSP_INTENSITY_QUANTITY_POWER`, `OSP_INTENSITY_QUANTITY_INTENSITY` (then default) and `OSP_INTENSITY_QUANTITY_RADIANCE` as `intensityQuantity` parameter value. In addition to the [general parameters](#) understood by all lights and the [photometric parameters](#) the spotlight supports the special parameters listed in the table.




Figure 3.4 – Angles used by the spotlight.

Setting the radius to a value greater than zero will result in soft shadows when the renderer uses stochastic sampling (like the [path tracer](#)). Additionally setting the inner radius will result in a ring instead of a disk emitting the light.

Table 3.27 – Special parameters accepted by the spotlight.


| Type | Name | Default | Description |
|-------|---------------|-----------|---|
| vec3f | position | (0, 0, 0) | the center of the spotlight |
| vec3f | direction | (0, 0, 1) | main emission direction of the spot |
| float | openingAngle | 180 | full opening angle (in degree) of the spot; outside of this cone is no illumination |
| float | penumbraAngle | 5 | size (angle in degree) of the “penumbra”, the region between the rim (of the illumination cone) and full intensity of the spot; should be smaller than half of openingAngle |
| float | radius | 0 | the size of the spotlight, the radius of a disk with normal direction |
| float | innerRadius | 0 | in combination with radius turns the disk into a ring |

3.5.5 Quad Light

The quad³ light is a planar, procedural area light source emitting uniformly on one side into the half-space. It is created by passing the type string “quad” to `ospNewLight`. The quad light supports only `OSP_INTENSITY_QUANTITY_SCALE` when `intensityDistribution` is set, or otherwise `OSP_INTENSITY_QUANTITY_POWER`, `OSP_INTENSITY_QUANTITY_INTENSITY` and `OSP_INTENSITY_QUANTITY_RADIANCE` (then default) as `intensityQuantity` parameter. In addition to the [general parameters](#) understood by all lights and the [photometric parameters](#) the quad light supports the following special parameters:

³ actually a parallelogram

| Type | Name | Default | Description |
|-------|----------|-----------|--|
| vec3f | position | (0, 0, 0) | position of one vertex of the quad light |
| vec3f | edge1 | (1, 0, 0) | vector to one adjacent vertex |
| vec3f | edge2 | (0, 1, 0) | vector to the other adjacent vertex |

Table 3.28 – Special parameters accepted by the quad light.**Figure 3.5** – Defining a quad light which emits toward the reader.

The emission side is determined by the cross product of $\text{edge1} \times \text{edge2}$, which is also the main emission direction for `intensityDistribution`. Note that only renderers that use stochastic sampling (like the path tracer) will compute soft shadows from the quad light. Other renderers will just sample the center of the quad light, which results in hard shadows.

3.5.6 Cylinder Light

The cylinder light is a cylindrical, procedural area light source emitting uniformly outwardly into the space beyond the boundary. It is created by passing the type string “cylinder” to `ospNewLight`. The cylinder light supports `OSP_INTENSITY_QUANTITY_POWER`, `OSP_INTENSITY_QUANTITY_INTENSITY` and `OSP_INTENSITY_QUANTITY_RADIANCE` (default) as `intensityQuantity` parameter.

`tyQuantity` parameter. In addition to the [general parameters](#) understood by all lights the cylinder light supports the following special parameters:

| Type | Name | Default | Description |
|-------|-----------|-----------|---------------------------------------|
| vec3f | position0 | (0, 0, 0) | position of the start of the cylinder |
| vec3f | position1 | (0, 0, 1) | position of the end of the cylinder |
| float | radius | 1 | radius of the cylinder |

Table 3.29 – Special parameters accepted by the cylinder light.

Note that only renderers that use stochastic sampling (like the path tracer) will compute soft shadows from the cylinder light. Other renderers will just sample the closest point on the cylinder light, which results in hard shadows.

3.5.7 HDRI Light

The HDRI light is a textured light source surrounding the scene and illuminating it from infinity. It is created by passing the type string “hdri” to `ospNewLight`. The values of the HDRI correspond to radiance and therefore the HDRI light only accepts `OSP_INTENSITY_QUANTITY_SCALE` as `intensityQuantity` parameter value. In addition to the [general parameters](#) the HDRI light supports the following special parameters:

Table 3.30 – Special parameters accepted by the HDRI light.

| Type | Name | Default | Description |
|------------|-----------|-----------|--|
| vec3f | up | (0, 1, 0) | up direction of the light |
| vec3f | direction | (0, 0, 1) | direction to which the center of the texture will be mapped to (analog to panoramic camera) |
| OSPTexture | map | | environment map in latitude / longitude format |




Figure 3.6 – Orientation and Mapping of an HDRI Light.

Note that the [SciVis renderer](#) only shows the HDRI light in the background (like an environment map) without computing illumination of the scene.

3.5.8 Ambient Light

The ambient light surrounds the scene and illuminates it from infinity with constant radiance (determined by combining the [parameters color and intensity](#)). It is created by passing the type string “ambient” to `ospNewLight`. The ambient light supports `OSP_INTENSITY_QUANTITY_RADIANC` and `OSP_INTENSITY_QUANTITY_IRRADIANCE` (default) as `intensityQuantity` parameter value.

Note that the [SciVis renderer](#) uses ambient lights to control the color and intensity of the computed ambient occlusion (AO).

3.5.9 Sun-Sky Light

The sun-sky light is a combination of a *distant* light for the sun and a procedural *hdri* light for the sky. It is created by passing the type string “`sunSky`” to `ospNewLight`. The sun-sky light surrounds the scene and illuminates it from infinity and can be used for rendering outdoor scenes. The radiance values are calculated using the Hošek-Wilkie sky model and solar radiance function. The underlying model of the sun-sky light returns radiance values and therefore the light only accepts `OSP_INTENSITY_QUANTITY_SCALE` as `intensityQuantity` parameter value. To rescale the returned radiance of the sky model the default value for the `intensity` parameter is set to `0.025`. In addition to the [general parameters](#) the following special parameters are supported:

Table 3.31 – Special parameters accepted by the `sunSky` light.

| Type | Name | Default | Description |
|--------------------|-------------------------------|-------------------------|--|
| <code>vec3f</code> | <code>up</code> | <code>(0, 1, 0)</code> | zenith of sky |
| <code>vec3f</code> | <code>direction</code> | <code>(0, -1, 0)</code> | main emission direction of the sun |
| <code>float</code> | <code>turbidity</code> | 3 | atmospheric turbidity due to particles, in [1–10] |
| <code>float</code> | <code>albedo</code> | 0.3 | ground reflectance, in [0–1] |
| <code>float</code> | <code>horizonExtension</code> | 0.01 | extend the sky dome by stretching the horizon, fraction of the lower hemisphere to cover, in [0–1] |

The lowest elevation for the sun is restricted to the horizon.

Note that the [SciVis renderer](#) only computes illumination from the sun (yet the sky is still shown in the background, like an environment map).

3.5.10 Emissive Objects

The [path tracer](#) will consider illumination by [geometries](#) which have a light emitting material assigned (for example the [Luminous](#) or [Principled](#) material).

3.6 Materials

Materials describe how light interacts with surfaces, they give objects their distinctive look. To create a new material of given type `type` call

```
OSPMaterial ospNewMaterial(const char *material_type);
```

The returned handle can then be used to assign the material to a given geometry with

```
void ospSetObject(OSPGeometricModel, "material", OSPMaterial);
```

3.6.1 OBJ Material

The OBJ material is the workhorse material supported by both the [SciVis renderer](#) and the [path tracer](#) (the [Ambient Occlusion renderer](#) only uses the `kd` and `d` parameter). It offers widely used common properties like diffuse and specular reflection and is based on the [MTL material format](#) of Lightwave’s OBJ scene files. To create an OBJ material pass the type string “`obj`” to `ospNewMaterial`. Its main parameters are

In particular when using the path tracer it is important to adhere to the principle of energy conservation, i.e., that the amount of light reflected by a surface is not larger than the light arriving. Therefore the path tracer issues a warning

| Type | Name | Default | Description |
|------------|----------|-----------|---|
| vec3f | kd | white 0.8 | diffuse color (linear RGB) |
| vec3f | ks | black | specular color (linear RGB) |
| float | ns | 10 | shininess (Phong exponent), usually in $[2-10^4]$ |
| float | d | opaque | opacity |
| vec3f | tf | black | transparency filter color (linear RGB) |
| OSPTexture | map_bump | NULL | normal map |

Table 3.32 – Main parameters of the OBJ material.

and renormalizes the color parameters if the sum of `kd`, `ks`, and `tf` is larger than one in any color channel. Similarly important to mention is that almost all materials of the real world reflect at most only about 80% of the incoming light. So even for a white sheet of paper or white wall paint do better not set `kd` larger than 0.8; otherwise rendering times are unnecessary long and the contrast in the final images is low (for example, the corners of a white room would hardly be discernible, as can be seen in the figure below).




Figure 3.7 – Comparison of diffuse rooms with 100% reflecting white paint (left) and realistic 80% reflecting white paint (right), which leads to higher overall contrast. Note that exposure has been adjusted to achieve similar brightness levels.

If present, the color component of [geometries](#) is also used for the diffuse color `kd` and the alpha component is also used for the opacity `d`.

Normal mapping can simulate small geometric features via the texture `map_bump`. The normals n in the normal map are with respect to the local tangential shading coordinate system and are encoded as $\frac{1}{2}(n + 1)$, thus a texel $(0.5, 0.5, 1)$ ⁴ represents the unperturbed shading normal $(0, 0, 1)$. Because of this encoding an sRGB gamma [texture](#) format is ignored and normals are always fetched as linear from a normal map. Note that the orientation of normal maps is important for a visually consistent look: by convention OSPRay uses a coordinate system with the origin in the lower left corner; thus a convexity will look green toward the top of the texture image (see also the example image of a normal map). If this is not the case flip the normal map vertically or invert its green channel.

Note that `tf` colored transparency is implemented in the SciVis and the path tracer but normal mapping with `map_bump` is currently supported in the path tracer only.

All parameters (except `tf`) can be textured by passing a [texture](#) handle, pre-fixed with “`map_`”. The fetched texels are multiplied by the respective parameter value. If only the texture is given (but not the corresponding parameter), only the texture is used (the default value of the parameter is *not* multiplied). The color textures `map_kd` and `map_ks` are typically in one of the sRGB gamma encoded formats, whereas textures `map_ns` and `map_d` are usually in a linear format (and only the first component is used). Additionally, all textures support [texture transformations](#).

⁴ respectively $(127, 127, 255)$ for 8 bit textures and $(32767, 32767, 65535)$ for 16 bit textures

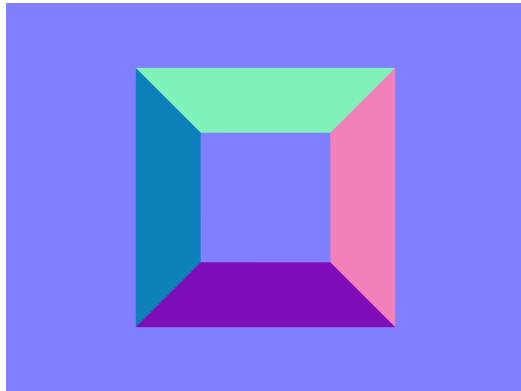


Figure 3.8 – Normal map representing an exalted square pyramidal frustum.



Figure 3.9 – Rendering of a OBJ material with wood textures.

3.6.2 Principled

The Principled material is the most complex material offered by the [path tracer](#), which is capable of producing a wide variety of materials (e.g., plastic, metal, wood, glass) by combining multiple different layers and lobes. It uses the GGX microfacet distribution with approximate multiple scattering for dielectrics and metals, uses the Oren-Nayar model for diffuse reflection, and is energy conserving. To create a Principled material, pass the type string “principled” to `osrNewMaterial`. Its parameters are listed in the table below.

All parameters can be textured by passing a [texture](#) handle, prefixed with “`map_`” (e.g., “`map_baseColor`”). [texture transformations](#) are supported as well.




Figure 3.10 – Rendering of a Principled coated brushed metal material with textured anisotropic rotation and a dust layer (sheen) on top.

3.6.3 CarPaint

The CarPaint material is a specialized version of the Principled material for rendering different types of car paints. To create a CarPaint material, pass the type

Table 3.33 – Parameters of the Principled material.

| Type | Name | Default | Description |
|-------|-------------------|-----------|---|
| vec3f | baseColor | white 0.8 | base reflectivity (diffuse and/or metallic, linear RGB) |
| vec3f | edgeColor | white | edge tint (metallic only, linear RGB) |
| float | metallic | 0 | mix between dielectric (diffuse and/or specular) and metallic (specular only with complex IOR) in [0–1] |
| float | diffuse | 1 | diffuse reflection weight in [0–1] |
| float | specular | 1 | specular reflection/transmission weight in [0–1] |
| bool | specularMetallic | true | whether specular influences metallic |
| float | ior | 1 | dielectric index of refraction |
| float | transmission | 0 | specular transmission weight in [0–1] |
| vec3f | transmissionColor | white | attenuated color due to transmission (Beer's law, linear RGB) |
| float | transmissionDepth | 1 | distance at which color attenuation is equal to transmissionColor |
| float | roughness | 0 | diffuse and specular roughness in [0–1], 0 is perfectly smooth |
| float | anisotropy | 0 | amount of specular anisotropy in [0–1] |
| float | rotation | 0 | rotation of the direction of anisotropy in [0–1], 1 is going full circle |
| float | normal | 1 | default normal map/scale for all layers |
| float | baseNormal | 1 | base normal map/scale (overrides default normal) |
| bool | thin | false | flag specifying whether the material is thin or solid |
| float | thickness | 1 | thickness of the material (thin only), affects the amount of color attenuation due to specular transmission |
| float | backlight | 0 | amount of diffuse transmission (thin only) in [0–2], 1 is 50% reflection and 50% transmission, 2 is transmission only |
| float | coat | 0 | clear coat layer weight in [0–1] |
| float | coatIor | 1.5 | clear coat index of refraction |
| vec3f | coatColor | white | clear coat color tint (linear RGB) |
| float | coatThickness | 1 | clear coat thickness, affects the amount of color attenuation |
| float | coatRoughness | 0 | clear coat roughness in [0–1], 0 is perfectly smooth |
| float | coatNormal | 1 | clear coat normal map/scale (overrides default normal) |
| float | sheen | 0 | sheen layer weight in [0–1] |
| vec3f | sheenColor | white | sheen color tint (linear RGB) |
| float | sheenTint | 0 | how much sheen is tinted from sheenColor toward baseColor |
| float | sheenRoughness | 0.2 | sheen roughness in [0–1], 0 is perfectly smooth |
| float | opacity | 1 | cut-out opacity/transparency, 1 is fully opaque |
| vec3f | emissiveColor | black | color (and intensity) of the emitted light |

string “carPaint” to `ospNewMaterial`. Its parameters are listed in the table below.

All parameters can be textured by passing a `texture` handle, prefixed with “`map_`” (e.g., “`map_baseColor`”). `texture transformations` are supported as well.

3.6.4 Metal

The `path tracer` offers a physical metal, supporting changing roughness and realistic color shifts at edges. To create a Metal material pass the type string “`metal`” to `ospNewMaterial`. Its parameters are

Table 3.34 – Parameters of the CarPaint material.

| Type | Name | Default | Description |
|-------|-----------------|-----------|--|
| vec3f | baseColor | white 0.8 | diffuse base reflectivity (linear RGB) |
| float | roughness | 0 | diffuse roughness in [0–1], 0 is perfectly smooth |
| float | normal | 1 | normal map/scale |
| vec3f | flakeColor | Aluminium | color of metallic flakes (linear RGB) |
| float | flakeDensity | 0 | density of metallic flakes in [0–1], 0 disables flakes, 1 fully covers the surface with flakes |
| float | flakeScale | 100 | scale of the flake structure, higher values increase the amount of flakes |
| float | flakeSpread | 0.3 | flake spread in [0–1] |
| float | flakeJitter | 0.75 | flake randomness in [0–1] |
| float | flakeRoughness | 0.3 | flake roughness in [0–1], 0 is perfectly smooth |
| float | coat | 1 | clear coat layer weight in [0–1] |
| float | coatIor | 1.5 | clear coat index of refraction |
| vec3f | coatColor | white | clear coat color tint (linear RGB) |
| float | coatThickness | 1 | clear coat thickness, affects the amount of color attenuation |
| float | coatRoughness | 0 | clear coat roughness in [0–1], 0 is perfectly smooth |
| float | coatNormal | 1 | clear coat normal map/scale |
| vec3f | flipflopColor | white | reflectivity of coated flakes at grazing angle, used together with coatColor produces a pearlescent paint (linear RGB) |
| float | flipflopFalloff | 1 | flip flop color falloff, 1 disables the flip flop effect |

Table 3.35 – Parameters of the Metal material.

| Type | Name | Default | Description |
|---------|-----------|-----------|--|
| vec3f[] | ior | Aluminium | <code>data</code> array of spectral samples of complex refractive index, each entry in the form (wavelength, eta, k), ordered by wavelength (which is in nm) |
| vec3f | eta | | RGB complex refractive index, real part |
| vec3f | k | | RGB complex refractive index, imaginary part |
| float | roughness | 0.1 | roughness in [0–1], 0 is perfect mirror |

The main appearance (mostly the color) of the Metal material is controlled by the physical parameters `eta` and `k`, the wavelength-dependent, complex index of refraction. These coefficients are quite counter-intuitive but can be found in [published measurements](#). For accuracy the index of refraction can be given as an array of spectral samples in `ior`, each sample a triplet of wavelength (in nm), eta, and k, ordered monotonically increasing by wavelength; OSPRay will then calculate the Fresnel in the spectral domain. Alternatively, `eta` and `k` can also be specified as approximated RGB coefficients; some examples are given in below table.

The `roughness` parameter controls the variation of microfacets and thus how polished the metal will look. The roughness can be modified by a `texture` `map_roughness` (`texture transformations` are supported as well) to create notable edging effects.




Figure 3.11 – Rendering of a pearlescent CarPaint material.

| Metal | eta | k |
|---------------|-----------------------|-----------------|
| Ag, Silver | (0.051, 0.043, 0.041) | (5.3, 3.6, 2.3) |
| Al, Aluminium | (1.5, 0.98, 0.6) | (7.6, 6.6, 5.4) |
| Au, Gold | (0.07, 0.37, 1.5) | (3.7, 2.3, 1.7) |
| Cr, Chromium | (3.2, 3.1, 2.3) | (3.3, 3.3, 3.1) |
| Cu, Copper | (0.1, 0.8, 1.1) | (3.5, 2.5, 2.4) |

Table 3.36 – Index of refraction of selected metals as approximated RGB coefficients, based on data from <https://refractiveindex.info/>.

3.6.5 Alloy

The [path tracer](#) offers an alloy material, which behaves similar to [Metal](#), but allows for more intuitive and flexible control of the color. To create an Alloy material pass the type string “alloy” to `ospNewMaterial`. Its parameters are

| Type | Name | Default | Description |
|-------|-----------|-----------|---|
| vec3f | color | white 0.9 | reflectivity at normal incidence (0 degree, linear RGB) |
| vec3f | edgeColor | white | reflectivity at grazing angle (90 degree, linear RGB) |
| float | roughness | 0.1 | roughness, in [0–1], 0 is perfect mirror |

Table 3.37 – Parameters of the Alloy material.

The main appearance of the Alloy material is controlled by the parameter `color`, while `edgeColor` influences the tint of reflections when seen at grazing angles (for real metals this is always 100% white). If present, the color component of `geometries` is also used for reflectivity at normal incidence `color`. As in [Metal](#) the `roughness` parameter controls the variation of microfacets and thus how polished the alloy will look. All parameters can be textured by passing a [texture](#) handle, prefixed with “`map_`”; [texture transformations](#) are supported as well.

3.6.6 Glass

The [path tracer](#) offers a realistic a glass material, supporting refraction and volumetric attenuation (i.e., the transparency color varies with the geometric thickness). To create a Glass material pass the type string “glass” to `ospNewMaterial`. Its parameters are

For convenience, the rather counter-intuitive physical attenuation coefficients will be calculated from the user inputs in such a way, that the `attenuationColor` will be the result when white light traveled through a glass of thickness `attenuationDistance`.




Figure 3.12 – Rendering of golden Metal material with textured roughness.




Figure 3.13 – Rendering of a fictional Alloy material with textured color.




Figure 3.14 – Rendering of a Glass material with orange attenuation.

| Type | Name | Default | Description |
|-------|---------------------|---------|--|
| float | eta | 1.5 | index of refraction |
| vec3f | attenuationColor | white | resulting color due to attenuation (linear RGB) |
| float | attenuationDistance | 1 | distance affecting attenuation |

Table 3.38 – Parameters of the Glass material.

3.6.7 ThinGlass

The [path tracer](#) offers a thin glass material useful for objects with just a single surface, most prominently windows. It models a thin, transparent slab, i.e., it behaves as if a second, virtual surface is parallel to the real geometric surface. The implementation accounts for multiple internal reflections between the interfaces (including attenuation), but neglects parallax effects due to its (virtual) thickness. To create a such a thin glass material pass the type string “`thinGlass`” to `ospNewMaterial`. Its parameters are

| Type | Name | Default | Description |
|-------|---------------------|---------|--|
| float | eta | 1.5 | index of refraction |
| vec3f | attenuationColor | white | resulting color due to attenuation (linear RGB) |
| float | attenuationDistance | 1 | distance affecting attenuation |
| float | thickness | 1 | virtual thickness |

Table 3.39 – Parameters of the Thin-Glass material.

For convenience the attenuation is controlled the same way as with the [Glass](#) material. Additionally, the color due to attenuation can be modulated with a [texture map_attenuationColor](#) ([texture transformations](#) are supported as well). If present, the color component of [geometries](#) is also used for the attenuation color. The `thickness` parameter sets the (virtual) thickness and allows for easy exchange of parameters with the (real) [Glass](#) material; internally just the ratio between `attenuationDistance` and `thickness` is used to calculate the resulting attenuation and thus the material appearance.




Figure 3.15 – Rendering of a ThinGlass material with red attenuation.

3.6.8 MetallicPaint

The [path tracer](#) offers a metallic paint material, consisting of a base coat with optional flakes and a clear coat. To create a [MetallicPaint](#) material pass the type string “`metallicPaint`” to `ospNewMaterial`. Its parameters are listed in the table below.




Figure 3.16 – Example image of a coloured window made with textured attenuation of the ThinGlass material.

| Type | Name | Default | Description |
|-------|-------------|-----------|---------------------------------------|
| vec3f | baseColor | white 0.8 | color of base coat (linear RGB) |
| float | flakeAmount | 0.3 | amount of flakes, in [0–1] |
| vec3f | flakeColor | Aluminium | color of metallic flakes (linear RGB) |
| float | flakeSpread | 0.5 | spread of flakes, in [0–1] |
| float | eta | 1.5 | index of refraction of clear coat |

Table 3.40 – Parameters of the MetallicPaint material.

The color of the base coat `baseColor` can be textured by a [texture map](#) `baseColor`, which also supports [texture transformations](#). If present, the color component of [geometries](#) is also used for the color of the base coat. Parameter `flakeAmount` controls the proportion of flakes in the base coat, so when setting it to 1 the `baseColor` will not be visible. The shininess of the metallic component is governed by `flakeSpread`, which controls the variation of the orientation of the flakes, similar to the `roughness` parameter of [Metal](#). Note that the effect of the metallic flakes is currently only computed on average, thus individual flakes are not visible.




Figure 3.17 – Rendering of a MetallicPaint material.

3.6.9 Luminous

The [path tracer](#) supports the Luminous material which emits light uniformly in all directions and which can thus be used to turn any geometric object into a light source. It is created by passing the type string “luminous” to `ospNewMaterial`. The amount of constant radiance that is emitted is determined by combining the general parameters of lights: [color](#) and [intensity](#) (which essentially means that parameter `intensityQuantity` is not needed because it is always `OSP_INTENSITY_QUANTITY_RADIANC`E).

| Type | Name | Default | Description |
|-------|--------------|---------|---|
| vec3f | color | white | color of the emitted light (linear RGB) |
| float | intensity | 1 | intensity of the light (a factor) |
| float | transparency | 0 | material transparency |

Table 3.41 – Parameters accepted by the Luminous material.

The emission can be textured by passing a `map_color` texture handle, [texture transformations](#) are supported as well.




Figure 3.18 – Rendering of a yellow Luminous material.

3.7 Texture

OSPRay currently implements two texture types (`texture2d` and `volume`) and is open for extension to other types by applications. More types may be added in future releases.

To create a new texture use

```
OSPTexture ospNewTexture(const char *type);
```

3.7.1 Texture2D

The `texture2d` texture type implements an image-based texture, where its parameters are as follows

| Type | Name | Description |
|---------------|----------|---|
| uint | format | OSPTextureFormat for the texture |
| uint | filter | default OSP_TEXTURE_FILTER_LINEAR, alternatively OSP_TEXTURE_FILTER_NEAREST |
| OSPData | data | the actual texel 2D <code>data</code> |
| uint / vec2ui | wrapMode | OSPTextureWrapMode for the texture coordinates s and t; supported wrap modes are: OSP_TEXTURE_WRAP_REPEAT (default) OSP_TEXTURE_WRAP_MIRRORED_REPEAT OSP_TEXTURE_WRAP_CLAMP_TO_EDGE |

Table 3.42 – Parameters of `texture2d` texture type.

The supported texture formats for `texture2d` are:

The size of the texture is inferred from the size of the 2D array `data`, which also needs have a compatible type to `format`. The texel data in `data` starts with

| Name | Description |
|---------------------|--|
| OSP_TEXTURE_RGBA8 | 8 bit [0–255] linear components red, green, blue, alpha |
| OSP_TEXTURE_SRGB8 | 8 bit sRGB gamma encoded color components, and linear alpha |
| OSP_TEXTURE_RGBA32F | 32 bit float components red, green, blue, alpha |
| OSP_TEXTURE_RGBA16F | 16 bit float components red, green, blue, alpha |
| OSP_TEXTURE_RGB8 | 8 bit [0–255] linear components red, green, blue |
| OSP_TEXTURE_SRGB8 | 8 bit sRGB gamma encoded components red, green, blue |
| OSP_TEXTURE_RGB32F | 32 bit float components red, green, blue |
| OSP_TEXTURE_RGB16F | 16 bit float components red, green, blue |
| OSP_TEXTURE_R8 | 8 bit [0–255] linear single component red |
| OSP_TEXTURE_RA8 | 8 bit [0–255] linear two components red, alpha |
| OSP_TEXTURE_L8 | 8 bit [0–255] gamma encoded luminance (replicated into red, green, blue) |
| OSP_TEXTURE_LA8 | 8 bit [0–255] gamma encoded luminance, and linear alpha |
| OSP_TEXTURE_RA32F | 32 bit float two component red, alpha |
| OSP_TEXTURE_R32F | 32 bit float single component red |
| OSP_TEXTURE_RA16F | 16 bit float two component red, alpha |
| OSP_TEXTURE_R16F | 16 bit float single component red |
| OSP_TEXTURE_RGBA16 | 16 bit [0–65535] linear components red, green, blue, alpha |
| OSP_TEXTURE_RGB16 | 16 bit [0–65535] linear components red, green, blue |
| OSP_TEXTURE_RA16 | 16 bit [0–65535] linear two components red, alpha |
| OSP_TEXTURE_R16 | 16 bit [0–65535] linear single component red |

Table 3.43 – Supported texture formats by `texture2d`, i.e., valid constants of type `OSPTextureFormat`.

the texels in the lower left corner of the texture image, like in OpenGL. Per default a texture fetch is filtered by performing bi-linear interpolation of the nearest 2×2 texels; if instead fetching only the nearest texel is desired (i.e., no filtering) then pass the `OSP_TEXTURE_FILTER_NEAREST` flag.

Texturing with `texture2d` image textures requires `geometries` with texture coordinates, e.g., a `mesh` with `vertex.texcoord` provided.

3.7.2 Volume Texture

The `volume` texture type implements texture lookups based on 3D object coordinates of the surface hit point on the associated geometry. If the given hit point is within the attached volume, the volume is sampled and classified with the transfer function attached to the volume. This implements the ability to visualize volume values (as colored by a transfer function) on arbitrary surfaces inside the volume (as opposed to an isosurface showing a particular value in the volume). Its parameters are as follows

`TextureVolume` can be used for implementing slicing of volumes with any geometry type. It enables coloring of the slicing geometry with a different transfer

| Type | Name | Description |
|----------------------|------------------|---------------------------------------|
| OSPVolume | volume | Volume used to generate color lookups |
| OSPTTransferFunction | transferFunction | transfer function applied to volume |

Table 3.44 – Parameters of `volume` texture type.

function than that of the sliced volume.

3.7.3 Texture Transformations

All materials with textures also offer to manipulate the placement of these textures with the help of texture transformations. If so, this convention shall be used: the following parameters are prefixed with “`texture_name.*`”).

| Type | Name | Description |
|----------|-------------|--|
| linear2f | transform | linear transformation (rotation, scale) |
| float | rotation | angle in degree, counterclockwise, around center |
| vec2f | scale | enlarge texture, relative to center (0.5, 0.5) |
| vec2f | translation | move texture in positive direction (right/up) |

Table 3.45 – Parameters to define 2D texture coordinate transformations.

Above parameters are combined into a single `affine2d` transformation matrix and the transformations are applied in the given order. Rotation, scale and translation are interpreted “texture centric”, i.e., their effect seen by an user are relative to the texture (although the transformations are applied to the texture coordinates).

| Type | Name | Description |
|----------|-----------|--|
| affine3f | transform | linear transformation (rotation, scale) plus translation |

Table 3.46 – Parameter to define 3D volume texture transformations.

Similarly, volume texture placement can also be modified by an `affine3f` transformation matrix.

3.8 Cameras

To create a new camera of given type `type` use

```
OSPCamera ospNewCamera(const char *type);
```

All cameras accept these parameters:

The camera is placed and oriented in the world with `position`, `direction` and `up`. Additionally, an extra transformation `transform` can be specified, which will only be applied to 3D vectors (i.e., `position`, `direction` and `up`), but does *not* affect any sizes (e.g., `nearClip`, `apertureRadius`, or `height`). The same holds for the array of transformations `motion.transform` to achieve camera motion blur (in combination with `time` and `shutter`).

OSPRay uses a right-handed coordinate system. The region of the camera sensor that is rendered to the image can be specified in normalized screen-space coordinates with `imageStart` (lower left corner) and `imageEnd` (upper right corner). This can be used, for example, to crop the image, to achieve asymmetrical view frusta, or to horizontally flip the image to view scenes which are specified in a left-handed coordinate system. Note that values outside the default

Table 3.47 – Parameters accepted by all cameras.

| Type | Name | Default | Description |
|------------|------------------------|--------------------|--|
| vec3f | position | (0, 0, 0) | position of the camera |
| vec3f | direction | (0, 0, 1) | main viewing direction of the camera |
| vec3f | up | (0, 1, 0) | up direction of the camera |
| affine3f | transform | identity | additional world-space transform, overridden by <code>motion.*</code> arrays |
| float | nearClip | 10^{-6} | near clipping distance |
| vec2f | imageStart | (0, 0) | start of image region (lower left corner) |
| vec2f | imageEnd | (1, 1) | end of image region (upper right corner) |
| affine3f[] | motion.transform | | additional uniformly distributed world-space transforms |
| vec3f[] | motion.scale | | additional uniformly distributed world-space scale, overridden by <code>motion.transform</code> |
| vec3f[] | motion.pivot | | additional uniformly distributed world-space translation which is applied before <code>motion.rotation</code> (i.e., the rotation center), overridden by <code>motion.transform</code> |
| quatf[] | motion.rotation | | additional uniformly distributed world-space quaternion rotation, overridden by <code>motion.transform</code> |
| vec3f[] | motion.translation | | additional uniformly distributed world-space translation, overridden by <code>motion.transform</code> |
| box1f | time | [0, 1] | time associated with first and last key in <code>motion.*</code> arrays |
| box1f | shutter | [0.5, 0.5] | start and end of shutter time (for motion blur), in [0, 1] |
| uint | shutterType | OSP_SHUTTER_GLOBAL | OSPShutterType for motion blur, also allowed are: OSP_SHUTTER_ROLLING_RIGHT OSP_SHUTTER_ROLLING_LEFT OSP_SHUTTER_ROLLING_DOWN OSP_SHUTTER_ROLLING_UP |
| float | rollingShutterDuration | 0 | for a rolling shutter (see <code>shutterType</code>) the “open” time per line, in [0, <code>shutter.upper-shutter.lower</code>] |

range of [0–1] are valid, which is useful to easily realize overscan or film gate, or to emulate a shifted sensor.

3.8.1 Perspective Camera

The perspective camera implements a simple thin lens camera for perspective rendering, supporting optionally depth of field and stereo rendering (with the [path tracer](#)). It is created by passing the type string “perspective” to `ospNewCamera`. In addition to the [general parameters](#) understood by all cameras the perspective camera supports the special parameters listed in the table below.

Note that when computing the aspect ratio a potentially set image region (using `imageStart` & `imageEnd`) needs to be regarded as well.

Table 3.48 – Additional parameters accepted by the perspective camera.

| Type | Name | Default | Description |
|-------|------------------------|-----------------|---|
| float | fovy | 60 | the field of view (angle in degree) of the frame's height |
| float | aspect | 1 | ratio of width by height of the frame (and image region) |
| float | apertureRadius | 0 | size of the aperture, controls the depth of field |
| float | focusDistance | 1 | distance at where the image is sharpest when depth of field is enabled |
| bool | architectural | false | vertical edges are projected to be parallel |
| uint | stereoMode | OSP_STEREO_NONE | OSPSTEREOMode for stereo rendering, also allowed are: OSP_STEREO_LEFT OSP_STEREO_RIGHT OSP_STEREO_SIDE_BY_SIDE OSP_STEREO_TOP_BOTTOM (left eye at top half) |
| float | interpupillaryDistance | 0.0635 | distance between left and right eye when stereo is enabled |

In architectural photography it is often desired for aesthetic reasons to display the vertical edges of buildings or walls vertically in the image as well, regardless of how the camera is tilted. Enabling the `architectural` mode achieves this by internally leveling the camera parallel to the ground (based on the `up` direction) and then shifting the lens such that the objects in direction `dir` are centered in the image. If finer control of the lens shift is needed use `imageStart` & `imageEnd`. Because the camera is now effectively leveled its image plane and thus the plane of focus is oriented parallel to the front of buildings, the whole façade appears sharp, as can be seen in the example images below. The resolution of the `framebuffer` is not altered by `imageStart`/`imageEnd`.

3.8.2 Orthographic Camera

The orthographic camera implements a simple camera with orthographic projection, without support for depth. It is created by passing the type string “`orthographic`” to `ospNewCamera`. In addition to the [general parameters](#) understood by all cameras the orthographic camera supports the following special parameters:

| Type | Name | Description |
|-------|--------|---|
| float | height | size of the camera's image plane in y, in world coordinates |
| float | aspect | ratio of width by height of the frame |

Table 3.49 – Additional parameters accepted by the orthographic camera.

For convenience the size of the camera sensor, and thus the extent of the scene that is captured in the image, can be controlled with the `height` parameter. The same effect can be achieved with `imageStart` and `imageEnd`, and both methods can be combined. In any case, the `aspect` ratio needs to be set accordingly to get an undistorted image.

3.8.3 Panoramic Camera

The panoramic camera implements a simple camera with support for stereo rendering. It captures the complete surrounding with a latitude / longitude mapping and thus the rendered images should best have a ratio of 2:1. A panoramic camera is created by passing the type string “`panoramic`” to `ospNewCamera`. It is

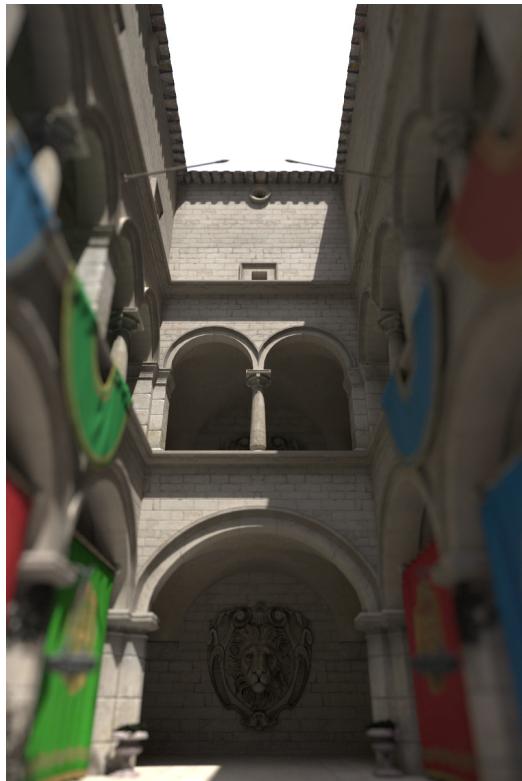


Figure 3.19 – Example image created with the perspective camera, featuring depth of field.

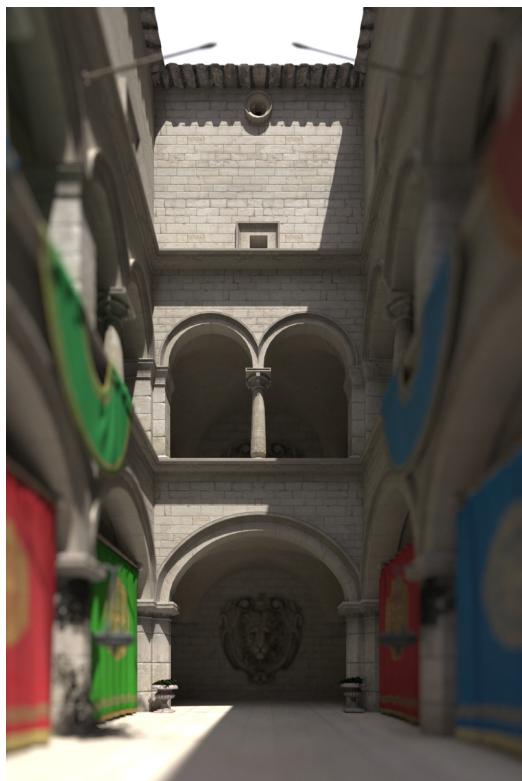


Figure 3.20 – Enabling the architectural flag corrects the perspective projection distortion, resulting in parallel vertical edges.

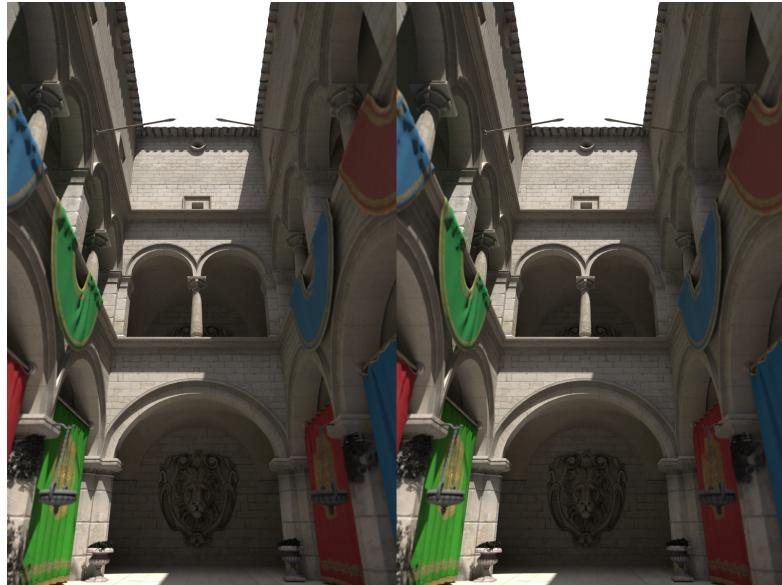


Figure 3.21 – Example 3D stereo image using `stereoMode = OSP_STEREO_SIDE_BY_SIDE`.

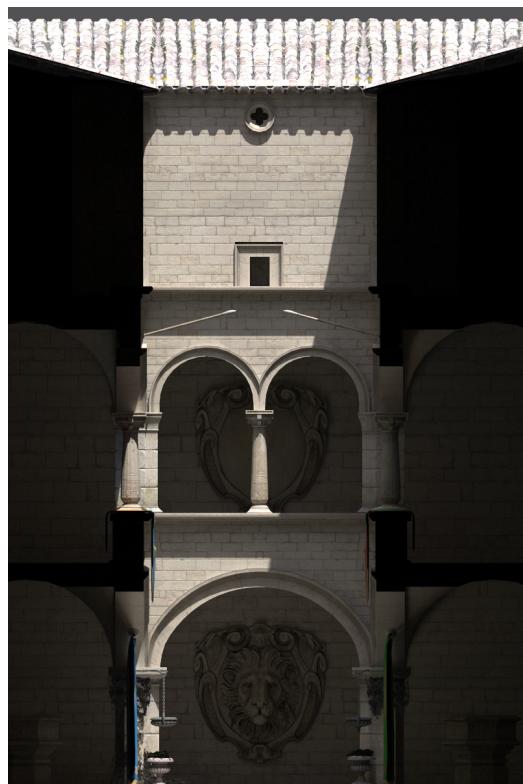


Figure 3.22 – Example image created with the orthographic camera.

placed and oriented in the scene by using the [general parameters](#) understood by all cameras.

Table 3.50 – Additional parameters accepted by the panoramic camera.

| Type | Name | Description |
|-------|------------------------|---|
| uint | stereoMode | OSPStereoMode for stereo rendering, possible values are: OSP_STEREO_NONE (default) OSP_STEREO_LEFT OSP_STEREO_RIGHT OSP_STEREO_SIDE_BY_SIDE OSP_STEREO_TOP_BOTTOM (left eye at top half) |
| float | interpupillaryDistance | distance between left and right eye when stereo is enabled, default 0.0635 |



Figure 3.23 – Latitude / longitude map created with the panoramic camera.

3.9 Scene Hierarchy

3.9.1 Groups

Groups in OSPRay represent collections of GeometricModels, VolumetricModels and Lights which share a common local-space coordinate system. To create a group call

```
OSPGroup ospNewGroup();
```

Groups take arrays of geometric models, volumetric models, clipping geometric models and lights, but they are all optional. In other words, there is no need to create empty arrays if there are no geometries, volumes or lights in the group.

By adding OSPGeometricModels to the `clippingGeometry` array a clipping geometry feature is enabled. Geometries assigned to this parameter will be used as clipping geometries. Any supported geometry can be used for clipping⁵, the only requirement is that it has to distinctly partition space into clipping and non-clipping one. The use of clipping geometry that is not closed or infinite could result in rendering artifacts. User can decide which part of space is clipped by changing shading normals orientation with the `invertNormals` flag of the [GeometricModel](#). All geometries and volumes assigned to `geometry` or `volume` will be clipped. All clipping geometries from all groups and [Instances](#) will be combined together – a union of these areas will be applied to all other objects in the [world](#).

⁵ including spheres, boxes, infinite planes, closed meshes, closed subdivisions and curves

Table 3.51 – Parameters understood by groups.

| Type | Name | Default | Description |
|----------------------|------------------|---------|--|
| OSPGeometricModel[] | geometry | NULL | data array of GeometricModels |
| OSPVolumetricModel[] | volume | NULL | data array of VolumetricModels |
| OSPGeometricModel[] | clippingGeometry | NULL | data array of GeometricModels used for clipping |
| OSPLight[] | light | NULL | data array of lights |
| bool | dynamicScene | false | tell Embree to use faster BVH build (slower ray traversal), otherwise optimized for faster ray traversal (slightly slower BVH build) |
| bool | compactMode | false | tell Embree to use a more compact BVH in memory by trading ray traversal performance |
| bool | robustMode | false | tell Embree to enable more robust ray intersection code paths (slightly slower) |

3.9.2 Instances

Instances in OSPRay represent a single group's placement into the world via a transform. To create and instance call

```
OSPIstance ospNewInstance(OSPGroup);
```

The passed group can be NULL as long as the group to be instanced is passed as a parameter. If both a group is specified on object creation and as a parameter, the parameter value is used. If the parameter value is later removed, the group object passed on object creation is again used.

Table 3.52 – Parameters understood by instances.

| Type | Name | Default | Description |
|------------|--------------------|----------|---|
| OSPGroup | group | | optional group object to be instanced |
| affine3f | transform | identity | world-space transform for all attached geometries and volumes, overridden by motion.* arrays |
| affine3f[] | motion.transform | | uniformly distributed world-space transforms |
| vec3f[] | motion.scale | | uniformly distributed world-space scale, overridden by motion.transform |
| vec3f[] | motion.pivot | | uniformly distributed world-space translation which is applied before motion.rotation (i.e., the rotation center), overridden by motion.transform |
| quatf[] | motion.rotation | | uniformly distributed world-space quaternion rotation, overridden by motion.transform |
| vec3f[] | motion.translation | | uniformly distributed world-space translation, overridden by motion.transform |
| box1f | time | [0, 1] | time associated with first and last key in motion.* arrays (for motion blur) |
| uint32 | id | -1u | optional user ID, for framebuffer channel OSP_FB_ID_INSTANCE |

3.9.3 World

Worlds are a container of scene data represented by [instances](#). To create an (empty) world call

```
OSPRWorld ospNewWorld();
```

Objects are placed in the world through an array of instances. Similar to [groups](#), the array of instances is optional: there is no need to create empty arrays if there are no instances (though there will be nothing to render).

Applications can query the world (axis-aligned) bounding box after the world has been committed. To get this information, call

```
OSPBounds ospGetBounds(OSPObject);
```

The result is returned in the provided [OSPBounds](#)⁶ struct:

```
typedef struct {
    float lower[3];
    float upper[3];
} OSPBounds;
```

⁶ [OSPBounds](#) has essentially the same layout as the [OSP_BOX3F](#) [OSPDataType](#).

This call can also take [OSPGroup](#) and [OSPIstance](#) as well: all other object types will return an empty bounding box.

Finally, Worlds can be configured with parameters for making various feature/performance trade-offs (similar to groups).

Table 3.53 – Parameters understood by worlds.

| Type | Name | Default | Description |
|--------------|--------------|---------|--|
| OSPIstance[] | instance | NULL | data array with handles of the instances |
| OSPLight[] | light | NULL | data array with handles of the lights |
| bool | dynamicScene | false | tell Embree to use faster BVH build (slower ray traversal), otherwise optimized for faster ray traversal (slightly slower BVH build) |
| bool | compactMode | false | tell Embree to use a more compact BVH in memory by trading ray traversal performance |
| bool | robustMode | false | tell Embree to enable more robust ray intersection code paths (slightly slower) |

3.10 Renderers

A renderer is the central object for rendering in OSPRay. Different renderers implement different features and support different materials. To create a new renderer of given type `type` use

```
OSPRenderer ospNewRenderer(const char *type);
```

General parameters of all renderers are

OSPRay's renderers support a feature called adaptive accumulation, which accelerates progressive [rendering](#) by stopping the rendering and refinement of image regions that have an estimated variance below the [varianceThreshold](#). This feature requires a [framebuffer](#) with an [OSP_FB_VARIANCE](#) channel.

Per default the background of the rendered image will be transparent black, i.e., the alpha channel holds the opacity of the rendered objects. This eases transparency-aware blending of the image with an arbitrary background image by the application (via the “over” operator, i.e., $ospray.rgb + appBackground.rgb \cdot (1 - ospray.alpha)$). The parameter [backgroundColor](#) or [map_backplate](#) can be used to already blend with a constant background color or backplate texture, respectively, (and alpha) during rendering.

Table 3.54 – Parameters understood by all renderers.

| Type | Name | Default | Description |
|-----------------------|-------------------|-----------------------|---|
| int | pixelSamples | 1 | samples per pixel, best results when a power of 2 |
| int | maxPathLength | 20 | maximum ray recursion depth |
| float | minContribution | 0.001 | sample contributions below this value will be neglected to speedup rendering |
| float | varianceThreshold | 0 | threshold for adaptive accumulation |
| float / vec3f / vec4f | backgroundColor | black, transparent | background color and alpha (linear A/RGB/RGBA), if no <code>map_backplate</code> is set |
| OSPTexture | map_backplate | | optional <code>texture</code> image used as background (use texture type <code>texture2d</code>) |
| OSPTexture | map_maxDepth | | optional screen-sized float <code>texture</code> with maximum far distance per pixel (use texture type <code>texture2d</code>) |
| OSPMaterial[] | material | | optional <code>data</code> array of <code>materials</code> which can be indexed by a <code>GeometricModel</code> 's <code>material</code> parameter |
| uint | pixelFilter | OSP_PIXELFILTER_GAUSS | OSPPixelFilterType to select the pixel filter used by the renderer for antialiasing. Possible pixel filters are listed below. |
| float | mipMapBias | 0 | bias for texture MIP-mapping, balancing between sharpness/aliasing and blurriness due to prefiltering |

OSPRay renderers support depth composition with images of other renderers, for example to incorporate help geometries of a 3D UI that were rendered with OpenGL. The screen-sized `texture` `map_maxDepth` must have format `OSP_TEXTURE_R32F` and flag `OSP_TEXTURE_FILTER_NEAREST`. The fetched values are used to limit the distance of primary rays, thus objects of other renderers can hide objects rendered by OSPRay.

OSPRay supports antialiasing in image space by using pixel filters, which are aligned around the center of a pixel. The size $w \times w$ of the filter depends on the selected filter type. The types of supported pixel filters are defined by the `OSPPixelFilterType` enum and can be set using the `pixelFilter` parameter.

Table 3.55 – Pixel filter types supported by OSPRay for antialiasing in image space.

| Name | Description |
|--|---|
| <code>OSP_PIXELFILTER_POINT</code> | a point filter only samples the center of the pixel, therefore the filter width is $w = 0$ |
| <code>OSP_PIXELFILTER_BOX</code> | a uniform box filter with a width of $w = 1$ |
| <code>OSP_PIXELFILTER_GAUSS</code> | a truncated, smooth Gaussian filter with a standard deviation of $\sigma = 0.5$ and a filter width of $w = 3$ |
| <code>OSP_PIXELFILTER_MITCHELL</code> | the Mitchell-Netravali filter with a width of $w = 4$ |
| <code>OSP_PIXELFILTER_BLACKMAN_HARRIS</code> | the Blackman-Harris filter with a width of $w = 3$ |

OSPRay also antialiases textures with prefiltering and MIP-mapping, which can be adjusted with parameter `mipMapBias`. For final frame rendering with a high number of `pixelSamples` or accumulated frames `mipMapBias` can be lowered (e.g., set to -0.5 or -2) to result in sharper textures which are essentially anisotropically filtered. Conversely, for preview rendering with just a single sample per pixel a higher `mipMapBias` of 1 or 2 can reduce texture aliasing and increase rendering speed.

3.10.1 SciVis Renderer

The SciVis renderer is a fast ray tracer for scientific visualization which supports volume rendering and ambient occlusion (AO). It is created by passing the type string “`scivis`” to `ospNewRenderer`. In addition to the [general parameters](#) understood by all renderers, the SciVis renderer supports the following parameters:

Table 3.56 – Special parameters understood by the SciVis renderer.

| Type | Name | Default | Description |
|-------|--------------------|-----------|--|
| bool | shadows | false | whether to compute (hard) shadows |
| int | aoSamples | 0 | number of rays per sample to compute ambient occlusion |
| float | aoDistance | 10^{20} | maximum distance to consider for ambient occlusion |
| float | volumeSamplingRate | 1 | sampling rate for volumes |
| bool | visibleLights | false | whether light sources are potentially visible (as in the path tracer , regarding each light’s <code>visible</code>) |

Note that the intensity (and color) of AO is deduced from an [ambient light](#) in the [lights](#) array.⁷ If `aoSamples` is zero (the default) then ambient lights cause ambient illumination (without occlusion).

⁷ If there are multiple ambient lights then their contribution is added.

3.10.2 Ambient Occlusion Renderer

This renderer supports only a subset of the features of the [SciVis renderer](#) to gain performance. As the name suggest its main shading method is ambient occlusion (AO), [lights](#) are *not* considered at all. Volume rendering is supported. The Ambient Occlusion renderer is created by passing the type string “`ao`” to `ospNewRenderer`. In addition to the [general parameters](#) understood by all renderers the following parameters are supported as well:

Table 3.57 – Special parameters understood by the Ambient Occlusion renderer.

| Type | Name | Default | Description |
|-------|--------------------|-----------|--|
| int | aoSamples | 1 | number of rays per sample to compute ambient occlusion |
| float | aoDistance | 10^{20} | maximum distance to consider for ambient occlusion |
| float | aoIntensity | 1 | ambient occlusion strength |
| float | volumeSamplingRate | 1 | sampling rate for volumes |

3.10.3 Path Tracer

The path tracer supports soft shadows, indirect illumination and realistic materials. This renderer is created by passing the type string “`pathtracer`” to

Table 3.58 – Special parameters understood by the path tracer.

| Type | Name | Default | Description |
|-------|---------------------------|----------|--|
| int | lightSamples | all | number of random light samples per path vertex, best results when a power of 2; per default all light sources are sampled |
| bool | limitIndirectLightSamples | true | after the first non-specular (i.e., diffuse and glossy) path vertex take (at most) a single light sample (instead of <code>lightSamples</code> many) |
| int | roulettePathLength | 5 | ray recursion depth at which to start Russian roulette termination |
| int | maxScatteringEvents | 20 | maximum number of non-specular (i.e., diffuse and glossy) bounces |
| float | maxContribution | ∞ | samples are clamped to this value before they are accumulated into the framebuffer |
| bool | backgroundRefraction | false | allow for alpha blending even if background is seen through refractive objects like glass |
| vec4f | shadowCatcherPlane | disabled | components of an invisible <code>plane</code> that captures shadow attenuations, which are blended with the background |

`ospNewRenderer`. In addition to the [general parameters](#) understood by all renderers the path tracer supports the following special parameters:

The path tracer requires that [materials](#) are assigned to [geometries](#), otherwise surfaces are treated as completely black.

The path tracer supports [volumes](#) with multiple scattering. The scattering albedo can be specified using the [transfer function](#). Extinction is assumed to be spectrally constant.

Rendering an object with an [HDRI light](#) results in realistic illumination, but the object looks detached and floats in space. As remedy a so-called Shadow Catcher can be used by setting the coefficients of a [plane](#) via `shadowCatcherPlane`, which will ground the object by computing matching (contact) shadows. The accuracy of the shadow estimation depends on the number of light samples per frame (via `lightSamples` and/or `pixelSamples`), in particular when using a high-contrast HDRI or multiple (colored) light sources. With just a single light sample the intensity, gradient and color of the shadows will potentially be quite off.

**Figure 3.24** – Chair rendered in an interior HDRI environment.

Applications can also perform their own compositing of the captured shadow, because the shadow is also baked into the alpha channel (which however means loosing the color tint of the shadow). To do so, the `backgroundColor` must be transparent (`alpha=0`) and should be black (otherwise the background color will bleed into the shadow); likewise, light sources should be set to `visible=false` to avoid them being visible in the background. When using the `denoiser` best enable `denoiseAlpha` as well.



Figure 3.25 – Same scene with enabled Shadow Catcher plane.



Figure 3.26 – The Shadow Catcher also bakes the (monochrome) shadow into the alpha channel, which can be used for further compositing.

3.11 Framebuffer

The framebuffer holds the rendered 2D image (and optionally auxiliary information associated with pixels). To create a new framebuffer object of given size `size` (in pixels), color format, and channels use

```
OSPFrameBuffer ospNewFrameBuffer(int size_x, int size_y,
    OSPFrameBufferFormat format = OSP_FB_SRGB,
    uint32_t frameBufferChannels = OSP_FB_COLOR);
```

The parameter `format` describes the format the color buffer has *on the host*, and the format that `ospMapFrameBuffer` will eventually return. Valid values are:

| Name | Description |
|----------------|---|
| OSP_FB_NONE | framebuffer will not be mapped by the application |
| OSP_FB_RGBA8 | 8 bit [0–255] linear component red, green, blue, alpha |
| OSP_FB_SRGB | 8 bit sRGB gamma encoded color components, and linear alpha |
| OSP_FB_RGBA32F | 32 bit float components red, green, blue, alpha |

Table 3.59 – Supported color formats of the framebuffer that can be passed to `ospNewFrameBuffer`, i.e., valid constants of type `OSPFrameBufferFormat`.

The parameter `frameBufferChannels` specifies which channels the framebuffer holds, and can be combined together by bitwise OR from the values of `OSPFrameBufferChannel` listed in the table below.

If a certain channel value is *not* specified, the given buffer channel will not be present. Note that OSPRay makes a clear distinction between the *external* format of the framebuffer and the internal one: The external format is the format the user specifies in the `format` parameter; it specifies what color format

| Name | Description |
|---------------------|---|
| OSP_FB_COLOR | RGB color including alpha |
| OSP_FB_DEPTH | distance to the camera, as linear 32 bit float; euclidean or projected distance is selected via parameter <code>projectedDepth</code> ; for multiple samples per pixel the closest is taken |
| OSP_FB_ACCUM | accumulation buffer for progressive refinement |
| OSP_FB_VARIANCE | for estimation of the current noise level if <code>OSP_FB_ACCUM</code> is also present, see rendering |
| OSP_FB_FIRST_NORMAL | accumulated world-space normal of the <i>first</i> hit, as <code>vec3f</code> |
| OSP_FB_NORMAL | accumulated world-space normal of the first <i>non-specular</i> hit (for denoising), as <code>vec3f</code> |
| OSP_FB_ALBEDO | accumulated material albedo (color without illumination) at the first non-specular hit (for denoising), as <code>vec3f</code> |
| OSP_FB_POSITION | accumulated world-space position of the first hit, as <code>vec3f</code> |
| OSP_FB_ID_PRIMITIVE | primitive index of the first hit, as <code>uint32</code> |
| OSP_FB_ID_OBJECT | geometric/volumetric model <code>id</code> , if specified, or index in group of first hit, as <code>uint32</code> |
| OSP_FB_ID_INSTANCE | user defined instance <code>id</code> , if specified, or instance index of first hit, as <code>uint32</code> |

Table 3.60 – Framebuffer channels constants (of type `OSPFrameBufferChannel`), naming optional information the framebuffer can store. These values can be combined by bitwise OR when passed to `ospNewFrameBuffer`.

OSPRay will eventually *return* the framebuffer to the application (when calling `ospMapFrameBuffer`): no matter what OSPRay uses internally, it will simply return a 2D array of pixels of that format, with possibly all kinds of reformatting, compression/decompression, etc., going on in-between the generation of the *internal* framebuffer and the mapping of the externally visible one.

In particular, `OSP_FB_NONE` is a perfectly valid pixel format for a framebuffer that an application will never map. For example, an application driving a display wall may well generate an intermediate framebuffer and eventually transfer its pixel to the individual displays using an `OSPIImageOperation` [image operation](#).

The application can map the given channel of a framebuffer – and thus access the stored pixel information – via

```
const void *ospMapFrameBuffer(OSPFrameBuffer, OSPFrameBufferChannel = OSP_FB_COLOR);
```

Note that `OSP_FB_ACCUM` or `OSP_FB_VARIANCE` cannot be mapped. The origin of the screen coordinate system in OSPRay is the lower left corner (as in OpenGL), thus the first pixel addressed by the returned pointer is the lower left pixel of the image.

A previously mapped channel of a framebuffer can be unmapped by passing the received pointer `mapped` to

```
void ospUnmapFrameBuffer(const void *mapped, OSPFrameBuffer);
```

The individual channels of a framebuffer can be cleared with

```
void ospResetAccumulation(OSPFrameBuffer);
```

This function will clear *all* accumulating buffers (`OSP_FB_VARIANCE`, `OSP_FB_NORMAL`, and `OSP_FB_ALBEDO`, if present) and resets the accumulation

counter `accumID`. It is unspecified if the existing color and depth buffers are physically cleared when `ospResetAccumulation` is called.

If `OSP_FB_VARIANCE` is specified, an estimate of the variance of the last accumulated frame can be queried with

```
float ospGetVariance(OSPFrameBuffer);
```

Note this value is only updated after synchronizing with `OSP_FRAME_FISHED`, as further described in [asynchronous rendering](#). The estimated variance can be used by the application as a quality indicator and thus to decide whether to stop or to continue progressive rendering.

The framebuffer takes a list of pixel operations to be applied to the image in sequence as an `OSPData`. The pixel operations will be run in the order they are in the array.

Table 3.61 – Parameters accepted by the framebuffer.

| Type | Name | Description |
|-----------------------------------|-----------------------------|---|
| <code>OSPIImageOperation[]</code> | <code>imageOperation</code> | ordered sequence of image operations |
| <code>int</code> | <code>targetFrames</code> | anticipated number of frames that will be accumulated for progressive refinement, used renderers to generate a blue noise sampling pattern; should be a power of 2, is always 1 without <code>OSP_FB_ACCUM</code> ; default 0 (disabled) |
| <code>bool</code> | <code>projectedDepth</code> | whether channel <code>OSP_FB_DEPTH</code> should hold the projected distance onto the main camera direction (which is the distance to the image plane for perspective camera and orthographic camera , can be negative for panoramic camera), default false (euclidean distance) |

If the total number of frames to be accumulated is known, then `targetFrames` should be set, because then renderers can generate more pleasing blue noise patterns. Accumulation stops when `targetFrames` is reached.

3.11.1 Image Operation

Image operations are functions that are applied to every pixel of a frame. Examples include post-processing, filtering, blending, tone mapping, or sending tiles to a display wall. To create a new pixel operation of given type `type` use

```
OSPIImageOperation ospNewImageOperation(const char *type);
```

3.11.1.1 Tone Mapper

The tone mapper is a pixel operation which implements a generic filmic tone mapping operator. Using the default parameters it approximates the Academy Color Encoding System (ACES). The tone mapper is created by passing the type string “tonemapper” to `ospNewImageOperation`. The tone mapping curve can be customized using the parameters listed in the table below.

3.11.1.2 Denoiser

OSPRay comes with a module that adds support for Intel® Open Image Denoise (OIDN). This is provided as an optional module as it creates an additional project dependency at compile time. The module implements a “denoiser” frame operation, which denoises the entire frame before the frame is completed. OIDN will automatically select the fastest device, using a GPU when available. The device selection be overridden by the environment variable `OIDN_DEFAULT_DEVICE`, possible values are `cpu`, `sycl`, `cuda`, `hip`, or a physical device ID

Table 3.62 – Parameters accepted by the tone mapper. The column “Filmic” lists alternative values for the popular “Uncharted 2” tone mapping curve (note that that curve includes an exposure bias to match 18% middle gray).

| Type | Name | Default | Filmic | Description |
|-------|-----------|---------|--------|--|
| float | exposure | 1.0 | 1.0 | amount of light per unit area |
| float | contrast | 1.6773 | 1.1759 | contrast (toe of the curve); typically is in [1–2] |
| float | shoulder | 0.9714 | 0.9746 | highlight compression (shoulder of the curve); typically is in [0.9–1] |
| float | midIn | 0.18 | 0.18 | mid-level anchor input; default is 18% gray |
| float | midOut | 0.18 | 0.18 | mid-level anchor output; default is 18% gray |
| float | hdrMax | 11.0785 | 6.3704 | maximum HDR input that is not clipped |
| bool | acesColor | true | false | apply the ACES color transforms |

Table 3.63 – Parameters accepted by the denoiser.

| Type | Name | Description |
|------|--------------|--|
| uint | quality | OSPDenoiserQuality for denoiser quality, default is OSP_DENOISER_QUALITY_MEDIUM: balanced quality/performance for interactive/real-time rendering; also allowed are: OSP_DENOISER_QUALITY_LOW: high performance, for interactive/real-time preview rendering OSP_DENOISER_QUALITY_HIGH: high quality, for final frame rendering |
| bool | denoiseAlpha | whether to denoise the alpha channel as well, default false |

3.12 Rendering

3.12.1 Asynchronous Rendering

Rendering is by default asynchronous (non-blocking), and is done by combining a framebuffer, renderer, camera, and world.

What to render and how to render it depends on the renderer’s parameters. If the framebuffer supports accumulation (i.e., it was created with OSP_FB_ACCUM) then successive calls to `ospRenderFrame` will progressively refine the rendered image (until `targetFrames` is reached).

To start an render task, use

```
OSPFuture ospRenderFrame(OSPFrameBuffer, OSPRenderer, OSPCamera, OSPWorld);
```

This returns an `OSPFuture` handle, which can be used to synchronize with the application, cancel, or query for progress of the running task. When `ospRenderFrame` is called, there is no guarantee when the associated task will begin execution.

Progress of a running frame can be queried with the following API function

```
float ospGetProgress(OSPFuture);
```

This returns the approximated progress of the task in [0-1].

Applications can cancel a currently running asynchronous operation via

```
void ospCancel(OSPFuture);
```

Applications can wait on the result of an asynchronous operation, or choose to only synchronize with a specific event. To synchronize with an `OSPFuture` use

```
void ospWait(OSPFuture, OSPSyncEvent = OSP_TASK_FINISHED);
```

The following are values which can be synchronized with the application

Table 3.64 – Supported events that can be passed to `ospWait`.

| Name | Description |
|---------------------|---|
| OSP_NONE_FINISHED | Do not wait for anything to be finished (immediately return from <code>ospWait</code>) |
| OSP_WORLD_COMMITTED | Wait for the world to be committed (not yet implemented) |
| OSP_WORLD_RENDERED | Wait for the world to be rendered, but not post-processing operations (Pixel/Tile/Frame Op) |
| OSP_FRAME_FINISHED | Wait for all rendering operations to complete |
| OSP_TASK_FINISHED | Wait on full completion of the task associated with the future. The underlying task may involve one or more of the above synchronization events |

Currently only rendering can be invoked asynchronously. However, future releases of OSPRay may add more asynchronous versions of API calls (and thus return `OSPFuture`).

Applications can query whether particular events are complete with

```
int ospIsReady(OSPFuture, OSPSyncEvent = OSP_TASK_FINISHED);
```

As the given running task runs (as tracked by the `OSPFuture`), applications can query a boolean [0, 1] result if the passed event has been completed.

Applications can query how long an async task ran with

```
float ospGetTaskDuration(OSPFuture);
```

This returns the wall clock execution time of the task in seconds. If the task is still running, this will block until the task is completed. This is useful for applications to query exactly how long an asynchronous task executed without the overhead of measuring both task execution + synchronization by the calling application.

3.12.2 Synchronous Rendering

For convenience in certain use cases, `ospray_util.h` provides a synchronous version of `ospRenderFrame`:

```
float ospRenderFrameBlocking(OSPFrameBuffer, OSPRenderer, OSPCamera, OSPWorld);
```

This version is the equivalent of:

```
ospRenderFrame
ospWait(f, OSP_TASK_FINISHED)
return ospGetVariance(fb)
```

This version is closest to `ospRenderFrame` from OSPRay v1.x.

3.12.3 Rendering and `ospCommit`

The use of either `ospRenderFrame` or `ospRenderFrameBlocking` requires that all objects in the scene being rendered have been committed before rendering occurs. If a call to `ospCommit` happens while a frame is rendered, the result is undefined behavior and should be avoided.

3.12.4 Picking

To get the world-space position of the geometry (if any) seen at [0–1] normalized screen-space pixel coordinates `screenPos_x` and `screenPos_y` use

```
void ospPick(OSPPickResult *,
    OSPFrameBuffer,
    OSPRenderer,
    OSPCamera,
    OSPWorld,
    float screenPos_x,
    float screenPos_y);
```

The result is returned in the provided `OSPPickResult` struct:

```
typedef struct {
    int hasHit;
    float worldPosition[3];
    OSPInstance instance;
    OSPGeometricModel model;
    uint32_t primID;
} OSPPickResult;
```

Note that `ospPick` considers exactly the same camera of the given renderer that is used to render an image, thus matching results can be expected. If the camera supports depth of field then the center of the lens and thus the center of the circle of confusion is used for picking. Note that the caller needs to `ospRelease` the `instance` and `model` handles of `OSPPickResult` once the information is not needed anymore.

Chapter 4

Modules and Devices

4.1 CPU

The CPU module is implicitly loaded and the `cpu` device is automatically used if no other options are specified.

4.2 GPU (Beta)

To use the GPU for rendering load the `gpu` module and select the `gpu` device:

```
./ospExamples --osp:load-modules=gpu --osp:device=gpu
```

or via explicit device creation by the application:

```
ospLoadModule("gpu");
OSPDevice dev = ospNewDevice("gpu");
ospDeviceCommit(dev);
ospSetCurrentDevice(dev);
```

| Type | Name | Description |
|--------|--------------------------|--------------|
| void * | <code>syclContext</code> | SYCL context |
| void * | <code>syclDevice</code> | SYCL device |

Table 4.1 – Parameters specific to the `gpu` device.

Applications can set their SYCL context and device to share device memory with OSPRay or to control which device should be used (e.g., in case multiple GPUs are present). If neither parameter is set, the `gpu` device will automatically create a context internally and select a GPU (that selection can be influenced via environment variable `ONEAPI_DEVICE_SELECTOR`, see [Intel oneAPI DPC++ Compiler documentation](#)).

Compile times for just in time compilation (JIT compilation) can be large. To resolve this issue we recommend enabling persistent JIT compilation caching inside your application before the SYCL device is created, by setting environment variables `SYCL_CACHE_PERSISTENT=1` (and optionally `SYCL_CACHE_DIR=<path>` to some proper directory where the JIT cache should get stored).

To reduce GPU memory allocation overhead when rendering scenes with many objects (geometries, instances, etc.), memory pooling should be enabled by setting the environment variable `SYCL_PI_LEVEL_ZERO_USM_ALLOCATION="1;0;shared:1M,0,2M"`. See [Intel oneAPI DPC++ Compiler documentation](#) for more details.

Known Issues

The following features are not implemented yet or are not working correctly on the GPU device:

- Multiple volumes in the scene
- Clipping
- Motion blur
- Subdivision surfaces
- Progress reporting via `ospGetProgress` or canceling the frame via `ospCancel`
- Picking via `ospPick`
- Adaptive accumulation via `OSP_FB_VARIANCE` and `varianceThreshold`
- Framebuffer channels `OSP_FB_ID_*` (id buffers)
- Experimental support for shared device-only data, works only for `structuredRegular` volume

There will be some delay on start-up as the kernel code is JIT compiled for the device, and similar pauses when changing the scene configuration, because the kernel specialized and re-compiled.

For some combination of compiler, GPU driver and scene the rendered images might show artifacts (e.g., vertical lines or small blocks).

4.3 Distributed Rendering with MPI

The purpose of OSPRay's MPI modules is to provide distributed rendering capabilities for OSPRay. The modules enables image- and data-parallel rendering across HPC clusters using MPI, allowing applications to transparently distribute rendering work, or to render data sets which are too large to fit in memory on a single machine.

OSPRay provides multiple MPI modules that expose different distributed rendering capabilities. The `mpi_offload` module provides image-parallel rendering through the `mpiOffload` device; it enables OSPRay applications written for local rendering to be replicated across multiple nodes to distribute the rendering work without code changes.

And the `mpi_distributed_cpu` and `mpi_distributed_gpu` modules provides data-parallel rendering through the `mpiDistributed` device, which allows MPI distributed applications to use OSPRay for distributed rendering. Each rank using the `mpiDistributed` device can render an independent piece of a global data set, or perform hybrid rendering where ranks partially or completely share data.

The `mpiDistributed` device's image-parallel rendering support can be used to accelerate data loading for image-parallel applications, where all ranks load the same data from a shared disk and then perform image-parallel rendering on the replicated data, as if the `mpiOffload` device were being used.

4.3.1 MPI Offload Rendering

The `mpiOffload` device can be used to distribute image rendering tasks across a cluster without requiring modifications to the application itself. Existing applications using OSPRay for local rendering simply be passed command line arguments to load the module and indicate that the `mpiOffload` device should be used for image-parallel rendering. To load the module, pass `--osp:load-modules=mpi_offload`, to select the `MPIOffloadDevice`, pass `--osp:device=mpiOffload`. For example, the `ospExamples` application can be run as:

```
mpirun -n <N> ./ospExamples --osp:load-modules=mpi_offload --osp:device=mpiOffload
```

and will automatically distribute the image rendering tasks among the corresponding N nodes. Note that in this configuration rank 0 will act as a master/application rank, and will run the user application code but not perform rendering locally. Thus, a minimum of 2 ranks are required, one master to run the application and one worker to perform the rendering. Running with 3 ranks for example would now distribute half the image rendering work to rank 1 and half to rank 2.

If more control is required over the placement of ranks to nodes, or you want to run a worker rank on the master node as well you can run the application and the `ospray_mpi_worker` program through MPI's MPMD mode. The `ospray_mpi_worker` will load the MPI module and select the offload device by default.

```
mpirun -n 1 ./ospExamples --osp:load-modules=mpi_offload --osp:device=mpiOffload \
: -n <N> ./ospray_mpi_worker
```

If initializing the `mpiOffload` device manually, or passing parameters through the command line, the following parameters can be set:

Table 4.2 – Parameters specific to the `mpiOffload` device.

| Type | Name | Default | Description |
|--------|-------------------------|---------|---|
| string | mpiMode | mpi | The mode to communicate with the worker ranks. <code>mpi</code> will assume you are launching the application and workers in the same <code>mpi</code> command (or split launch command). <code>mpi</code> is the only supported mode |
| uint | maxCommandBufferEntries | 8192 | Set the max number of commands to buffer before submitting the command buffer to the workers |
| uint | commandBufferSize | 512 MiB | Set the max command buffer size to allow. Units are in MiB. Max size is 1.8 GiB |
| uint | maxInlineContentSize | 32 MiB | Set the max size of an OSPData which can be inline'd into the command buffer instead of being sent separately. Max size is half the commandBufferSize. Units are in MiB |

The `maxCommandBufferEntries`, `commandBufferSize`, and `maxInlineContentSize` can also be set via the environment variables: `OSPRAY_MPI_MAX_COMMAND_BUFFER_ENTRIES`, `OSPRAY_MPI_COMMAND_BUFFER_SIZE`, and `OSPRAY_MPI_MAX_INLINE_DATA_SIZE`, respectively.

The `mpiOffload` device uses a dynamic load balancer by default. If you wish to use a static load balancer you can do so by setting the `OSPRAY_STATIC_BALANCER` environment variable to 1.

For the worker ranks to create GPU devices instead of the default CPU devices set the environment variable `OSPRAY_MPI_DISTRIBUTED_GPU`, e.g.,

```
export OSPRAY_MPI_DISTRIBUTED_GPU=1
```

or

```
mpiexec -genv OSPRAY_MPI_DISTRIBUTED_GPU 1 \
-n 1 ./ospExamples --osp:load-modules=mpi_offload --osp:device=mpiOffload \
: -n 2 ./ospray_mpi_worker
```

The `mpiOffload` device does not support multiple init/shutdown cycles. Thus, to run `ospBenchmark` for this device make sure to exclude the init/shutdown test by passing `--benchmark_filter=-ospInit_ospShutdown` through the command line.

4.3.2 MPI Distributed Rendering

While MPI Offload rendering is used to transparently distribute rendering work without requiring modification to the application, MPI Distributed rendering is targeted at use of OSPRay within MPI-parallel applications. The MPI distributed device can be selected by loading the `mpi_distributed_cpu` module for CPU rendering or `mpi_distributed_gpu` for GPU rendering, and manually creating and using an instance of the `mpiDistributed` device, for example:

```
ospLoadModule("mpi_distributed_cpu");

OSPDevice mpiDevice = ospNewDevice("mpiDistributed");
ospDeviceCommit(mpiDevice);
ospSetCurrentDevice(mpiDevice);
```

Your application can either initialize MPI before-hand, ensuring that `MPI_THREAD_SERIALIZED` or higher is supported, or allow the device to initialize MPI on commit. Thread multiple support is required if your application will make MPI calls while rendering asynchronously with OSPRay. When using the distributed device each rank can specify independent local data using the OSPRay API, as if rendering locally. However, when calling `ospRenderFrameAsync` the ranks will work collectively to render the data. The distributed device supports both image-parallel, where the data is replicated, and data-parallel, where the data is distributed, rendering modes. The `mpiDistributed` device will by default use each rank in `MPI_COMM_WORLD` as a render worker; however, it can also take a specific MPI communicator to use as the world communicator. Only those ranks in the specified communicator will participate in rendering.

Table 4.3 – Parameters specific to the `mpiDistributed` device.

| Type | Name | Default | Description |
|--------|-------------------|-----------------------------|---|
| void * | worldCommunicator | <code>MPI_COMM_WORLD</code> | The MPI communicator which the OSPRay workers should treat as their world |

Table 4.4 – Parameters specific to the distributed OSPWorld.

| Type | Name | Default | Description |
|----------------------|--------|---------|--|
| <code>box3f[]</code> | region | NULL | A list of bounding boxes which bound the owned local data to be rendered by the rank |

Table 4.5 – Parameters specific to the `mpiRaycast` renderer.

| Type | Name | Default | Description |
|-------|--------------------|-----------|--|
| int | aoSamples | 0 | The number of AO samples to take per-pixel |
| float | aoDistance | 10^{20} | The AO ray length to use. Note that if the AO ray would have crossed a rank boundary and ghost geometry is not available, there will be visible artifacts in the shading |
| float | volumeSamplingRate | 1 | sampling rate for volumes |

4.3.2.1 Image Parallel Rendering in the MPI Distributed Device

If all ranks specify exactly the same data, the distributed device can be used for image-parallel rendering. This works identical to the offload device, except that

the MPI-aware application is able to load data in parallel on each rank rather than loading on the master and shipping data out to the workers. When a parallel file system is available, this can improve data load times. Image-parallel rendering is selected by specifying the same data on each rank, and using any of the existing local renderers (e.g., `scivis`, `pathtracer`). See [ospMPIDistribTutorialRepli-cated](#) for an example.

4.3.2.2 Data Parallel Rendering in the MPI Distributed Device

The MPI Distributed device also supports data-parallel rendering with sort-last compositing. Each rank can specify a different piece of data, as long as the bounding boxes of each rank's data are non-overlapping. The rest of the scene setup is similar to local rendering; however, for distributed rendering only the `mpiRay-cast` renderer is supported. This renderer implements a subset of the `scivis` rendering features which are suitable for implementation in a distributed environment.

By default the aggregate bounding box of the instances in the local world will be used as the bounds of that rank's data. However, when using ghost zones for volume interpolation, geometry or ambient occlusion, each rank's data can overlap. To clip these non-owned overlap regions out a set of regions (the `region` parameter) can pass as a parameter to the `OSPWorld` being rendered. Each rank can specify one or more non-overlapping `box3f`'s which bound the portions of its local data which it is responsible for rendering. See the [ospMPIDistribTutori-alVolume](#) for an example.

Finally, the MPI distributed device also supports hybrid-parallel rendering, where multiple ranks can share a single piece of data. For each shared piece of data the rendering work will be assigned image-parallel among the ranks. Partially-shared regions are determined by finding those ranks specifying data with the same bounds (matching regions) and merging them. See the [ospMPIDistribTutorialPartialRepl](#) for an example.

Picking on Distributed Data in the MPI Distributed Device

Calling `ospPick` in the distributed device will find and return the closest global object at the screen position on the rank that owns that object. The other ranks will report no hit. Picking in the distributed device takes into account data clipping applied through the `regions` parameter to avoid picking ghost data.

4.3.3 Interaction with User Modules

The MPI Offload rendering mode trivially supports user modules, with the caveat that attempting to share data directly with the application (e.g., passing a `void *` or other tricks to the module) will not work in a distributed environment. Instead, use the `ospNewSharedData` API to share data from the application with OSPRay, which will in turn be copied over the network to the workers.

The MPI Distributed device also supports user modules, as all that is required for compositing the distributed data are the bounds of each rank's local data.

4.4 MultiDevice

The multidevice module is an experimental OSPRay device type that renders images by delegating off pixel tiles to a number of internal delegate OSPRay devices.

If you wish to try it set the `OSPRAY_NUM_SUBDEVICES` environmental variable to the number of subdevices you want to create and tell OSPRay to both load the `multidevice_cpu` extension and create a multidevice for rendering instead of the default CPU device.

One example in a bash like shell is as follows:

```
OSPRAY_NUM_SUBDEVICES=6 ./ospTutorial --osp:load-modules=multidevice_cpu --osp:device=multidevice
```

Note that the multidevice currently does not support OSPImageOperations
for denoising nor tone mapping.

Chapter 5

Tutorials

5.1 ospTutorial

A minimal working example demonstrating how to use OSPRay can be found at [apps/tutorials/ospTutorial.c](#)¹.

An example of building `ospTutorial.c` with CMake can be found in [apps/tutorials/ospTutorialFindospray/](#).

To build the tutorial on Linux, build it in a build directory with

```
gcc -std=c99 .../apps/ospTutorial/ospTutorial.c \
-I ..\ospray\include -L . -lospray -Wl,-rpath,. -o ospTutorial
```

On Windows build it can be build manually in a “`build_directory\$Configuration`” directory with

```
cl ..\..\apps\ospTutorial\ospTutorial.c -I ..\..\ospray\include -I ..\..\ospray.lib
```

Running `ospTutorial` will create two images of two triangles, rendered with the Scientific Visualization renderer with full Ambient Occlusion. The first image `firstFrame.ppm` shows the result after one call to `ospRenderFrame` – jagged edges and noise in the shadow can be seen. Calling `ospRenderFrame` multiple times enables progressive refinement, resulting in antialiased edges and converged shadows, shown after ten frames in the second image `accumulatedFrames.ppm`.




Figure 5.1 – First frame.

¹ A C++ version that uses the C++ convenience wrappers of OSPRay’s C99 API via `include/ospray/ospray.cpp.h` is available at [apps/tutorials/ospTutorial.cpp](#).




Figure 5.2 – After accumulating ten frames.

5.2 ospExamples

Apart from tutorials, OSPRay comes with a C++ app called `ospExamples` which is an elaborate easy-to-use tutorial, with a single interface to try various OSPRay features. It is aimed at providing users with multiple simple scenes composed of basic geometry types, lights, volumes etc. to get started with OSPRay quickly.

`ospExamples` app runs a `GLFWOSPRayWindow` instance that manages instances of the camera, framebuffer, renderer and other OSPRay objects necessary to render an interactive scene. The scene is rendered on a `GLFW` window with an `imgui` GUI controls panel for the user to manipulate the scene at runtime.

The application is located in `apps/ospExamples/` directory and can be built with CMake. It can be run from the build directory via:

```
./ospExamples <command-line-parameter>
```

The command line parameter is optional however.




Figure 5.3 – `ospExamples` application with default boxes scene.

5.2.1 Scenes

Different scenes can be selected from the `scenes` dropdown and each scene corresponds to an instance of a special `detail::Builder` struct. Example builders are located in `apps/common/ospray_testing/builders/`. These builders provide a usage guide for the OSPRay scene hierarchy and OSPRay API in the form of cpp wrappers. They instantiate and manage objects for the specific scene like `cpp::Geometry`, `cpp::Volume`, `cpp::Light` etc.

The `detail::Builder` base struct is mostly responsible for setting up OSPray world and objects common in all scenes (for example lighting and ground plane), which can be conveniently overridden in the derived builders.

5.2.2 Renderer

This app comes with four `renderer` options: `scivis`, `pathtracer`, `ao` and `debug`. The app provides some common rendering controls like `pixelSamples` and other more specific to the renderer type like `aoSamples` for the `scivis` and `ao` renderer or `maxPathLength` for the `pathtracer`.

The sun-sky lighting can be used in a sample scene by enabling the `renderSunSky` option of the `pathtracer` renderer. It allows the user to change `turbidity` and `sunDirection`.




Figure 5.4 – Rendering an evening sky with the `renderSunSky` option.

5.3 ospMPIDistribTutorial

A minimal working example demonstrating how to use OSPray for rendering distributed data can be found at [modules/mpi/tutorials/ospMPIDistribTutorial.c²](#).

The compilation process via CMake is the similar to [apps/tutorials/ospTutorialFindospray/](#), with the addition of finding and linking MPI.

To build the tutorial on Linux, build it in a build directory with

```
mpicc -std=c99 ../modules/mpi/tutorials/ospMPIDistribTutorial.c
-I ..../ospray/include -L . -lospray -Wl,-rpath,. -o ospMPIDistribTutorial
```

On Windows build it can be build manually in a “`build_directory\$Configuration`” directory with

```
cl ..\..\modules\mpi\tutorials\ospMPIDistribTutorial.c -I ..\..\ospray\include -I ..\..\ospray.lib
```

The MPI module does not need to be linked explicitly, as it is loaded as a module at runtime.

²A C++ version that uses the C++ convenience wrappers of OSPray’s C99 API via `include/ospray/ospray_cpp.h` is available at [modules/mpi/tutorials/ospMPIDistribTutorial.cpp](#).




Figure 5.5 – The first frame output by the `ospMPIDistribTutorial` or C++ tutorial with 4 ranks.




Figure 5.6 – The accumulated frame output by the `ospMPIDistribTutorial` or C++ tutorial with 4 ranks.

5.4 `ospMPIDistribTutorialSpheres` and `ospMPIDistribTutorialVolume`

The spheres and volume distributed tutorials are built as part of the MPI tutorials when building OSPRay with the MPI module and tutorials enabled. These tutorials demonstrate using OSPRay to render distributed data sets where each rank owns some subregion of the data, and displaying the output in an interactive window. The domain is distributed in a grid among the processes, each of which will generate a subset of the data corresponding to its subdomain.

In `ospMPIDistribTutorialSpheres`, each process generates a set of spheres within its assigned domain. The spheres are colored from dark to light blue, where lighter colors correspond to higher ranks.




Figure 5.7 – Running `ospMPIDistribTutorialSpheres` on 4 ranks.

In `ospMPIDistribTutorialVolume`, each process generates a subbrick of

volume data, which is colored by its rank.




Figure 5.8 – Running `ospMPIDistribTutorialVolume` on 4 ranks.

5.5 ospMPIDistribTutorialPartialRepl

The partially replicated MPI tutorial demonstrates how to use OSPRay’s distributed rendering capabilities to render data sets that are partially replicated among the processes. Each pair of ranks generates the same volume brick, allowing them to subdivide the rendering workload between themselves. For example, when run with two ranks, each will generate the same brick and be responsible for rendering half of the image tiles it projects to. When run with four ranks, the pairs of ranks 0,1 and 2,3 will generate the same data and divide the rendering workload for that data among themselves.

The image work subdivision happens automatically, based on which ranks specify the same bounding box for their data, as demonstrated in the tutorial.

The partially replicated distribution is useful to support load-balanced rendering of data sets that are too large to be fully replicated among the processes, but are small enough to be partially replicated among them.

5.6 ospMPIDistribTutorialReplicated

The replicated MPI tutorial demonstrates how OSPRay’s distributed rendering capabilities can be used to render data sets that are fully replicated among the ranks with advanced illumination effects. In this case, although the processes are run MPI parallel, each rank specifies the exact same data. OSPRay’s MPI parallel renderer will detect that the data is replicated in this case and use the same image-parallel rendering algorithms employed in the MPI offload rendering configuration to render the data. This image-parallel rendering algorithm supports all rendering configurations that are used in local rendering, e.g., path tracing, to provide high-quality images.

The replicated MPI tutorial supports the same scenes and parameters as the `ospExamples` app described above.

This mode can be useful when high-quality rendering is desired and it is possible to copy the entire data set on to each rank, or to accelerate loading of a large model by leveraging a parallel file system.

© 2013 Intel Corporation

Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries.

*Other names and brands may be claimed as the property of others.

Performance varies by use, configuration and other factors. Learn more at www.Intel.com/PerformanceIndex.

Intel optimizations, for Intel compilers or other products, may not optimize to the same degree for non-Intel products.