

INTEL® OSPRAY

AN OPEN, SCALABLE, PARALLEL, RAY TRACING BASED RENDERING ENGINE FOR HIGH-FIDELITY VISUALIZATION

Version 1.8.1
February 20, 2019

Contents

1	OSPRay Overview	4
1.1	OSPRay Support and Contact	4
1.2	Version History	4
2	Building OSPRay from Source	17
2.1	Prerequisites	17
2.2	Compiling OSPRay on Linux and Mac OS X	18
2.3	Compiling OSPRay on Windows	19
3	Finding an OSPRay install with CMake	20
4	OSPRay API	21
4.1	Initialization and Shutdown	21
4.1.1	Command Line Arguments	21
4.1.2	Manual Device Instantiation	21
4.1.3	Environment Variables	23
4.1.4	Error Handling and Status Messages	23
4.1.5	Loading OSPRay Extensions at Runtime	24
4.1.6	Shutting Down OSPRay	24
4.2	Objects	24
4.2.1	Parameters	25
4.2.2	Data	26
4.3	Volumes	26
4.3.1	Structured Volume	26
4.3.2	Adaptive Mesh Refinement (AMR) Volume	28
4.3.3	Unstructured Volumes	29
4.3.4	Transfer Function	29
4.4	Geometries	30
4.4.1	Triangle Mesh	30
4.4.2	Quad Mesh	30
4.4.3	Subdivision	31
4.4.4	Spheres	32
4.4.5	Cylinders	32
4.4.6	Streamlines	33
4.4.7	Curves	34
4.4.8	Isosurfaces	34
4.4.9	Slices	34
4.4.10	Instances	34
4.5	Renderer	35
4.5.1	SciVis Renderer	35
4.5.2	Path Tracer	36
4.5.3	Model	36
4.5.4	Lights	37
4.5.5	Materials	39

4.5.6	Texture	48
4.5.7	Texture2D Transformations	49
4.5.8	Cameras	49
4.5.9	Picking	52
4.6	Framebuffer	54
	Pixel Operation	55
4.7	Rendering	56
	Progress and Cancel	56
5	Parallel Rendering with MPI	57
5.1	Prerequisites for MPI Mode	57
5.2	Enabling the MPI Module in your Build	57
5.3	Modes of Using OSPRay’s MPI Features	57
5.4	Running an Application with the “offload” Device	58
5.5	Running an Application with the “distributed” Device	59
6	Scene graph	61
6.1	Hierarchy Structure	61
6.2	Traversals	62
6.3	Thread Safety	63
7	Examples	64
7.1	Simple Tutorial	64
7.2	Mini-App Tutorials	65
7.3	Example Viewer	65
7.4	Distributed Viewer	66
7.5	Demos	67

Chapter 1

OSPRay Overview

Intel® OSPRay is an open source, scalable, and portable ray tracing engine for high-performance, high-fidelity visualization on Intel Architecture CPUs. OSPRay is part of the [Intel Rendering Framework](#) and is released under the permissive [Apache 2.0 license](#).

The purpose of OSPRay is to provide an open, powerful, and easy-to-use rendering library that allows one to easily build applications that use ray tracing based rendering for interactive applications (including both surface- and volume-based visualizations). OSPRay is completely CPU-based, and runs on anything from laptops, to workstations, to compute nodes in HPC systems.

OSPRay internally builds on top of [Intel Embree](#) and [ISPC \(Intel SPMD Program Compiler\)](#), and fully exploits modern instruction sets like Intel SSE4, AVX, AVX2, and AVX-512 to achieve high rendering performance, thus a CPU with support for at least SSE4.1 is required to run OSPRay.

1.1 OSPRay Support and Contact

OSPRay is under active development, and though we do our best to guarantee stable release versions a certain number of bugs, as-yet-missing features, inconsistencies, or any other issues are still possible. Should you find any such issues please report them immediately via [OSPRay's GitHub Issue Tracker](#) (or, if you should happen to have a fix for it, you can also send us a pull request); for missing features please contact us via email at ospray@googlegroups.com.

For recent news, updates, and announcements, please see our complete [news/updates](#) page.

Join our [mailing list](#) to receive release announcements and major news regarding OSPRay.

1.2 Version History

1.2.1 Changes in v1.8.1:

- CMake bugfix to remove full paths to dependencies in packages

1.2.2 Changes in v1.8.0:

- This will be the last minor revision of OSPRay. Future development effort in the `devel` branch will be dedicated to v2.0 API changes and may break existing v1.x applications.
 - This will also be the last version of OSPRay to include `ospray_sg` and the Example Viewer. Users which depend on these should instead

- use the separate OSPRay Studio project, where `ospray_sg` will be migrated.
- We will continue to support patch releases of v1.8.x in case of any reported bugs
- Refactored CMake to use newer CMake concepts
 - All targets are now exported in OSPRay installs and can be consumed by client projects with associated includes, libraries, and definitions
 - OSPRay now requires CMake v3.1 to build
 - See documentation for more details
- Added new minimal tutorial apps to replace the more complex Example Viewer
- Added new “subdivision” geometry type to support subdivision surfaces
- Added support for texture formats L8, LA8 (gamma-encoded luminance), and RA8 (linear two component). Note that the enum `OSP_TEXTURE_FORMAT_INVALID` changed its value, thus recompilation may be necessary.
- Automatic epsilon handling, which removes the “`epsilon`” parameter on all renderers
- Normals in framebuffer channel `OSP_FB_NORMAL` are now in world-space
- Added support for Intel® Open Image Denoise to the Example Viewer
 - This same integration will soon be ported to OSPRay Studio
- Fixed artifacts for scaled instances of spheres, cylinders and streamlines
- Improvements to precision of intersections with cylinders and streamlines
- Fixed Quadlight: the emitting side is now indeed in direction $\text{edge1} \times \text{edge2}$

1.2.3 Changes in v1.7.3:

- Make sure a “default” device can always be created
- Fixed `ospNewTexture2D` (completely implementing old behavior)
- Cleanup any shared object handles from the OS created from `ospLoadModule()`

1.2.4 Changes in v1.7.2:

- Fixed issue in `mpi_offload` device where `ospRelease` would sometimes not correctly free objects
- Fixed issue in `ospray_sg` where structured volumes would not properly release the underlying OSPRay object handles

1.2.5 Changes in v1.7.1:

- Fixed issue where the Principled material would sometimes show up incorrectly as black
- Fixed issue where some headers were missing from install packages

1.2.6 Changes in v1.7.0:

- Generalized texture interface to support more than classic 2D image textures, thus `OSPTexture2D` and `ospNewTexture2D` are now deprecated, use the new API call `ospNewTexture("texture2d")` instead
 - Added new volume texture type to visualize volume data on arbitrary geometry placed inside the volume
- Added new framebuffer channels `OSP_FB_NORMAL` and `OSP_FB_ALBEDO`

- Applications can get information about the progress of rendering the current frame, and optionally cancel it, by registering a callback function via `ospSetProgressFunc()`
- Lights are not tied to the renderer type, so a new function `ospNewLight3()` was introduced to implement this. Please convert all uses of `ospNewLight()` and/or `ospNewLight2()` to `ospNewLight3()`
- Added sheenTint parameter to Principled material
- Added baseNormal parameter to Principled material
- Added low-discrepancy sampling to path tracer
- The `spp` parameter on the renderer no longer supports values less than 1, instead applications should render to a separate, lower resolution framebuffer during interaction to achieve the same behavior

1.2.7 Changes in v1.6.1:

- Many bug fixes
 - Quad mesh interpolation and sampling
 - Normal mapping in path tracer materials
 - Memory corruption with partly emitting mesh lights
 - Logic for setting thread affinity

1.2.8 Changes in v1.6.0:

- Updated ispc device to use Embree3 (minimum required Embree version is 3.1)
- Added new `ospShutdown()` API function to aid in correctness and determinism of OSPRay API cleanup
- Added "Principled" and "CarPaint" materials to the path tracer
- Improved flexibility of the tone mapper
- Improvements to unstructured volume
 - Support for wedges (in addition to tets and hexes)
 - Support for implicit isosurface geometry
 - Support for cell-centered data (as an alternative to per-vertex data)
 - Added an option to precompute normals, providing a memory/performance trade-off for applications
- Implemented "quads" geometry type to handle quads directly
- Implemented the ability to set "void" cell values in all volume types: when NaN is present as a volume's cell value the volume sample will be ignored by the SciVis renderer
- Fixed support for color strides which were not multiples of `sizeof(float)`
- Added support for RGBA8 color format to spheres, which can be set by specifying the "color_format" parameter as `OSP_UCHAR4`, or passing the "color" array through an `OSPData` of type `OSP_UCHAR4`.
- Added ability to configure Embree scene flags via `OSPModel` parameters
- `ospFreeFrameBuffer()` has been deprecated in favor of using `ospRelease()` to free framebuffer handles
- Fixed memory leak caused by incorrect parameter reference counts in ispc device
- Fixed occasional crashes in the `mpi_offload` device on shutdown
- Various improvements to sample apps and `ospray_sg`
 - Added new generator nodes, allowing the ability to inject programmatically generated scene data (only C++ for now)
 - Bug fixes and improvements to enhance stability and usability

1.2.9 Changes in v1.5.0:

- Unstructured tetrahedral volume now generalized to also support hexahedral data, now called `unstructured_volume`
- New function for creating materials (`ospNewMaterial2()`) which takes the renderer type string, not a renderer instance (the old version is now deprecated)
- New tonemapper `PixelOp` for tone mapping final frames
- Streamlines now support per-vertex radii and smooth interpolation
- `ospray_sg` headers are now installed alongside the SDK
- Core OSPRay `ispc` device now implemented as a module
 - Devices which implement the public API are no longer required to link the dependencies to core OSPRay (e.g., Embree v2.x)
 - By default, `ospInit()` will load the `ispc` module if a device was not created via `--osp:mpi` or `--osp:device:[name]`
- MPI devices can now accept an existing world communicator instead of always creating their own
- Added ability to control ISPC specific optimization flags via CMake options. See the various `ISPC_FLAGS_*` variables to control which flags get used
- Enhancements to sample applications
 - `ospray_sg` (and thus `ospExampleViewer/ospBenchmark`) can now be extended with new scene data importers via modules or the SDK
 - Updated `ospTutorial` examples to properly call `ospRelease()`
 - New options in the `ospExampleViewer` GUI to showcase new features (sRGB framebuffers, tone mapping, etc.)
- General bug fixes
 - Fixes to geometries with multiple emissive materials
 - Improvements to precision of ray-sphere intersections

1.2.10 Changes in v1.4.3:

- Several bug fixes
 - Fixed potential issue with static initialization order
 - Correct compiler flags for Debug config
 - Spheres `postIntersect` shading is now 64-bit safer

1.2.11 Changes in v1.4.2:

- Several cleanups and bug fixes
 - Fixed memory leak where the Embree BVH was never released when an `OSPModel` was released
 - Fixed a crash when API logging was enabled in certain situations
 - Fixed a crash in MPI mode when creating lights without a renderer
 - Fixed an issue with camera lens samples not initialized when `spp <= 0`
 - Fixed an issue in `ospExampleViewer` when specifying multiple data files
- The C99 tutorial is now indicated as the default; the C++ wrappers do not change the semantics of the API (memory management) so the C99 version should be considered first when learning the API

1.2.12 Changes in v1.4.1:

- Several cleanups and bug fixes
 - Improved precision of ray intersection with streamlines, spheres, and cylinder geometries
 - Fixed address overflow in framebuffer, in practice image size is now limited only by available memory
 - Fixed several deadlocks and race conditions
 - Fix shadow computation in SciVis renderer, objects behind light sources do not occlude anymore
 - No more image jittering with MPI rendering when no accumulation buffer is used
- Improved path tracer materials
 - Additionally support RGB eta/k for Metal
 - Added Alloy material, a “metal” with textured color
- Minimum required Embree version is now v2.15

1.2.13 Changes in v1.4.0:

- New adaptive mesh refinement (AMR) and unstructured tetrahedral volume types
- Dynamic load balancing is now implemented for the `mpi_offload` device
- Many improvements and fixes to the available path tracer materials
 - Specular lobe of OBJMaterial uses Blinn-Phong for more realistic highlights
 - Metal accepts spectral samples of complex refraction index
 - ThinGlass behaves consistent to Glass and can texture attenuation color
- Added Russian roulette termination to path tracer
- SciVis OBJMaterial accepts texture coordinate transformations
- Applications can now access depth information in MPI distributed uses of OSPRay (both `mpi_offload` and `mpi_distributed` devices)
- Many robustness fixes for both the `mpi_offload` and `mpi_distributed` devices through improvements to the `mpi_common` and `mpi_maml` infrastructure libraries
- Major sample app cleanups
 - `ospray_sg` library is the new basis for building apps, which is a scene graph implementation
 - Old (unused) libraries have been removed: miniSG, command line, importer, loaders, and scripting
 - Some removed functionality (such as scripting) may be reintroduced in the new infrastructure later, though most features have remained and have been improved
 - Optional improved texture loading has been transitioned from ImageMagick to OpenImageIO
- Many cleanups, bug fixes, and improvements to `ospray_common` and other support libraries
- This will be the last release in which we support MSVC12 (Visual Studio 2013). Future releases will require VS2015 or newer

1.2.14 Changes in v1.3.1:

- Improved robustness of OSPRay CMake `find_package` config

- Fixed bugs related to CMake configuration when using the OSPRay SDK from an install
- Fixed issue with Embree library when installing with `OSPRAY_INSTALL_DEPENDENCIES` enabled

1.2.15 Changes in v1.3.0:

- New MPI distributed device to support MPI distributed applications using OSPRay collectively for “in-situ” rendering (currently in “alpha”)
 - Enabled via new `mpi_distributed` device type
 - Currently only supports `raycast` renderer, other renderers will be supported in the future
 - All API calls are expected to be exactly replicated (object instances and parameters) except scene data (geometries and volumes)
 - The original MPI device is now called the `mpi_offload` device to differentiate between the two implementations
- Support of Intel® AVX-512 for next generation Intel® Xeon® processor (codename Skylake), thus new minimum ISPC version is 1.9.1
- Thread affinity of OSPRay’s tasking system can now be controlled via either device parameter `setAffinity`, or command line parameter `osp:setaffinity`, or environment variable `OSPRAY_SET_AFFINITY`
- Changed behavior of the background color in the SciVis renderer: `bgColor` now includes alpha and is always blended (no `backgroundEnabled` anymore). To disable the background do not set `bgColor`, or set it to transparent black (0, 0, 0, 0)
- Geometries “spheres” and “cylinders” now support texture coordinates
- The GLUT- and Qt-based demo viewer applications have been replaced by an example viewer with minimal dependencies
 - Building the sample applications now requires GCC 4.9 (previously 4.8) for features used in the C++ standard library; OSPRay itself can still be built with GCC 4.8
 - The new example viewer based on `ospray::sg` (called `ospExampleViewerSg`) is the single application we are consolidating to, `ospExampleViewer` will remain only as a deprecated viewer for compatibility with the old `ospGlutViewer` application
- Deprecated `ospCreateDevice()`; use `ospNewDevice()` instead
- Improved error handling
 - Various API functions now return an `OSPError` value
 - `ospDeviceSetStatusFunc()` replaces the deprecated `ospDeviceSetErrorMsgFunc()`
 - New API functions to query the last error (`ospDeviceGetLastErrorCode()` and `ospDeviceGetLastErrorMsg()`) or to register an error callback with `ospDeviceSetErrorFunc()`
 - Fixed bug where exceptions could leak to C applications

1.2.16 Changes in v1.2.1:

- Various bug fixes related to MPI distributed rendering, ISPC issues on Windows, and other build related issues

1.2.17 Changes in v1.2.0:

- Added support for volumes with `voxelType short` (16-bit signed integers). Applications need to recompile, because `OSPDataType` has been re-enumerated

- Removed SciVis renderer parameter `aoWeight`, the intensity (and now color as well) of AO is controlled via “ambient” lights. If `aoSamples` is zero (the default) then ambient lights cause ambient illumination (without occlusion)
- New SciVis renderer parameter `aoTransparencyEnabled`, controlling whether object transparency is respected when computing ambient occlusion (disabled by default, as it is considerable slower)
- Implement normal mapping for SciVis renderer
- Support of emissive (and illuminating) geometries in the path tracer via new material “Luminous”
- Lights can optionally made invisible by using the new parameter `isVisible` (only relevant for path tracer)
- OSPRay Devices are now extendable through modules and the SDK
 - Devices can be created and set current, creating an alternative method for initializing the API
 - New API functions for committing parameters on Devices
- Removed support for the first generation Intel® Xeon Phi™ coprocessor (codename Knights Corner)
- Other minor improvements, updates, and bug fixes
 - Updated Embree required version to v2.13.0 for added features and performance
 - New API function `ospDeviceSetErrorMsgFunc()` to specify a callback for handling message outputs from OSPRay
 - Added ability to remove user set parameter values with new `ospRemoveParam()` API function
 - The MPI device is now provided via a module, removing the need for having separately compiled versions of OSPRay with and without MPI
 - OSPRay build dependencies now only get installed if `OSPRAY_INSTALL_DEPENDENCIES` CMake variable is enabled

1.2.18 Changes in v1.1.2:

- Various bug fixes related to normalization, epsilons and debug messages

1.2.19 Changes in v1.1.1:

- Fixed support of first generation Intel Xeon Phi coprocessor (codename Knights Corner) and the COI device
- Fix normalization bug that caused rendering artifacts

1.2.20 Changes in v1.1.0:

- New “scivis” renderer features
 - Single sided lighting (enabled by default)
 - Many new volume rendering specific features
 - * Adaptive sampling to help improve the correctness of rendering high-frequency volume data
 - * Pre-integration of transfer function for higher fidelity images
 - * Ambient occlusion
 - * Volumes can cast shadows
 - * Smooth shading in volumes
 - * Single shading point option for accelerated shading
- Added preliminary support for adaptive accumulation in the MPI device

- Camera specific features
 - Initial support for stereo rendering with the perspective camera
 - Option `architectural` in perspective camera, rectifying vertical edges to appear parallel
 - Rendering a subsection of the full view with `imageStart`/`imageEnd` supported by all cameras
- This will be our last release supporting the first generation Intel Xeon Phi coprocessor (codename Knights Corner)
 - Future major and minor releases will be upgraded to the latest version of Embree, which no longer supports Knights Corner
 - Depending on user feedback, patch releases are still made to fix bugs
- Enhanced output statistics in `ospBenchmark` application
- Many fixes to the OSPRay SDK
 - Improved CMake detection of compile-time enabled features
 - Now distribute OSPRay configuration and ISPC CMake macros
 - Improved SDK support on Windows
- OSPRay library can now be compiled with `-Wall` and `-Wextra` enabled
 - Tested with GCC v5.3.1 and Clang v3.8
 - Sample applications and modules have not been fixed yet, thus applications which build OSPRay as a CMake subproject should disable them with `-DOSPRAY_ENABLE_APPS=OFF` and `-DOSPRAY_ENABLE_MODULES=OFF`
- Minor bug fixes, improvements, and cleanups
 - Regard shading normal when bump mapping
 - Fix internal CMake naming inconsistencies in macros
 - Fix missing API calls in C++ wrapper classes
 - Fix crashes on MIC
 - Fix thread count initialization bug with TBB
 - CMake optimizations for faster configuration times
 - Enhanced support for scripting in both `ospGlutViewer` and `ospBenchmark` applications

1.2.21 Changes in v1.0.0:

- New OSPRay SDK
 - OSPRay internal headers are now installed, enabling applications to extend OSPRay from a binary install
 - CMake macros for OSPRay and ISPC configuration now a part of binary releases
 - * CMake clients use them by calling `include(${OSPRAY_USE_FILE})` in their CMake code after calling `find_package(ospray)`
- New OSPRay C++ wrapper classes
 - * These act as a thin layer on top of OSPRay object handles, where multiple wrappers will share the same underlying handle when assigned, copied, or moved
 - * New OSPRay objects are only created when a class instance is explicitly constructed
 - * C++ users are encouraged to use these over the `ospray.h` API
- Complete rework of sample applications

- New shared code for parsing the `commandline`
- Save/load of transfer functions now handled through a separate library which does not depend on Qt
- Added `ospCvtParaViewTfcn` utility, which enables `ospVolumeViewer` to load color maps from ParaView
- GLUT based sample viewer updates
 - * Rename of `ospModelViewer` to `ospGlutViewer`
 - * GLUT viewer now supports volume rendering
 - * Command mode with preliminary scripting capabilities, enabled by pressing ‘:’ key (not available when using Intel C++ Compiler (icc))
- Enhanced support of sample applications on Windows
- New minimum ISPC version is 1.9.0
- Support of Intel® AVX-512 for second generation Intel Xeon Phi processor (codename Knights Landing) is now a part of the `OSPRAY_BUILD_ISA` CMake build configuration
 - Compiling AVX-512 requires `icc` to be enabled as a build option
- Enhanced error messages when `ospLoadModule()` fails
- Added `OSP_FB_RGB32F` support in the `DistributedFrameBuffer`
- Updated Glass shader in the path tracer
- Many miscellaneous cleanups, bug fixes, and improvements

1.2.22 Changes in v0.10.1:

- Fixed support of first generation Intel Xeon Phi coprocessor (codename Knights Corner)
- Restored missing implementation of `ospRemoveVolume()`

1.2.23 Changes in v0.10.0:

- Added new tasking options: `Cilk`, `Internal`, and `Debug`
 - Provides more ways for OSPRay to interact with calling application tasking systems
 - * `Cilk`: Use Intel® Cilk™ Plus language extensions (`icc` only)
 - * `Internal`: Use hand written OSPRay tasking system
 - * `Debug`: All tasks are run in serial (useful for debugging)
 - In most cases, Intel Threading Building Blocks (Intel TBB) remains the fastest option
- Added support for adaptive accumulation and stopping
 - `ospRenderFrame` now returns an estimation of the variance in the rendered image if the framebuffer was created with the `OSP_FB_VARIANCE` channel
 - If the renderer parameter `varianceThreshold` is set, progressive refinement concentrates on regions of the image with a variance higher than this threshold
- Added support for volumes with `voxelType ushort` (16-bit unsigned integers)
- `OSPTexture2D` now supports sRGB formats – actually most images are stored in sRGB. As a consequence the API call `ospNewTexture2D()` needed to change to accept the new `OSPTextureFormat` parameter
- Similarly, OSPRay’s framebuffer types now also distinguishes between linear and sRGB 8-bit formats. The new types are `OSP_FB_NONE`, `OSP_FB_RGB8`, `OSP_FB_SRGB8`, and `OSP_FB_RGBA32F`

- Changed “scivis” renderer parameter defaults
 - All shading (AO + shadows) must be explicitly enabled
- OSPRay can now use a newer, pre-installed Embree enabled by the new `OSPRAY_USE_EXTERNAL_EMBREE` CMake option
- New `ospcommon` library used to separately provide math types and OS abstractions for both OSPRay and sample applications
 - Removes extra dependencies on internal Embree math types and utility functions
 - `ospray.h` header is now C99 compatible
- Removed loaders module, functionality remains inside of `ospVolumeViewer`
- Many miscellaneous cleanups, bug fixes, and improvements
 - Fixed data distributed volume rendering bugs when using less blocks than workers
 - Fixes to CMake `find_package()` config
 - Fix bug in `GhostBlockBrickVolume` when using doubles
 - Various robustness changes made in CMake to make it easier to compile OSPRay

1.2.24 Changes in v0.9.1:

- Volume rendering now integrated into the “scivis” renderer
 - Volumes are rendered in the same way the “dvr” volume renderer renders them
 - Ambient occlusion works with implicit isosurfaces, with a known visual quality/performance trade-off
- Intel Xeon Phi coprocessor (codename Knights Corner) COI device and build infrastructure restored (volume rendering is known to still be broken)
- New support for CPack built OSPRay binary redistributable packages
- Added support for HDRI lighting in path tracer
- Added `ospRemoveVolume()` API call
- Added ability to render a subsection of the full view into the entire framebuffer in the perspective camera
- Many miscellaneous cleanups, bug fixes, and improvements
 - The depthbuffer is now correctly populated by in the “scivis” renderer
 - Updated default renderer to be “ao1” in `ospModelViewer`
 - Trianglemesh `postIntersect` shading is now 64-bit safe
 - Texture2D has been reworked, with many improvements and bug fixes
 - Fixed bug where MPI device would freeze while rendering frames with Intel TBB
 - Updates to CMake with better error messages when Intel TBB is missing

1.2.25 Changes in v0.9.0:

The OSPRay v0.9.0 release adds significant new features as well as API changes.

- Experimental support for data-distributed MPI-parallel volume rendering
- New SciVis-focused renderer (“raytracer” or “scivis”) combining functionality of “obj” and “ao” renderers
 - Ambient occlusion is quite flexible: dynamic number of samples, maximum ray distance, and weight

- Updated Embree version to v2.7.1 with native support for Intel AVX-512 for triangle mesh surface rendering on the Intel Xeon Phi processor (codename Knights Landing)
- OSPRay now uses C++11 features, requiring up to date compiler and standard library versions (GCC v4.8.0)
- Optimization of volume sampling resulting in volume rendering speedups of up to 1.5×
- Updates to path tracer
 - Reworked material system
 - Added texture transformations and colored transparency in OBJ material
 - Support for alpha and depth components of framebuffer
- Added thinlens camera, i.e., support for depth of field
- Tasking system has been updated to use Intel Threading Building Blocks (Intel TBB)
- The `ospGet*`() API calls have been deprecated and will be removed in a subsequent release

1.2.26 Changes in v0.8.3:

- Enhancements and optimizations to path tracer
 - Soft shadows (light sources: sphere, cone, extended spot, quad)
 - Transparent shadows
 - Normal mapping (OBJ material)
- Volume rendering enhancements
 - Expanded material support
 - Support for multiple lights
 - Support for double precision volumes
 - Added `ospSampleVolume()` API call to support limited probing of volume values
- New features to support compositing externally rendered content with OSPRay-rendered content
 - Renderers support early ray termination through a maximum depth parameter
 - New OpenGL utility module to convert between OSPRay and OpenGL depth values
- Added panoramic and orthographic camera types
- Proper CMake-based installation of OSPRay and CMake `find_package()` support for use in external projects
- Experimental Windows support
- Deprecated `ospNewTriangleMesh()`; use `ospNewGeometry("triangles")` instead
- Bug fixes and cleanups throughout the codebase

1.2.27 Changes in v0.8.2:

- Initial support for Intel AVX-512 and the Intel Xeon Phi processor (code-name Knights Landing)
- Performance improvements to the volume renderer
- Incorporated implicit slices and isosurfaces of volumes as core geometry types
- Added support for multiple disjoint volumes to the volume renderer
- Improved performance of `ospSetRegion()`, reducing volume load times

- Improved large data handling for the `shared_structured_volume` and `block_bricked_volume` volume types
- Added support for DDS horizon data to the seismic module
- Initial support in the Qt viewer for volume rendering
- Updated to ISPC 1.8.2
- Various bug fixes, cleanups and documentation updates throughout the codebase

1.2.28 Changes in v0.8.1:

- The volume renderer and volume viewer can now be run MPI parallel (data replicated) using the `--osp:mpi` command line option
- Improved performance of volume grid accelerator generation, reducing load times for large volumes
- The volume renderer and volume viewer now properly handle multiple isosurfaces
- Added small example tutorial demonstrating how to use OSPRay
- Several fixes to support older versions of GCC
- Bug fixes to `ospSetRegion()` implementation for arbitrarily shaped regions and setting large volumes in a single call
- Bug fix for geometries with invalid bounds; fixes streamline and sphere rendering in some scenes
- Fixed bug in depth buffer generation

1.2.29 Changes in v0.8.0:

- Incorporated early version of a new Qt-based viewer to eventually unify (and replace) the existing simpler GLUT-based viewers
- Added new path tracing renderer (`ospray/render/pathtracer`), roughly based on the Embree sample path tracer
- Added new features to the volume renderer
 - Gradient shading (lighting)
 - Implicit isosurfacing
 - Progressive refinement
 - Support for regular grids, specified with the `gridOrigin` and `gridSpacing` parameters
 - New `shared_structured_volume` volume type that allows voxel data to be provided by applications through a shared data buffer
 - New API call to set (sub-)regions of volume data (`ospSetRegion()`)
- Added a subsampling-mode, enabled with a negative `spp` parameter; the first frame after scene changes is rendered with reduced resolution, increasing interactivity
- Added multi-target ISA support: OSPRay will now select the appropriate ISA at runtime
- Added support for the Stanford SEP file format to the seismic module
- Added `--osp:numthreads <n>` command line option to restrict the number of threads OSPRay creates
- Various bug fixes, cleanups and documentation updates throughout the codebase

1.2.30 Changes in v0.7.2:

- Build fixes for older versions of GCC and Clang
- Fixed time series support in `ospVolumeViewer`
- Corrected memory management for shared data buffers
- Updated to ISPC 1.8.1

- Resolved issue in XML parser

Chapter 2

Building OSPRay from Source

The latest OSPRay sources are always available at the [OSPRay GitHub repository](#). The default master branch should always point to the latest tested bugfix release.

2.1 Prerequisites

OSPRay currently supports Linux, Mac OS X, and Windows. In addition, before you can build OSPRay you need the following prerequisites:

- You can clone the latest OSPRay sources via:

```
git clone https://github.com/ospray/ospray.git
```

- To build OSPRay you need [CMake](#), any form of C++11 compiler (we recommend using GCC, but also support Clang and the [Intel® C++ Compiler \(icc\)](#)), and standard Linux development tools. To build the example viewers, you should also have some version of OpenGL.
- Additionally you require a copy of the [Intel® SPMD Program Compiler \(ISPC\)](#), version 1.9.1 or later. Please obtain a release of ISPC from the [ISPC downloads page](#). The build system looks for ISPC in the PATH and in the directory right “next to” the checked-out OSPRay sources.¹ Alternatively set the CMake variable `ISPC_EXECUTABLE` to the location of the ISPC compiler.
- Per default OSPRay uses the [Intel® Threading Building Blocks](#) (TBB) as tasking system, which we recommend for performance and flexibility reasons. Alternatively you can set CMake variable `OSPRAY_TASKING_SYSTEM` to `OpenMP`, `Internal`, or `Cilk` (icc only).
- OSPRay also heavily uses Intel [Embree](#), installing version 3.2 or newer is required. If Embree is not found by CMake its location can be hinted with the variable `embree_DIR`. NOTE: Windows users should use Embree v3.2.2 or later.
- If available OSPRay’s [Example Viewer] can be compiled with support for Intel [Open Image Denoise](#) by enabling `OSPRAY_APPS_ENABLE_DENOISER`. You may need to hint the location of the library with the CMake variable `OpenImageDenoise_DIR`.

¹ For example, if OSPRay is in `~/Projects/ospray`, ISPC will also be searched in `~/Projects/ispc-v1.9.2-linux`

Depending on your Linux distribution you can install these dependencies using `yum` or `apt-get`. Some of these packages might already be installed or might have slightly different names.

Type the following to install the dependencies using `yum`:

```
sudo yum install cmake.x86_64  
sudo yum install tbb.x86_64 tbb-devel.x86_64
```

Type the following to install the dependencies using apt-get:

```
sudo apt-get install cmake-curses-gui  
sudo apt-get install libtbb-dev
```

Under Mac OS X these dependencies can be installed using [MacPorts](#):

```
sudo port install cmake tbb
```

Under Windows please directly use the appropriate installers for [CMake](#), [TBB](#), [ISPC](#) (for your Visual Studio version) and [Embree](#).

2.2 Compiling OSPRay on Linux and Mac OS X

Assume the above requisites are all fulfilled, building OSPRay through CMake is easy:

- Create a build directory, and go into it

```
mkdir ospray/build  
cd ospray/build
```

(We do recommend having separate build directories for different configurations such as release, debug, etc.).

- The compiler CMake will use will default to whatever the CC and CXX environment variables point to. Should you want to specify a different compiler, run cmake manually while specifying the desired compiler. The default compiler on most linux machines is gcc, but it can be pointed to clang instead by executing the following:

```
cmake -DCMAKE_CXX_COMPILER=clang++ -DCMAKE_C_COMPILER=clang ..
```

CMake will now use Clang instead of GCC. If you are ok with using the default compiler on your system, then simply skip this step. Note that the compiler variables cannot be changed after the first cmake or ccmake run.

- Open the CMake configuration dialog

```
ccmake ..
```

- Make sure to properly set build mode and enable the components you need, etc.; then type 'configure' and 'generate'. When back on the command prompt, build it using

```
make
```

- You should now have `libospray.so` as well as a set of example application. You can test your version of OSPRay using any of the examples on the [OSPRay Demos and Examples](#) page.

2.3 Compiling OSPRay on Windows

On Windows using the CMake GUI (`cmake-gui.exe`) is the most convenient way to configure OSPRay and to create the Visual Studio solution files:

- Browse to the OSPRay sources and specify a build directory (if it does not exist yet CMake will create it).
- Click “Configure” and select as generator the Visual Studio version you have (OSPRay needs Visual Studio 14 2015 or newer), for Win64 (32 bit builds are not supported by OSPRay), e.g., “Visual Studio 15 2017 Win64”.
- If the configuration fails because some dependencies could not be found then follow the instructions given in the error message, e.g., set the variable `embree_DIR` to the folder where Embree was installed.
- Optionally change the default build options, and then click “Generate” to create the solution and project files in the build directory.
- Open the generated `OSPRay.sln` in Visual Studio, select the build configuration and compile the project.

Alternatively, OSPRay can also be built without any GUI, entirely on the console. In the Visual Studio command prompt type:

```
cd path\to\ospray
mkdir build
cd build
cmake -G "Visual Studio 15 2017 Win64" [-D VARIABLE=value] ..
cmake --build . --config Release
```

Use `-D` to set variables for CMake, e.g., the path to Embree with `"-D embree_DIR=path\to\embree"`.

You can also build only some projects with the `--target` switch. Additional parameters after `--` will be passed to `msbuild`. For example, to build in parallel only the OSPRay library without the example applications use

```
cmake --build . --config Release --target ospray -- /m
```

Chapter 3

Finding an OSPRay install with CMake

Client applications using OSPRay can find it with CMake's `find_package()` command. For example,

```
find_package(ospray 1.8.0 REQUIRED)
```

finds OSPRay via OSPRay's configuration file `osprayConfig.cmake`¹. Once found, the following is all that is required to use OSPRay:

```
target_link_libraries(${client_target} ospray::ospray)
```

This will automatically propagate all required include paths, linked libraries, and compiler definitions to the client CMake target (either an executable or library).

Advanced users may want to link to additional targets which are exported in OSPRay's CMake config, which includes all installed modules. All targets built with OSPRay are exported in the `ospray::` namespace, therefore all targets locally used in the OSPRay source tree can be accessed from an install. For example, `ospray_common` can be consumed directly via the `ospray::ospray_common` target. All targets have their libraries, includes, and definitions attached to them for public consumption (please [report bugs](#) if this is broken!).

¹ This file is usually in `${install_location}/[lib|lib64]/cmake/ospray-${version}/`. If CMake does not find it automatically, then specify its location in variable `ospray_DIR` (either an environment variable or CMake variable).

Chapter 4

OSPRay API

To access the OSPRay API you first need to include the OSPRay header

```
#include "ospray/ospray.h"
```

where the API is compatible with C99 and C++.

4.1 Initialization and Shutdown

To use the API, OSPRay must be initialized with a “device”. A device is the object which implements the API. Creating and initializing a device can be done in either of two ways: command line arguments or manually instantiating a device.

4.1.1 Command Line Arguments

The first is to do so by giving OSPRay the command line from `main()` by calling

```
OSPError ospInit(int *argc, const char **argv);
```

OSPRay parses (and removes) its known command line parameters from your application’s `main` function. For an example see the [tutorial](#). For possible error codes see section [Error Handling and Status Messages](#). It is important to note that the arguments passed to `ospInit()` are processed in order they are listed. The following parameters (which are prefixed by convention with “`--osp:`”) are understood:

4.1.2 Manual Device Instantiation

The second method of initialization is to explicitly create the device yourself, and possibly set parameters. This method looks almost identical to how other [objects](#) are created and used by OSPRay (described in later sections). The first step is to create the device with

```
OSPDevice ospNewDevice(const char *type);
```

where the `type` string maps to a specific device implementation. OSPRay always provides the “default” device, which maps to a local CPU rendering device. If it is enabled in the build, you can also use “`mpi`” to access the MPI multi-node rendering device (see [Parallel Rendering with MPI](#) section for more information). Once a device is created, you can call

```
void ospDeviceSet1i(OSPDevice, const char *id, int val);
void ospDeviceSetString(OSPDevice, const char *id, const char *val);
void ospDeviceSetVoidPtr(OSPDevice, const char *id, void *val);
```

Table 4.1 – Command line parameters accepted by OSPRay’s `ospInit`.

Parameter	Description
<code>--osp:debug</code>	enables various extra checks and debug output, and disables multi-threading
<code>--osp:numthreads <n></code>	use n threads instead of per default using all detected hardware threads
<code>--osp:loglevel <n></code>	set logging level, default 0; increasing n means increasingly verbose log messages
<code>--osp:verbose</code>	shortcut for <code>--osp:loglevel 1</code>
<code>--osp:vv</code>	shortcut for <code>--osp:loglevel 2</code>
<code>--osp:module:<name></code>	load a module during initialization; equivalent to calling <code>ospLoadModule(name)</code>
<code>--osp:mpi</code>	enables MPI mode for parallel rendering with the <code>mpi_offload</code> device, to be used in conjunction with <code>mpirun</code> ; this will automatically load the “ <code>mpi</code> ” module if it is not yet loaded or linked
<code>--osp:mpi-offload</code>	same as <code>--osp:mpi</code>
<code>--osp:mpi-distributed</code>	same as <code>--osp:mpi</code> , but will create an <code>mpi_distributed</code> device instead; Note that this will likely require application changes to work properly
<code>--osp:logoutput <dst></code>	convenience for setting where status messages go; valid values for dst are <code>cerr</code> and <code>cout</code>
<code>--osp:erroroutput <dst></code>	convenience for setting where error messages go; valid values for dst are <code>cerr</code> and <code>cout</code>
<code>--osp:device:<name></code>	use name as the type of device for OSPRay to create; e.g., <code>--osp:device:default</code> gives you the default local device; Note if the device to be used is defined in a module, remember to pass <code>--osp:module:<name></code> first
<code>--osp:setaffinity <n></code>	if 1, bind software threads to hardware threads; 0 disables binding; default is 1 on KNL and 0 otherwise

Table 4.2 – Parameters shared by all devices.

Type	Name	Description
int	<code>numThreads</code>	number of threads which OSPRay should use
int	<code>logLevel</code>	logging level
string	<code>logOutput</code>	convenience for setting where status messages go; valid values are <code>cerr</code> and <code>cout</code>
string	<code>errorOutput</code>	convenience for setting where error messages go; valid values are <code>cerr</code> and <code>cout</code>
int	<code>debug</code>	set debug mode; equivalent to <code>logLevel=2</code> and <code>numThreads=1</code>
int	<code>setAffinity</code>	bind software threads to hardware threads if set to 1; 0 disables binding omitting the parameter will let OSPRay choose

to set parameters on the device. The following parameters can be set on all devices:

Once parameters are set on the created device, the device must be committed with

```
void ospDeviceCommit(OSPDevice);
```

To use the newly committed device, you must call

```
void ospSetCurrentDevice(OSPDevice);
```

This then sets the given device as the object which will respond to all other OSPRay API calls.

Users can change parameters on the device after initialization (from either method above), by calling

```
OSPDevice ospGetCurrentDevice();
```

This function returns the handle to the device currently used to respond to OSPRay API calls, where users can set/change parameters and recommit the device. If changes are made to the device that is already set as the current device, it does not need to be set as current again.

4.1.3 Environment Variables

Finally, OSPRay's generic device parameters can be overridden via environment variables for easy changes to OSPRay's behavior without needing to change the application (variables are prefixed by convention with "OSPRAY_"):

Variable	Description
OSPRAY_THREADS	equivalent to --osp:numthreads
OSPRAY_LOG_LEVEL	equivalent to --osp:loglevel
OSPRAY_LOG_OUTPUT	equivalent to --osp:logoutput
OSPRAY_ERROR_OUTPUT	equivalent to --osp:erroroutput
OSPRAY_DEBUG	equivalent to --osp:debug
OSPRAY_SET_AFFINITY	equivalent to --osp:setaffinity

Table 4.3 – Environment variables interpreted by OSPRay.

4.1.4 Error Handling and Status Messages

The following errors are currently used by OSPRay:

Name	Description
OSP_NO_ERROR	no error occurred
OSP_UNKNOWN_ERROR	an unknown error occurred
OSP_INVALID_ARGUMENT	an invalid argument was specified
OSP_INVALID_OPERATION	the operation is not allowed for the specified object
OSP_OUT_OF_MEMORY	there is not enough memory to execute the command
OSP_UNSUPPORTED_CPU	the CPU is not supported (minimum ISA is SSE4.1)

Table 4.4 – Possible error codes, i.e., valid named constants of type OSPError.

These error codes are either directly return by some API functions, or are recorded to be later queried by the application via

```
OSPError ospDeviceGetLastErrorCode(OSPDevice);
```

A more descriptive error message can be queried by calling

```
const char* ospDeviceGetLastErrorMsg(OSPDevice);
```

Alternatively, the application can also register a callback function of type

```
typedef void (*OSPErrorFunc)(OSPError, const char* errorDetails);
```

via

```
void ospDeviceSetErrorFunc(OSPDevice, OSPErrorFunc);
```

to get notified when errors occur.

Applications may be interested in messages which OSPRay emits, whether for debugging or logging events. Applications can call

```
void ospDeviceSetStatusFunc(OSPDevice, OSPStatusFunc);
```

in order to register a callback function of type

```
typedef void (*OSPStatusFunc)(const char* messageText);
```

which OSPRay will use to emit status messages. By default, OSPRay uses a callback which does nothing, so any output desired by an application will require that a callback is provided. Note that callbacks for C++ `std::cout` and `std::cerr` can be alternatively set through `ospInit()` or the `OSPRAY_LOG_OUTPUT` environment variable.

4.1.5 Loading OSPRay Extensions at Runtime

OSPRay's functionality can be extended via plugins, which are implemented in shared libraries. To load plugin name from `libospray_module_<name>.so` (on Linux and Mac OS X) or `ospray_module_<name>.dll` (on Windows) use

```
OSPError ospLoadModule(const char *name);
```

Modules are searched in OS-dependent paths. `ospLoadModule` returns `OSP_NO_ERROR` if the plugin could be successfully loaded.

4.1.6 Shutting Down OSPRay

When the application is finished using OSPRay (typically on application exit), the OSPRay API should be finalized with

```
void ospShutdown();
```

This API call ensures that the current device is cleaned up appropriately. Due to static object allocation having non-deterministic ordering, it is recommended that applications call `ospShutdown()` before the calling application process terminates.

4.2 Objects

All entities of OSPRay (the renderer, volumes, geometries, lights, cameras, ...) are a specialization of `OSPObject` and share common mechanism to deal with parameters and lifetime.

An important aspect of object parameters is that parameters do not get passed to objects immediately. Instead, parameters are not visible at all to objects until they get explicitly committed to a given object via a call to

```
void ospCommit(OSPObject);
```

at which time all previously additions or changes to parameters are visible at the same time. If a user wants to change the state of an existing object (e.g., to change the origin of an already existing camera) it is perfectly valid to do so, as long as the changed parameters are recommitted.

The commit semantic allow for batching up multiple small changes, and specifies exactly when changes to objects will occur. This is important to ensure performance and consistency for devices crossing a PCI bus, or across a network.

In our MPI implementation, for example, we can easily guarantee consistency among different nodes by MPI barrier'ing on every commit.

Note that OSPRay uses reference counting to manage the lifetime of all objects, so one cannot explicitly “delete” any object. Instead, to indicate that the application does not need and does not access the given object anymore, call

```
void ospRelease(OSPObject);
```

This decreases its reference count and if the count reaches 0 the object will automatically get deleted. Passing NULL is not an error.

4.2.1 Parameters

Parameters allow to configure the behavior of and to pass data to objects. However, objects do *not* have an explicit interface for reasons of high flexibility and a more stable compile-time API. Instead, parameters are passed separately to objects in an arbitrary order, and unknown parameters will simply be ignored. The following functions allow adding various types of parameters with name `id` to a given object:

```
// add a C-string (zero-terminated char *) parameter
void ospSetString(OSPObject, const char *id, const char *s);

// add an object handle parameter to another object
void ospSetObject(OSPObject, const char *id, OSPObject object);

// add an untyped pointer -- this will *ONLY* work in local rendering!
void ospSetVoidPtr(OSPObject, const char *id, void *v);

// add scalar and vector integer and float parameters
void ospSetf (OSPObject, const char *id, float x);
void ospSet1f (OSPObject, const char *id, float x);
void ospSet1i (OSPObject, const char *id, int32_t x);
void ospSet2f (OSPObject, const char *id, float x, float y);
void ospSet2fv(OSPObject, const char *id, const float *xy);
void ospSet2i (OSPObject, const char *id, int x, int y);
void ospSet2iv(OSPObject, const char *id, const int *xy);
void ospSet3f (OSPObject, const char *id, float x, float y, float z);
void ospSet3fv(OSPObject, const char *id, const float *xyz);
void ospSet3i (OSPObject, const char *id, int x, int y, int z);
void ospSet3iv(OSPObject, const char *id, const int *xyz);
void ospSet4f (OSPObject, const char *id, float x, float y, float z, float w);
void ospSet4fv(OSPObject, const char *id, const float *xyzw);

// additional functions to pass vector integer and float parameters in C++
void ospSetVec2f(OSPObject, const char *id, const vec2f &v);
void ospSetVec2i(OSPObject, const char *id, const vec2i &v);
void ospSetVec3f(OSPObject, const char *id, const vec3f &v);
void ospSetVec3i(OSPObject, const char *id, const vec3i &v);
void ospSetVec4f(OSPObject, const char *id, const vec4f &v);
```

Users can also remove parameters that have been explicitly set via an `ospSet` call. Any parameters which have been removed will go back to their default value during the next commit unless a new parameter was set after the parameter was removed. The following API function removes the named parameter from the given object:

```
void ospRemoveParam(OSPObject, const char *id);
```

4.2.2 Data

There is also the possibility to aggregate many values of the same type into an array, which then itself can be used as a parameter to objects. To create such a new data buffer, holding `numItems` elements of the given type, from the initialization data pointed to by `source` and optional creation flags, use

```
OSPData ospNewData(size_t numItems,
                    OSPDataType,
                    const void *source,
                    const uint32_t dataCreationFlags = 0);
```

The call returns an `OSPData` handle to the created array. The flag `OSP_DATA_SHARED_BUFFER` indicates that the buffer can be shared with the application. In this case the calling program guarantees that the `source` pointer will remain valid for the duration that this data array is being used. The enum type `OSPDataType` describes the different data types that can be represented in OSPRay; valid constants are listed in the table below.

To add a data array as parameter named `id` to another object call

```
void ospSetData(OSPObject, const char *id, OSPData);
```

4.3 Volumes

Volumes are volumetric data sets with discretely sampled values in 3D space, typically a 3D scalar field. To create a new volume object of given type `type` use

```
OSPVolume ospNewVolume(const char *type);
```

The call returns `NULL` if that type of volume is not known by OSPRay, or else an `OSPVolume` handle.

The common parameters understood by all volume variants are summarized in the table below.

Note that if `voxelRange` is not provided for a volume then OSPRay will compute it based on the voxel data, which may result in slower data updates.

4.3.1 Structured Volume

Structured volumes only need to store the values of the samples, because their addresses in memory can be easily computed from a 3D position. A common type of structured volumes are regular grids. OSPRay supports two variants that differ in how the volumetric data for the regular grids is specified.

The first variant shares the voxel data with the application. Such a volume type is created by passing the type string “`shared_structured_volume`” to `ospNewVolume`. The voxel data is laid out in memory in xyz-order¹ and provided to the volume via a `data` buffer parameter named “`voxelData`”.

The second regular grid variant is optimized for rendering performance: data locality in memory is increased by arranging the voxel data in smaller blocks. This volume type is created by passing the type string “`block_bricked_volume`” to `ospNewVolume`. Because of this rearrangement of voxel data it cannot be shared with the application anymore, but has to be transferred to OSPRay via

```
OSPError ospSetRegion(OSPVolume, void *source,
                      const vec3i &regionCoords,
                      const vec3i &regionSize);
```

¹ For consecutive memory addresses the `x`-index of the corresponding voxel changes the quickest.

Type/Name	Description
OSP_DEVICE	API device object reference
OSP_VOID_PTR	void pointer
OSP_DATA	data reference
OSP_OBJECT	generic object reference
OSP_CAMERA	camera object reference
OSP_FRAMEBUFFER	framebuffer object reference
OSP_LIGHT	light object reference
OSP_MATERIAL	material object reference
OSP_TEXTURE	texture object reference
OSP_RENDERER	renderer object reference
OSP_MODEL	model object reference
OSP_GEOMETRY	geometry object reference
OSP_VOLUME	volume object reference
OSP_TRANSFER_FUNCTION	transfer function object reference
OSP_PIXEL_OP	pixel operation object reference
OSP_STRING	C-style zero-terminated character string
OSP_CHAR	8 bit signed character scalar
OSP_UCHAR	8 bit unsigned character scalar
OSP_UCHAR[234]	... and [234]-element vector
OSP USHORT	16 bit unsigned integer scalar
OSP_INT	32 bit signed integer scalar
OSP_INT[234]	... and [234]-element vector
OSP_UINT	32 bit unsigned integer scalar
OSP_UINT[234]	... and [234]-element vector
OSP_LONG	64 bit signed integer scalar
OSP_LONG[234]	... and [234]-element vector
OSP ULONG	64 bit unsigned integer scalar
OSP ULONG[234]	... and [234]-element vector
OSP_FLOAT	32 bit single precision floating-point scalar
OSP_FLOAT[234]	... and [234]-element vector
OSP_FLOAT3A	... and 3-element vector with padding (same size as an OSP_FLOAT4)
OSP_DOUBLE	64 bit double precision floating-point scalar

Table 4.5 – Valid named constants for OSPDataType.

The voxel data pointed to by `source` is copied into the given volume starting at position `regionCoords`, must be of size `regionSize` and be placed in memory in xyz-order. Note that OSPRay distinguishes between volume data and volume parameters. This function must be called only after all volume parameters (in particular `dimensions` and `voxelType`, see below) have been set and *before* `ospCommit(volume)` is called. If necessary then memory for the volume is allocated on the first call to this function.

The common parameters understood by both structured volume variants are summarized in the table below.

Table 4.6 – Configuration parameters shared by all volume types.

Type	Name	Default	Description
OSPTTransferFunction	transferFunction		transfer function to use
vec2f	voxelRange		minimum and maximum of the scalar values
bool	gradientShadingEnabled	false	volume is rendered with surface shading wrt. to normalized gradient
bool	preIntegration	false	use pre-integration for transfer function lookups
bool	singleShade	true	shade only at the point of maximum intensity
bool	adaptiveSampling	true	adapt ray step size based on opacity
float	adaptiveScalar	15	modifier for adaptive step size
float	adaptiveMaxSamplingRate	2	maximum sampling rate for adaptive sampling
float	samplingRate	0.125	sampling rate of the volume (this is the minimum step size for adaptive sampling)
vec3f	specular	gray 0.3	specular color for shading
vec3f	volumeClippingBoxLower	disabled	lower coordinate (in object-space) to clip the volume values
vec3f	volumeClippingBoxUpper	disabled	upper coordinate (in object-space) to clip the volume values

Table 4.7 – Additional configuration parameters for structured volumes.

Type	Name	Default	Description
vec3i	dimensions		number of voxels in each dimension (x, y, z)
string	voxelType		data type of each voxel, currently supported are: “uchar” (8 bit unsigned integer) “short” (16 bit signed integer) “ushort” (16 bit unsigned integer) “float” (32 bit single precision floating point) “double” (64 bit double precision floating point)
vec3f	gridOrigin	(0, 0, 0)	origin of the grid in world-space
vec3f	gridSpacing	(1, 1, 1)	size of the grid cells in world-space

4.3.2 Adaptive Mesh Refinement (AMR) Volume

AMR volumes are specified as a list of bricks, which are levels of refinement in potentially overlapping regions. There can be any number of refinement levels and any number of bricks at any level of refinement. An AMR volume type is created by passing the type string “amr_volume” to `ospNewVolume`.

Applications should first create an `OSPData` array which holds information about each brick. The following structure is used to populate this array (found in `ospray.h`):

```
struct amr_brick_info
{
    box3i bounds;
    int refinementLevel;
    float cellWidth;
};
```

Then for each brick, the application should create an `OSPData` array of `OSPData` handles, where each handle is the data per-brick. Currently we only support float voxels.

Table 4.8 – Additional configuration parameters for AMR volumes.

Type	Name	Default	Description
<code>vec3f</code>	<code>gridOrigin</code>	<code>(0, 0, 0)</code>	origin of the grid in world-space
<code>vec3f</code>	<code>gridSpacing</code>	<code>(1, 1, 1)</code>	size of the grid cells in world-space
<code>string</code>	<code>amrMethod</code>	<code>current</code>	sampling method; valid values are “finest”, “current”, or “octant”
<code>string</code>	<code>voxelType</code>	<code>undefined</code>	data type of each voxel, currently supported are: “uchar” (8 bit unsigned integer) “short” (16 bit signed integer) “ushort” (16 bit unsigned integer) “float” (32 bit single precision floating point) “double” (64 bit double precision floating point)
<code>OSPData</code>	<code>brickInfo</code>		array of info defining each brick
<code>OSPData</code>	<code>brickData</code>		array of handles to per-brick voxel data

Lastly, note that the `gridOrigin` and `gridSpacing` parameters act just like the structured volume equivalent, but they only modify the root (coarsest level) of refinement.

4.3.3 Unstructured Volumes

Unstructured volumes can contain tetrahedral, wedge, or hexahedral cell types, and are defined by three arrays: vertices, corresponding field values, and eight indices per cell (first four are -1 for tetrahedral cells, first two are -2 for wedge cells). An unstructured volume type is created by passing the type string “`unstructured_volume`” to `ospNewVolume`.

Field values can be specified per-vertex (`field`) or per-cell (`cellField`). If both values are set, `cellField` takes precedence.

Similar to `triangle mesh`, each tetrahedron is formed by a group of indices into the vertices. For each vertex, the corresponding (by array index) data value will be used for sampling when rendering. Note that the index order for each tetrahedron does not matter, as OSPRay internally calculates vertex normals to ensure proper sampling and interpolation.

For wedge cells, each wedge is formed by a group of six indices into the vertices and data value. Vertex ordering is the same as `VTK_WEDGE`: three bottom vertices counterclockwise, then top three counterclockwise.

For hexahedral cells, each hexahedron is formed by a group of eight indices into the vertices and data value. Vertex ordering is the same as `VTK_HEXAHEDRON`: four bottom vertices counterclockwise, then top four counterclockwise.

4.3.4 Transfer Function

Transfer functions map the scalar values of volumes to color and opacity and thus they can be used to visually emphasize certain features of the volume. To create a new transfer function of given type use

```
OSPTransferFunction ospNewTransferFunction(const char *type);
```

Table 4.9 – Additional configuration parameters for unstructured volumes.

Type	Name	Default	Description
vec3f[]	vertices		data array of vertex positions
float[]	field		data array of vertex data values to be sampled
float[]	cellField		data array of cell data values to be sampled
vec4i[]	indices		data array of tetrahedra indices (into vertices and field)
string	hexMethod	planar	“planar” (faster, assumes planar sides) or “nonplanar”
bool	precomputedNormals	true	whether to accelerate by precomputing, at a cost of 72 bytes/cell

The call returns `NULL` if that type of transfer functions is not known by OSPRay, or else an `OSPTransferFunction` handle to the created transfer function. That handle can be assigned to a volume as parameter “`transferFunction`” using `ospSetObject`.

One type of transfer function that is supported by OSPRay is the linear transfer function, which interpolates between given equidistant colors and opacities. It is created by passing the string “`piecewise_linear`” to `ospNewTransferFunction` and it is controlled by these parameters:

Type	Name	Description
vec3f[]	colors	data array of RGB colors
float[]	opacities	data array of opacities
vec2f	valueRange	domain (scalar range) this function maps from

Table 4.10 – Parameters accepted by the linear transfer function.

4.4 Geometries

Geometries in OSPRay are objects that describe surfaces. To create a new geometry object of given type use

```
OSPGeometry ospNewGeometry(const char *type);
```

The call returns `NULL` if that type of geometry is not known by OSPRay, or else an `OSPGeometry` handle.

4.4.1 Triangle Mesh

A traditional triangle mesh (indexed face set) geometry is created by calling `ospNewGeometry` with type string “`triangles`”. Once created, a triangle mesh recognizes the following parameters:

The `vertex` and `index` arrays are mandatory to create a valid triangle mesh.

4.4.2 Quad Mesh

A mesh consisting of quads is created by calling `ospNewGeometry` with type string “`quads`”. Once created, a quad mesh recognizes the following parameters:

The `vertex` and `index` arrays are mandatory to create a valid quad mesh. A quad is internally handled as a pair of two triangles, thus mixing triangles and quad is supported by encoding a triangle as a quad with the last two vertex indices being identical (`w=z`).

Type	Name	Description
vec3f(a)[]	vertex	data array of vertex positions
vec3f(a)[]	vertex.normal	data array of vertex normals
vec4f[] / vec3fa[]	vertex.color	data array of vertex colors (RGBA/RGB)
vec2f[]	vertex.texcoord	data array of vertex texture coordinates
vec3i(a)[]	index	data array of triangle indices (into the vertex array(s))

Table 4.11 – Parameters defining a triangle mesh geometry.

Type	Name	Description
vec3f(a)[]	vertex	data array of vertex positions
vec3f(a)[]	vertex.normal	data array of vertex normals
vec4f[] / vec3fa[]	vertex.color	data array of vertex colors (RGBA/RGB)
vec2f[]	vertex.texcoord	data array of vertex texture coordinates
vec4i[]	index	data array of quad indices (into the vertex array(s))

Table 4.12 – Parameters defining a quad mesh geometry.

4.4.3 Subdivision

A mesh consisting of subdivision surfaces, created by specifying a geometry of type “subdivision”. Once created, a subdivision recognizes the following parameters:

Type	Name	Description
vec3f[]	vertex	data array of vertex positions
vec4f[]	vertex.color	data array of vertex colors (RGBA)
vec2f[]	vertex.texcoord	data array of vertex texture coordinates
float	level	global level of tessellation, default is 5
uint[]/vec4i[]	index	data array of indices (into the vertex array(s))
float[]	index.level	data array of per-edge levels of tessellation, overrides global level
uint[]	face	data array holding the number of indices/edges (3 to 15) per face
vec2i[]	edgeCrease.index	data array of edge crease indices
float[]	edgeCrease.weight	data array of edge crease weights
uint[]	vertexCrease.inde	x data array of vertex crease indices
float[]	vertexCrease.weig	ht data array of vertex crease weights

Table 4.13 – Parameters defining a Subdivision geometry.

The `vertex` and `index` arrays are mandatory to create a valid subdivision surface. If no `face` array is present then a pure quad mesh is assumed (and indices must be of type `vec4i`). Optionally supported are edge and vertex creases.

4.4.4 Spheres

A geometry consisting of individual spheres, each of which can have an own radius, is created by calling `ospNewGeometry` with type string “spheres”. The spheres will not be tessellated but rendered procedurally and are thus perfectly round. To allow a variety of sphere representations in the application this geometry allows a flexible way of specifying the data of center position and radius within a `data` array:

Table 4.14 – Parameters defining a spheres geometry.

Type	Name	Default	Description
float	radius	0.01	radius of all spheres (if <code>offset_radius</code> is not used)
OSPData	spheres	NULL	memory holding the spatial <code>data</code> of all spheres
int	bytes_per_sphere	16	size (in bytes) of each sphere within the <code>spheres</code> array
int	offset_center	0	offset (in bytes) of each sphere’s “vec3f center” position (in object-space) within the <code>spheres</code> array
int	offset_radius	-1	offset (in bytes) of each sphere’s “float radius” within the <code>spheres</code> array (-1 means disabled and use <code>radius</code>)
int	offset_colorID	-1	offset (in bytes) of each sphere’s “int colorID” within the <code>spheres</code> array (-1 means disabled and use the shared material color)
vec4f[] / vec3f(a)[] / vec4uc	color	NULL	<code>data</code> array of colors (RGBA/RGB), color is constant for each sphere
int	color_offset	0	offset (in bytes) to the start of the color data in <code>color</code>
int	color_format	<code>color.data_type</code>	the format of the color data. Can be one of: <code>OSP_FLOAT4</code> , <code>OSP_FLOAT3</code> , <code>OSP_FLOAT3A</code> or <code>OSP_UCHAR4</code> . Defaults to the type of data in <code>color</code>
int	color_stride	<code>sizeof(color_format)</code>	stride (in bytes) between each color element in the <code>color</code> array. Defaults to the size of a single element of type <code>color_format</code>
vec2f[]	texcoord	NULL	<code>data</code> array of texture coordinates, coordinate is constant for each sphere

4.4.5 Cylinders

A geometry consisting of individual cylinders, each of which can have an own radius, is created by calling `ospNewGeometry` with type string “cylinders”. The cylinders will not be tessellated but rendered procedurally and are thus perfectly round. To allow a variety of cylinder representations in the application this geometry allows a flexible way of specifying the data of offsets for start position, end position and radius within a `data` array. All parameters are listed in the table below.

For texturing each cylinder is seen as a 1D primitive, i.e., a line segment: the 2D texture coordinates at its vertices v0 and v1 are linearly interpolated.

Table 4.15 – Parameters defining a cylinders geometry.

Type	Name	Default	Description
float	radius	0.01	radius of all cylinders (if offset_radius is not used)
OSPData	cylinders	NULL	memory holding the spatial data of all cylinders
int	bytes_per_cylinder	24	size (in bytes) of each cylinder within the cylinders array
int	offset_v0	0	offset (in bytes) of each cylinder's "vec3f v0" position (the start vertex, in object-space) within the cylinders array
int	offset_v1	12	offset (in bytes) of each cylinder's "vec3f v1" position (the end vertex, in object-space) within the cylinders array
int	offset_radius	-1	offset (in bytes) of each cylinder's "float radius" within the cylinders array (-1 means disabled and use radius instead)
vec4f[] / vec3f(a)[]	color	NULL	data array of colors (RGBA/RGB), color is constant for each cylinder
OSPData	texcoord	NULL	data array of texture coordinates, in pairs (each a vec2f at vertex v0 and v1)

4.4.6 Streamlines

A geometry consisting of multiple streamlines is created by calling `ospNewGeometry` with type string “streamlines”. The streamlines are internally assembled either from connected (and rounded) cylinder segments, or represented as Bézier curves; they are thus always perfectly round. The parameters defining this geometry are listed in the table below.

Table 4.16 – Parameters defining a streamlines geometry.

Type	Name	Description
float	radius	global radius of all streamlines (if per-vertex radius is not used), default 0.01
bool	smooth	enable curve interpolation, default off (always on if per-vertex radius is used)
vec3fa[] / vec4f[]	vertex	data array of all vertex position (and optional radius) for <i>all</i> streamlines
vec4f[]	vertex.color	data array of corresponding vertex colors (RGBA)
float[]	vertex.radius	data array of corresponding vertex radius
int32[]	index	data array of indices to the first vertex of a link

Each streamline is specified by a set of (aligned) control points in `vertex`. If `smooth` is disabled and a constant `radius` is used for all streamlines then all vertices belonging to the same logical streamline are connected via [cylinders](#), with additional [spheres](#) at each vertex to create a continuous, closed surface. Otherwise, streamlines are represented as Bézier curves, smoothly interpolating the vertices. This mode supports per-vertex varying radii (either given in `vertex.radius`, or in the 4th component of a `vec4f` `vertex`), but is slower and consumes more memory. Additionally, the radius needs to be smaller than the curvature radius of the Bézier curve at each location on the curve.

A streamlines geometry can contain multiple disjoint streamlines, each streamline is specified as a list of segments (or links) referenced via `index`: each entry `e` of the `index` array points the first vertex of a link (`vertex[index[e]]`) and the second vertex of the link is implicitly the directly following one (`vertex[index[e]+1]`). For example, two streamlines of vertices (A-B-C-D) and (E-F-G), respectively, would internally correspond to five links (A-B, B-C, C-D,

E-F, and F-G), and would be specified via an array of vertices [A,B,C,D,E,F,G], plus an array of link indices [0,1,2,4,5].

4.4.7 Curves

A geometry consisting of multiple curves is created by calling `ospNewGeometry` with type string “curves”. The parameters defining this geometry are listed in the table below.

Table 4.17 – Parameters defining a curves geometry.

Type	Name	Description
string	curveType	“flat” (ray oriented), “round” (circular cross section), “ribbon” (normal oriented flat curve)
string	curveBasis	“linear”, “bezier”, “bspline”, “hermite”
vec4f[]	vertex	<code>data</code> array of vertex position and radius
int32[]	index	<code>data</code> array of indices to the first vertex or tangent of a curve segment
vec3f[]	vertex.normal	<code>data</code> array of curve normals (only for “ribbon” curves)
vec3f[]	vertex.tangent	<code>data</code> array of curve tangents (only for “hermite” curves)

See Embree documentation for discussion of curve types and data formatting.

4.4.8 Isosurfaces

OSPRay can directly render multiple isosurfaces of a volume without first tessellating them. To do so create an isosurfaces geometry by calling `ospNewGeometry` with type string “isosurfaces”. Each isosurface will be colored according to the provided volume’s [transfer function](#).

Type	Name	Description
float[]	isovalues	<code>data</code> array of isovalues
OSPVolume	volume	handle of the volume to be isosurfaced

Table 4.18 – Parameters defining an isosurfaces geometry.

4.4.9 Slices

One tool to highlight notable features of volumetric data is to visualize 2D cuts (or slices) by placing planes into the volume. Such a slices geometry is created by calling `ospNewGeometry` with type string “slices”. The planes are defined by the coefficients (a, b, c, d) of the plane equation $ax + by + cz + d = 0$. Each slice is colored according to the provided volume’s [transfer function](#).

Type	Name	Description
vec4f[]	planes	<code>data</code> array with plane coefficients for all slices
OSPVolume	volume	handle of the volume that will be sliced

Table 4.19 – Parameters defining a slices geometry.

4.4.10 Instances

OSPRay supports instancing via a special type of geometry. Instances are created by transforming another given `model` `modelToInstantiate` with the given affine transformation `transform` by calling

```
OSPGeometry ospNewInstance(OSPModel modelToInstantiate, const affine3f &transform);
```

4.5 Renderer

A renderer is the central object for rendering in OSPRay. Different renderers implement different features and support different materials. To create a new renderer of given type use

```
OSPRenderer ospNewRenderer(const char *type);
```

The call returns NULL if that type of renderer is not known, or else an OSPRenderer handle to the created renderer. General parameters of all renderers are

Table 4.20 – Parameters understood by all renderers.

Type	Name	Default	Description
OSPModel	model		the model to render
OSPCamera	camera		the camera to be used for rendering
OSPLight[]	lights		data array with handles of the lights
int	spp	1	samples per pixel
int	maxDepth	20	maximum ray recursion depth
float	minContribution	0.001	sample contributions below this value will be neglected to speedup rendering
float	varianceThreshold	0	threshold for adaptive accumulation

OSPRay's renderers support a feature called adaptive accumulation, which accelerates progressive [rendering](#) by stopping the rendering and refinement of image regions that have an estimated variance below the [varianceThreshold](#). This feature requires a [framebuffer](#) with an [OSP_FB_VARIANCE](#) channel.

4.5.1 SciVis Renderer

The SciVis renderer is a fast ray tracer for scientific visualization which supports volume rendering and ambient occlusion (AO). It is created by passing the type string "scivis" or "raytracer" to [ospNewRenderer](#). In addition to the [general parameters](#) understood by all renderers the SciVis renderer supports the following special parameters:

Note that the intensity (and color) of AO is deduced from an [ambient light](#) in the [lights](#) array.² If [aoSamples](#) is zero (the default) then ambient lights cause ambient illumination (without occlusion).

Per default the background of the rendered image will be transparent black, i.e., the alpha channel holds the opacity of the rendered objects. This eases transparency-aware blending of the image with an arbitrary background image by the application. The parameter [bgColor](#) can be used to already blend with a constant background color (and alpha) during rendering.

The SciVis renderer supports depth composition with images of other renderers, for example to incorporate help geometries of a 3D UI that were rendered with OpenGL. The screen-sized [texture](#) [maxDepthTexture](#) must have format [OSP_TEXTURE_R32F](#) and flag [OSP_TEXTURE_FILTER_NEAREST](#). The fetched values are used to limit the distance of primary rays, thus objects of other renderers can hide objects rendered by OSPRay.

² If there are multiple ambient lights then their contribution is added

Table 4.21 – Special parameters understood by the SciVis renderer.

Type	Name	Default	Description
bool	shadowsEnabled	false	whether to compute (hard) shadows
int	aoSamples	0	number of rays per sample to compute ambient occlusion
float	aoDistance	10^{20}	maximum distance to consider for ambient occlusion
bool	aoTransparencyEnabled	false	whether object transparency is respected when computing ambient occlusion (slower)
bool	oneSidedLighting	true	if true, backfacing surfaces (wrt. light source) receive no illumination
float / vec3f / vec4f	bgColor	black, transparent	background color and alpha (RGBA)
OSPTexture	maxDepthTexture	NULL	screen-sized float texture with maximum far distance per pixel (use texture type ‘texture2d’)

4.5.2 Path Tracer

The path tracer supports soft shadows, indirect illumination and realistic materials. This renderer is created by passing the type string “pathtracer” to `ospNewRenderer`. In addition to the [general parameters](#) understood by all renderers the path tracer supports the following special parameters:

Table 4.22 – Special parameters understood by the path tracer.

Type	Name	Default	Description
int	rouletteDepth	5	ray recursion depth at which to start Russian roulette termination
float	maxContribution	∞	samples are clamped to this value before they are accumulated into the framebuffer
OSPTexture	backplate	NULL	texture image used as background, replacing visible lights in infinity (e.g., the HDRI light)

The path tracer requires that [materials](#) are assigned to [geometries](#), otherwise surfaces are treated as completely black.

4.5.3 Model

Models are a container of scene data. They can hold the different [geometries](#) and [volumes](#) as well as references to (and [instances](#) of) other models. A model is associated with a single logical acceleration structure. To create an (empty) model call

```
OSPModel ospNewModel();
```

The call returns an `OSPModel` handle to the created model. To add an already created geometry or volume to a model use

```
void ospAddGeometry(OSPModel, OSPGeometry);
void ospAddVolume(OSPModel, OSPVolume);
```

An existing geometry or volume can be removed from a model with

```
void ospRemoveGeometry(OSPModel, OSPGeometry);
void ospRemoveVolume(OSPModel, OSPVolume);
```

Finally, Models can be configured with parameters for making various feature/performance trade-offs:

Table 4.23 – Parameters understood by Models

Type	Name	Default	Description
bool	dynamicScene	false	use RTC_SCENE_DYNAMIC flag (faster BVH build, slower ray traversal), otherwise uses RTC_SCENE_STATIC flag (faster ray traversal, slightly slower BVH build)
bool	compactMode	false	tell Embree to use a more compact BVH in memory by trading ray traversal performance
bool	robustMode	false	tell Embree to enable more robust ray intersection code paths (slightly slower)

4.5.4 Lights

To create a new light source of given type type use

```
OSPLight ospNewLight3(const char *type);
```

The call returns NULL if that type of light is not known by the renderer, or else an OSPLight handle to the created light source. All light sources³ accept the following parameters:

Type	Name	Default	Description
vec3f(a)	color	white	color of the light
float	intensity	1	intensity of the light (a factor)
bool	isVisible	true	whether the light can be directly seen

The following light types are supported by most OSPRay renderers.

4.5.4.1 Directional Light / Distant Light

The distant light (or traditionally the directional light) is thought to be far away (outside of the scene), thus its light arrives (almost) as parallel rays. It is created by passing the type string “distant” to ospNewLight3. In addition to the general parameters understood by all lights the distant light supports the following special parameters:

Type	Name	Description
vec3f(a)	direction	main emission direction of the distant light
float	angularDiameter	apparent size (angle in degree) of the light

Setting the angular diameter to a value greater than zero will result in soft shadows when the renderer uses stochastic sampling (like the path tracer). For instance, the apparent size of the sun is about 0.53°.

³ The HDRI light is an exception, it knows about intensity, but not about color.

Table 4.24 – Parameters accepted by all lights.

Table 4.25 – Special parameters accepted by the distant light.

4.5.4.2 Point Light / Sphere Light

The sphere light (or the special case point light) is a light emitting uniformly in all directions. It is created by passing the type string “sphere” to `ospNewLight3`. In addition to the [general parameters](#) understood by all lights the sphere light supports the following special parameters:

Type	Name	Description
<code>vec3f(a)</code>	<code>position</code>	the center of the sphere light, in world-space
<code>float</code>	<code>radius</code>	the size of the sphere light

Table 4.26 – Special parameters accepted by the sphere light.

Setting the radius to a value greater than zero will result in soft shadows when the renderer uses stochastic sampling (like the [path tracer](#)).

4.5.4.3 Spotlight

The spotlight is a light emitting into a cone of directions. It is created by passing the type string “spot” to `ospNewLight3`. In addition to the [general parameters](#) understood by all lights the spotlight supports the special parameters listed in the table.

Table 4.27 – Special parameters accepted by the spotlight.

Type	Name	Description
<code>vec3f(a)</code>	<code>position</code>	the center of the spotlight, in world-space
<code>vec3f(a)</code>	<code>direction</code>	main emission direction of the spot
<code>float</code>	<code>openingAngle</code>	full opening angle (in degree) of the spot; outside of this cone is no illumination
<code>float</code>	<code>penumbraAngle</code>	size (angle in degree) of the “penumbra”, the region between the rim (of the illumination cone) and full intensity of the spot; should be smaller than half of <code>openingAngle</code>
<code>float</code>	<code>radius</code>	the size of the spotlight, the radius of a disk with normal <code>direction</code>




Figure 4.1 – Angles used by the spotlight.

Setting the radius to a value greater than zero will result in soft shadows when the renderer uses stochastic sampling (like the [path tracer](#)).

4.5.4.4 Quad Light

The quad⁴ light is a planar, procedural area light source emitting uniformly on one side into the half-space. It is created by passing the type string “quad” to `ospNewLight3`. In addition to the [general parameters](#) understood by all lights the quad light supports the following special parameters:

⁴ actually a parallelogram

Type	Name	Description
vec3f(a)	position	world-space position of one vertex of the quad light
vec3f(a)	edge1	vector to one adjacent vertex
vec3f(a)	edge2	vector to the other adjacent vertex




Table 4.28 – Special parameters accepted by the quad light.

Figure 4.2 – Defining a quad light which emits toward the reader.

The emission side is determined by the cross product of $\text{edge1} \times \text{edge2}$. Note that only renderers that use stochastic sampling (like the path tracer) will compute soft shadows from the quad light. Other renderers will just sample the center of the quad light, which results in hard shadows.

4.5.4.5 HDRI Light

The HDRI light is a textured light source surrounding the scene and illuminating it from infinity. It is created by passing the type string “hdri” to `ospNewLight3`. In addition to the [parameter intensity](#) the HDRI light supports the following special parameters:

Table 4.29 – Special parameters accepted by the HDRI light.

Type	Name	Description
vec3f(a)	up	up direction of the light in world-space
vec3f(a)	dir	direction to which the center of the texture will be mapped to (analog to panoramic camera)
OSPTexture	map	environment map in latitude / longitude format

Note that the currently only the [path tracer](#) supports the HDRI light.

4.5.4.6 Ambient Light

The ambient light surrounds the scene and illuminates it from infinity with constant radiance (determined by combining the [parameters color and intensity](#)). It is created by passing the type string “ambient” to `ospNewLight3`.

Note that the [SciVis renderer](#) uses ambient lights to control the color and intensity of the computed ambient occlusion (AO).

4.5.4.7 Emissive Objects

The [path tracer](#) will consider illumination by [geometries](#) which have a light emitting material assigned (for example the [Luminous](#) material).

4.5.5 Materials

Materials describe how light interacts with surfaces, they give objects their distinctive look. To let the given renderer create a new material of given type type call




Figure 4.3 – Orientation and Mapping of an HDRI Light.

```
OSPMaterial ospNewMaterial2(const char *renderer_type, const char *material_type);
```

The call returns NULL if the material type is not known by the renderer type, or else an OSPMaterial handle to the created material. The handle can then be used to assign the material to a given geometry with

```
void ospSetMaterial(OSPGeometry, OSPMaterial);
```

4.5.5.1 OBJ Material

The OBJ material is the workhorse material supported by both the [SciVis renderer](#) and the [path tracer](#). It offers widely used common properties like diffuse and specular reflection and is based on the [MTL material format](#) of Lightwave's OBJ scene files. To create an OBJ material pass the type string "OBJMaterial" to `ospNewMaterial2`. Its main parameters are

Type	Name	Default	Description
vec3f	Kd	white 0.8	diffuse color
vec3f	Ks	black	specular color
float	Ns	10	shininess (Phong exponent), usually in [2-10 ⁴]
float	d	opaque	opacity
vec3f	Tf	black	transparency filter color
OSPTexture	map_Bump	NULL	normal map

Table 4.30 – Main parameters of the OBJ material.

In particular when using the path tracer it is important to adhere to the principle of energy conservation, i.e., that the amount of light reflected by a surface is not larger than the light arriving. Therefore the path tracer issues a warning and renormalizes the color parameters if the sum of Kd, Ks, and Tf is larger than one in any color channel. Similarly important to mention is that almost all materials of the real world reflect at most only about 80% of the incoming light. So even for a white sheet of paper or white wall paint do better not set Kd larger than 0.8; otherwise rendering times are unnecessary long and the contrast in the final images is low (for example, the corners of a white room would hardly be discernible, as can be seen in the figure below).





Figure 4.4 – Comparison of diffuse rooms with 100% reflecting white paint (left) and realistic 80% reflecting white paint (right), which leads to higher overall contrast. Note that exposure has been adjusted to achieve similar brightness levels.

If present, the color component of [geometries](#) is also used for the diffuse color K_d and the alpha component is also used for the opacity d .

Note that currently only the path tracer implements colored transparency with T_f .

Normal mapping can simulate small geometric features via the texture `map_Bump`. The normals n in the normal map are wrt. the local tangential shading coordinate system and are encoded as $1/2(n + 1)$, thus a texel $(0.5, 0.5, 1)$ ⁵ represents the unperturbed shading normal $(0, 0, 1)$. Because of this encoding an sRGB gamma [texture](#) format is ignored and normals are always fetched as linear from a normal map. Note that the orientation of normal maps is important for a visually consistent look: by convention OSPRay uses a coordinate system with the origin in the lower left corner; thus a convexity will look green toward the top of the texture image (see also the example image of a normal map). If this is not the case flip the normal map vertically or invert its green channel.



⁵ respectively $(127, 127, 255)$ for 8 bit textures

Figure 4.5 – Normal map representing an exalted square pyramidal frustum.

All parameters (except T_f) can be textured by passing a [texture](#) handle, pre-fixed with “`map_`”. The fetched texels are multiplied by the respective parameter value. Texturing requires [geometries](#) with texture coordinates, e.g., a [triangle mesh](#) with `vertex.texcoord` provided. The color textures `map_Kd` and `map_Ks` are typically in one of the sRGB gamma encoded formats, whereas textures `map_Ns` and `map_d` are usually in a linear format (and only the first component is used). Additionally, all textures support [texture transformations](#).




Figure 4.6 – Rendering of a OBJ material with wood textures.

4.5.5.2 Principled

The Principled material is the most complex material offered by the [path tracer](#), which is capable of producing a wide variety of materials (e.g., plastic, metal, wood, glass) by combining multiple different layers and lobes. It uses the GGX microfacet distribution with approximate multiple scattering for dielectrics and

metals, uses the Oren-Nayar model for diffuse reflection, and is energy conserving. To create a Principled material, pass the type string “Principled” to `ospNewMaterial2`. Its parameters are listed in the table below.

Table 4.31 – Parameters of the Principled material.

Type	Name	Default	Description
vec3f	baseColor	white 0.8	base reflectivity (diffuse and/or metallic)
vec3f	edgeColor	white	edge tint (metallic only)
float	metallic	0	mix between dielectric (diffuse and/or specular) and metallic (specular only with complex IOR) in [0–1]
float	diffuse	1	diffuse reflection weight in [0–1]
float	specular	1	specular reflection/transmission weight in [0–1]
float	ior	1	dielectric index of refraction
float	transmission	0	specular transmission weight in [0–1]
vec3f	transmissionColor	white	attenuated color due to transmission (Beer’s law)
float	transmissionDepth	1	distance at which color attenuation is equal to transmissionColor
float	roughness	0	diffuse and specular roughness in [0–1], 0 is perfectly smooth
float	anisotropy	0	amount of specular anisotropy in [0–1]
float	rotation	0	rotation of the direction of anisotropy in [0–1], 1 is going full circle
float	normal	1	default normal map/scale for all layers
float	baseNormal	1	base normal map/scale (overrides default normal)
bool	thin	false	flag specifying whether the material is thin or solid
float	thickness	1	thickness of the material (thin only), affects the amount of color attenuation due to specular transmission
float	backlight	0	amount of diffuse transmission (thin only) in [0–2], 1 is 50% reflection and 50% transmission, 2 is transmission only
float	coat	0	clear coat layer weight in [0–1]
float	coatIor	1.5	clear coat index of refraction
vec3f	coatColor	white	clear coat color tint
float	coatThickness	1	clear coat thickness, affects the amount of color attenuation
float	coatRoughness	0	clear coat roughness in [0–1], 0 is perfectly smooth
float	coatNormal	1	clear coat normal map/scale (overrides default normal)
float	sheen	0	sheen layer weight in [0–1]
vec3f	sheenColor	white	sheen color tint
float	sheenTint	0	how much sheen is tinted from sheenColor toward baseColor
float	sheenRoughness	0.2	sheen roughness in [0–1], 0 is perfectly smooth
float	opacity	1	cut-out opacity/transparency, 1 is fully opaque

All parameters can be textured by passing a `texture` handle, suffixed with “Map” (e.g., “`baseColorMap`”); [texture transformations](#) are supported as well.

4.5.5.3 CarPaint

The CarPaint material is a specialized version of the Principled material for rendering different types of car paints. To create a CarPaint material, pass the type string “CarPaint” to `ospNewMaterial2`. Its parameters are listed in the table below.

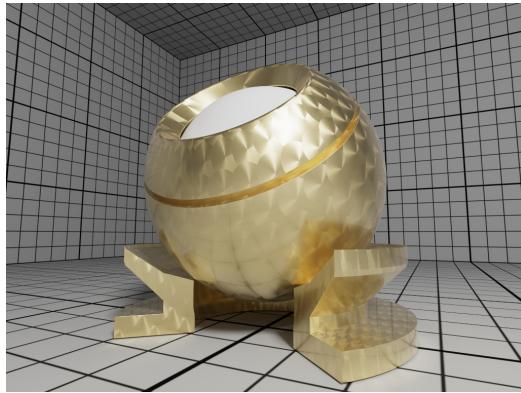


Figure 4.7 – Rendering of a Principled coated brushed metal material with textured anisotropic rotation and a dust layer (sheen) on top.

Table 4.32 – Parameters of the CarPaint material.

Type	Name	Default	Description
vec3f	baseColor	white 0.8	diffuse base reflectivity
float	roughness	0	diffuse roughness in [0–1], 0 is perfectly smooth
float	normal	1	normal map/scale
float	flakeDensity	0	density of metallic flakes in [0–1], 0 disables flakes, 1 fully covers the surface with flakes
float	flakeScale	100	scale of the flake structure, higher values increase the amount of flakes
float	flakeSpread	0.3	flake spread in [0–1]
float	flakeJitter	0.75	flake randomness in [0–1]
float	flakeRoughness	0.3	flake roughness in [0–1], 0 is perfectly smooth
float	coat	1	clear coat layer weight in [0–1]
float	coatIor	1.5	clear coat index of refraction
vec3f	coatColor	white	clear coat color tint
float	coatThickness	1	clear coat thickness, affects the amount of color attenuation
float	coatRoughness	0	clear coat roughness in [0–1], 0 is perfectly smooth
float	coatNormal	1	clear coat normal map/scale
vec3f	flipflopColor	white	reflectivity of coated flakes at grazing angle, used together with coatColor produces a pearlescent paint
float	flipflopFalloff	1	flip flop color falloff, 1 disables the flip flop effect

All parameters can be textured by passing a [texture](#) handle, suffixed with “Map” (e.g., “baseColorMap”); [texture transformations](#) are supported as well.

4.5.5.4 Metal

The [path tracer](#) offers a physical metal, supporting changing roughness and realistic color shifts at edges. To create a Metal material pass the type string “Metal” to `ospNewMaterial2`. Its parameters are

The main appearance (mostly the color) of the Metal material is controlled by the physical parameters `eta` and `k`, the wavelength-dependent, complex index of refraction. These coefficients are quite counterintuitive but can be found in [published measurements](#). For accuracy the index of refraction can be given as an array of spectral samples in `ior`, each sample a triplet of wavelength (in nm), `eta`, and `k`, ordered monotonically increasing by wavelength; OSPRay will then calculate the Fresnel in the spectral domain. Alternatively, `eta` and `k` can also be specified as approximated RGB coefficients; some examples are given in below




Figure 4.8 – Rendering of a pearlescent CarPaint material.

Table 4.33 – Parameters of the Metal material.

Type	Name	Default	Description
vec3f[]	ior	Aluminium	<code>data</code> array of spectral samples of complex refractive index, each entry in the form (wavelength, eta, k), ordered by wavelength (which is in nm)
vec3f	eta		RGB complex refractive index, real part
vec3f	k		RGB complex refractive index, imaginary part
float	roughness	0.1	roughness in [0–1], 0 is perfect mirror

table.

Metal	eta	k
Ag, Silver	(0.051, 0.043, 0.041)	(5.3, 3.6, 2.3)
Al, Aluminium	(1.5, 0.98, 0.6)	(7.6, 6.6, 5.4)
Au, Gold	(0.07, 0.37, 1.5)	(3.7, 2.3, 1.7)
Cr, Chromium	(3.2, 3.1, 2.3)	(3.3, 3.3, 3.1)
Cu, Copper	(0.1, 0.8, 1.1)	(3.5, 2.5, 2.4)

Table 4.34 – Index of refraction of selected metals as approximated RGB coefficients, based on data from <https://refractiveindex.info/>.

The roughness parameter controls the variation of microfacets and thus how polished the metal will look. The roughness can be modified by a `texture map_roughness` (`texture transformations` are supported as well) to create notable edging effects.

4.5.5.5 Alloy

The `path tracer` offers an alloy material, which behaves similar to `Metal`, but allows for more intuitive and flexible control of the color. To create an Alloy material pass the type string “Alloy” to `ospNewMaterial2`. Its parameters are

Type	Name	Default	Description
vec3f	color	white 0.9	reflectivity at normal incidence (0 degree)
vec3f	edgeColor	white	reflectivity at grazing angle (90 degree)
float	roughness	0.1	roughness, in [0–1], 0 is perfect mirror

Table 4.35 – Parameters of the Alloy material.

The main appearance of the Alloy material is controlled by the parameter `color`, while `edgeColor` influences the tint of reflections when seen at grazing




Figure 4.9 – Rendering of golden Metal material with textured roughness.

angles (for real metals this is always 100% white). If present, the color component of [geometries](#) is also used for reflectivity at normal incidence [color](#). As in [Metal](#) the [roughness](#) parameter controls the variation of microfacets and thus how polished the alloy will look. All parameters can be textured by passing a [texture](#) handle, prefixed with “`map_`”; [texture transformations](#) are supported as well.




Figure 4.10 – Rendering of a fictional Alloy material with textured color.

4.5.5.6 Glass

The [path tracer](#) offers a realistic a glass material, supporting refraction and volumetric attenuation (i.e., the transparency color varies with the geometric thickness). To create a Glass material pass the type string “Glass” to `ospNewMaterial2`. Its parameters are

Type	Name	Default	Description
float	eta	1.5	index of refraction
vec3f	attenuationColor	white	resulting color due to attenuation
float	attenuationDistance	1	distance affecting attenuation

Table 4.36 – Parameters of the Glass material.

For convenience, the rather counterintuitive physical attenuation coefficients will be calculated from the user inputs in such a way, that the [attenuationColor](#) will be the result when white light traveled trough a glass of thickness [attenuationDistance](#).

4.5.5.7 ThinGlass

The [path tracer](#) offers a thin glass material useful for objects with just a single surface, most prominently windows. It models a thin, transparent slab, i.e., it




Figure 4.11 – Rendering of a Glass material with orange attenuation.

behaves as if a second, virtual surface is parallel to the real geometric surface. The implementation accounts for multiple internal reflections between the interfaces (including attenuation), but neglects parallax effects due to its (virtual) thickness. To create a such a thin glass material pass the type string “ThinGlass” to `ospNewMaterial2`. Its parameters are

Type	Name	Default	Description
float	eta	1.5	index of refraction
vec3f	attenuationColor	white	resulting color due to attenuation
float	attenuationDistance	1	distance affecting attenuation
float	thickness	1	virtual thickness

For convenience the attenuation is controlled the same way as with the [Glass](#) material. Additionally, the color due to attenuation can be modulated with a [texture map_attenuationColor](#) ([texture transformations](#) are supported as well). If present, the color component of [geometries](#) is also used for the attenuation color. The `thickness` parameter sets the (virtual) thickness and allows for easy exchange of parameters with the (real) [Glass](#) material; internally just the ratio between `attenuationDistance` and `thickness` is used to calculate the resulting attenuation and thus the material appearance.




Table 4.37 – Parameters of the ThinGlass material.

Figure 4.12 – Rendering of a ThinGlass material with red attenuation.

4.5.5.8 MetallicPaint

The [path tracer](#) offers a metallic paint material, consisting of a base coat with optional flakes and a clear coat. To create a `MetallicPaint` material pass the type string “`MetallicPaint`” to `ospNewMaterial2`. Its parameters are listed in the table below.




Figure 4.13 – Example image of a coloured window made with textured attenuation of the ThinGlass material.

Type	Name	Default	Description
vec3f	baseColor	white 0.8	color of base coat
float	flakeAmount	0.3	amount of flakes, in [0–1]
vec3f	flakeColor	Aluminium	color of metallic flakes
float	flakeSpread	0.5	spread of flakes, in [0–1]
float	eta	1.5	index of refraction of clear coat

Table 4.38 – Parameters of the Metallic-Paint material.

The color of the base coat `baseColor` can be textured by a [texture map](#) `baseColor`, which also supports [texture transformations](#). If present, the color component of `geometries` is also used for the color of the base coat. Parameter `flakeAmount` controls the proportion of flakes in the base coat, so when setting it to 1 the `baseColor` will not be visible. The shininess of the metallic component is governed by `flakeSpread`, which controls the variation of the orientation of the flakes, similar to the `roughness` parameter of [Metal](#). Note that the effect of the metallic flakes is currently only computed on average, thus individual flakes are not visible.




Figure 4.14 – Rendering of a Metallic-Paint material.

4.5.5.9 Luminous

The [path tracer](#) supports the Luminous material which emits light uniformly in all directions and which can thus be used to turn any geometric object into a light source. It is created by passing the type string “Luminous” to `ospNewMaterial2`. The amount of constant radiance that is emitted is determined by combining the general parameters of lights: [color](#) and [intensity](#).




Figure 4.15 – Rendering of a yellow Luminous material.

4.5.6 Texture

OSPRay currently implements two texture types (`texture2d` and `volume`) and is open for extension to other types by applications. More types may be added in future releases.

To create a new texture use

```
OSPTexture ospNewTexture(const char *type);
```

The call returns `NULL` if the texture could not be created with the given parameters, or else an `OSPTexture` handle to the created texture.

4.5.6.1 Texture2D

The `texture2d` texture type implements an image-based texture, where its parameters are as follows

Type	Name	Description
<code>vec2i</code>	<code>size</code>	size of the textures
<code>int</code>	<code>type</code>	<code>OSPTextureFormat</code> for the texture
<code>int</code>	<code>flags</code>	special attribute flags for this texture, currently only responds to <code>OSP_TEXTURE_FILTER_NEAREST</code> or no flags
<code>OSPData</code>	<code>data</code>	the actual texel data

Table 4.39 – Parameters of `texture2d` texture type

The supported texture formats for `texture2d` are:

The texel data addressed by `source` starts with the texels in the lower left corner of the texture image, like in OpenGL. Per default a texture fetch is filtered by performing bi-linear interpolation of the nearest 2×2 texels; if instead fetching only the nearest texel is desired (i.e., no filtering) then pass the `OSP_TEXTURE_FILTER_NEAREST` flag.

4.5.6.2 TextureVolume

The `volume` texture type implements texture lookups based on 3D world coordinates of the surface hit point on the associated geometry. If the given hit point is within the attached volume, the volume is sampled and classified with the transfer function attached to the volume. This implements the ability to visualize volume values (as colored by its transfer function) on arbitrary surfaces inside the volume (as opposed to an isosurface showing a particular value in the volume). Its parameters are as follows

Name	Description
OSP_TEXTURE_RGBA8	8 bit [0–255] linear components red, green, blue, alpha
OSP_TEXTURE_SRGB8	8 bit sRGB gamma encoded color components, and linear alpha
OSP_TEXTURE_RGBA32F	32 bit float components red, green, blue, alpha
OSP_TEXTURE_RGB8	8 bit [0–255] linear components red, green, blue
OSP_TEXTURE_SRGB	8 bit sRGB gamma encoded components red, green, blue
OSP_TEXTURE_RGB32F	32 bit float components red, green, blue
OSP_TEXTURE_R8	8 bit [0–255] linear single component
OSP_TEXTURE_RA8	8 bit [0–255] linear two component
OSP_TEXTURE_L8	8 bit [0–255] gamma encoded luminance
OSP_TEXTURE_LA8	8 bit [0–255] gamma encoded luminance, and linear alpha
OSP_TEXTURE_R32F	32 bit float single component

Table 4.40 – Supported texture formats by `texture2D`, i.e., valid constants of type `OSPTextureFormat`.

Type	Name	Description
OSPVolume	volume	volume used to generate color lookups

Table 4.41 – Parameters of volume texture type

4.5.7 Texture2D Transformations

All materials with textures also offer to manipulate the placement of these textures with the help of texture transformations. If so, this convention shall be used. The following parameters (prefixed with “`texture_name.`”) are combined into one transformation matrix:

Type	Name	Description
vec4f	transform	interpreted as 2×2 matrix (linear part), column-major
float	rotation	angle in degree, counterclockwise, around center
vec2f	scale	enlarge texture, relative to center (0.5, 0.5)
vec2f	translation	move texture in positive direction (right/up)

Table 4.42 – Parameters to define texture coordinate transformations.

The transformations are applied in the given order. Rotation, scale and translation are interpreted “texture centric”, i.e., their effect seen by an user are relative to the texture (although the transformations are applied to the texture coordinates).

4.5.8 Cameras

To create a new camera of given type use

```
OSPCamera ospNewCamera(const char *type);
```

The call returns `NULL` if that type of camera is not known, or else an `OSPCamera` handle to the created camera. All cameras accept these parameters:

The camera is placed and oriented in the world with `pos`, `dir` and `up`. OSPRay uses a right-handed coordinate system. The region of the camera sensor that is

Type	Name	Description
vec3f(a)	pos	position of the camera in world-space
vec3f(a)	dir	main viewing direction of the camera
vec3f(a)	up	up direction of the camera
float	nearClip	near clipping distance
vec2f	imageStart	start of image region (lower left corner)
vec2f	imageEnd	end of image region (upper right corner)

Table 4.43 – Parameters accepted by all cameras.

rendered to the image can be specified in normalized screen-space coordinates with `imageStart` (lower left corner) and `imageEnd` (upper right corner). This can be used, for example, to crop the image, to achieve asymmetrical view frusta, or to horizontally flip the image to view scenes which are specified in a left-handed coordinate system. Note that values outside the default range of [0–1] are valid, which is useful to easily realize overscan or film gate, or to emulate a shifted sensor.

4.5.8.1 Perspective Camera

The perspective camera implements a simple thinlens camera for perspective rendering, supporting optionally depth of field and stereo rendering, but no motion blur. It is created by passing the type string “perspective” to `ospNewCamera`. In addition to the [general parameters](#) understood by all cameras the perspective camera supports the special parameters listed in the table below.

Table 4.44 – Parameters accepted by the perspective camera.

Type	Name	Description
float	fovy	the field of view (angle in degree) of the frame’s height
float	aspect	ratio of width by height of the frame (and image region)
float	apertureRadius	size of the aperture, controls the depth of field
float	focusDistance	distance at where the image is sharpest when depth of field is enabled
bool	architectural	vertical edges are projected to be parallel
int	stereoMode	0: no stereo (default), 1: left eye, 2: right eye, 3: side-by-side
float	interpupillaryDistance	distance between left and right eye when stereo is enabled

Note that when computing the aspect ratio a potentially set image region (using `imageStart` & `imageEnd`) needs to be regarded as well.

In architectural photography it is often desired for aesthetic reasons to display the vertical edges of buildings or walls vertically in the image as well, regardless of how the camera is tilted. Enabling the `architectural` mode achieves this by internally leveling the camera parallel to the ground (based on the up direction) and then shifting the lens such that the objects in direction `dir` are centered in the image. If finer control of the lens shift is needed use `imageStart` & `imageEnd`. Because the camera is now effectively leveled its image plane and thus the plane of focus is oriented parallel to the front of buildings, the whole façade appears sharp, as can be seen in the example images below.

4.5.8.2 Orthographic Camera

The orthographic camera implements a simple camera with orthographic projection, without support for depth of field or motion blur. It is created by passing the

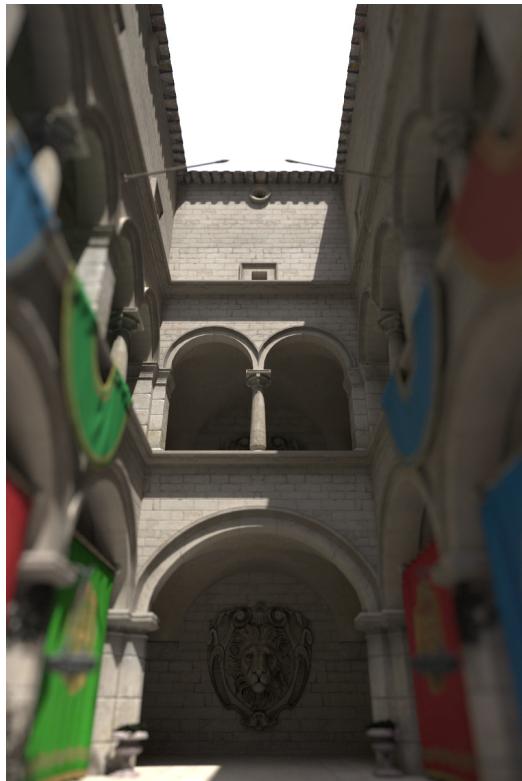


Figure 4.16 – Example image created with the perspective camera, featuring depth of field.

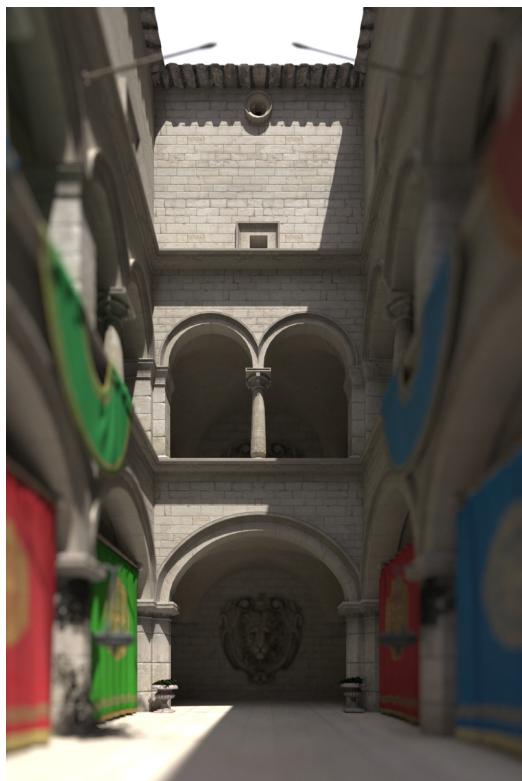


Figure 4.17 – Enabling the architectural flag corrects the perspective projection distortion, resulting in parallel vertical edges.



Figure 4.18 – Example 3D stereo image using `stereoMode` side-by-side.

type string “orthographic” to `ospNewCamera`. In addition to the [general parameters](#) understood by all cameras the orthographic camera supports the following special parameters:

Type	Name	Description
float	height	size of the camera’s image plane in y, in world coordinates
float	aspect	ratio of width by height of the frame

Table 4.45 – Parameters accepted by the orthographic camera.

For convenience the size of the camera sensor, and thus the extent of the scene that is captured in the image, can be controlled with the `height` parameter. The same effect can be achieved with `imageStart` and `imageEnd`, and both methods can be combined. In any case, the `aspect` ratio needs to be set accordingly to get an undistorted image.

4.5.8.3 Panoramic Camera

The panoramic camera implements a simple camera without support for motion blur. It captures the complete surrounding with a latitude / longitude mapping and thus the rendered images should best have a ratio of 2:1. A panoramic camera is created by passing the type string “panoramic” to `ospNewCamera`. It is placed and oriented in the scene by using the [general parameters](#) understood by all cameras.

4.5.9 Picking

To get the world-space position of the geometry (if any) seen at [0–1] normalized screen-space pixel coordinates `screenPos` use

```
void ospPick(OSPPickResult*, OSPRender, const vec2f &screenPos);
```

The result is returned in the provided `OSPPickResult` struct:

```
typedef struct {
    vec3f position; // the position of the hit point (in world-space)
    bool hit;       // whether or not a hit actually occurred
} OSPPickResult;
```




Figure 4.19 – Example image created with the orthographic camera.



Figure 4.20 – Latitude / longitude map created with the panoramic camera.

Note that `ospPick` considers exactly the same camera of the given renderer that is used to render an image, thus matching results can be expected. If the camera supports depth of field then the center of the lens and thus the center of the circle of confusion is used for picking.

4.6 Framebuffer

The framebuffer holds the rendered 2D image (and optionally auxiliary information associated with pixels). To create a new framebuffer object of given size `size` (in pixels), color format, and channels use

```
OSPFrameBuffer ospNewFrameBuffer(const vec2i &size,
                                const OSPFrameBufferFormat format = OSP_FB_SRGB,
                                const uint32_t frameBufferChannels = OSP_FB_COLOR);
```

The parameter `format` describes the format the color buffer has *on the host*, and the format that `ospMapFrameBuffer` will eventually return. Valid values are:

Name	Description
OSP_FB_NONE	framebuffer will not be mapped by the application
OSP_FB_RGBA8	8 bit [0–255] linear component red, green, blue, alpha
OSP_FB_SRGB	8 bit sRGB gamma encoded color components, and linear alpha
OSP_FB_RGBA32F	32 bit float components red, green, blue, alpha

The parameter `frameBufferChannels` specifies which channels the framebuffer holds, and can be combined together by bitwise OR from the values of `OSPFrameBufferChannel` listed in the table below.

Name	Description
OSP_FB_COLOR	RGB color including alpha
OSP_FB_DEPTH	euclidean distance to the camera (<i>not</i> to the image plane), as linear 32 bit float
OSP_FB_ACCUM	accumulation buffer for progressive refinement
OSP_FB_VARIANCE	for estimation of the current noise level if <code>OSP_FB_ACCUM</code> is also present, see rendering
OSP_FB_NORMAL	accumulated world-space normal of the first hit, as <code>vec3f</code>
OSP_FB_ALBEDO	accumulated material albedo (color without illumination) at the first hit, as <code>vec3f</code>

If a certain channel value is *not* specified, the given buffer channel will not be present. Note that OSPRay makes a clear distinction between the *external* format of the framebuffer and the internal one: The external format is the format the user specifies in the `format` parameter; it specifies what color format OSPRay will eventually *return* the framebuffer to the application (when calling `ospMapFrameBuffer`): no matter what OSPRay uses internally, it will simply return a 2D array of pixels of that format, with possibly all kinds of reformatting, compression/decompression, etc., going on in-between the generation of the *internal* framebuffer and the mapping of the externally visible one.

Table 4.46 – Supported color formats of the framebuffer that can be passed to `ospNewFrameBuffer`, i.e., valid constants of type `OSPFrameBufferFormat`.

Table 4.47 – Framebuffer channels constants (of type `OSPFrameBufferChannel`), naming optional information the framebuffer can store. These values can be combined by bitwise OR when passed to `ospNewFrameBuffer` or `ospFrameBufferClear`.

In particular, OSP_FB_NONE is a perfectly valid pixel format for a framebuffer that an application will never map. For example, an application driving a display wall may well generate an intermediate framebuffer and eventually transfer its pixel to the individual displays using an OSPPixelOp [pixel operation](#).

The application can map the given channel of a framebuffer – and thus access the stored pixel information – via

```
const void *ospMapFrameBuffer(OSPFrameBuffer,
                               const OSPFrameBufferChannel = OSP_FB_COLOR);
```

Note that OSP_FB_ACCUM or OSP_FB_VARIANCE cannot be mapped. The origin of the screen coordinate system in OSPRay is the lower left corner (as in OpenGL), thus the first pixel addressed by the returned pointer is the lower left pixel of the image.

A previously mapped channel of a framebuffer can be unmapped by passing the received pointer mapped to

```
void ospUnmapFrameBuffer(const void *mapped, OSPFrameBuffer);
```

The individual channels of a framebuffer can be cleared with

```
void ospFrameBufferClear(OSPFrameBuffer, const uint32_t frameBufferChannels);
```

When selected, OSP_FB_COLOR will clear the color buffer to black (0, 0, 0), OSP_FB_DEPTH will clear the depth buffer to `inf`. OSP_FB_ACCUM will clear *all* accumulating buffers (OSP_FB_VARIANCE, OSP_FB_NORMAL, and OSP_FB_ALBEDO, if present) and resets the accumulation counter `accumID`.

Pixel Operation

Pixel operations are functions that are applied to every pixel that gets written into a framebuffer. Examples include post-processing, filtering, blending, tone mapping, or sending tiles to a display wall. To create a new pixel operation of given type `type` use

```
OSPPixelOp ospNewPixelOp(const char *type);
```

The call returns `NULL` if that type is not known, or else an OSPPixelOp handle to the created pixel operation.

To set a pixel operation to the given framebuffer use

```
void ospSetPixelOp(OSPFrameBuffer, OSPPixelOp);
```

Tone Mapper

The tone mapper is a pixel operation which implements a generic filmic tone mapping operator. Using the default parameters it approximates the Academy Color Encoding System (ACES). The tone mapper is created by passing the type string “tonemapper” to `ospNewPixelOp`. The tone mapping curve can be customized using the parameters listed in the table below.

To use the popular “Uncharted 2” filmic tone mapping curve instead, set the parameters to the values listed in the table below.

Table 4.48 – Parameters accepted by the tone mapper.

Type	Name	Default	Description
float	contrast	1.6773	contrast (toe of the curve); typically is in [1–2]
float	shoulder	0.9714	highlight compression (shoulder of the curve); typically is in [0.9–1]
float	midIn	0.18	mid-level anchor input; default is 18% gray
float	midOut	0.18	mid-level anchor output; default is 18% gray
float	hdrMax	11.0785	maximum HDR input that is not clipped
bool	acesColor	true	apply the ACES color transforms

Name	Value
contrast	1.1759
shoulder	0.9746
midIn	0.18
midOut	0.18
hdrMax	6.3704
acesColor	false

Table 4.49 – Filmic tone mapping curve parameters. Note that the curve includes an exposure bias to match 18% middle gray.

4.7 Rendering

To render a frame into the given framebuffer with the given renderer use

```
float ospRenderFrame(OSPFrameBuffer, OSPRenderer,
                      const uint32_t frameBufferChannels = OSP_FB_COLOR);
```

The third parameter specifies what channel(s) of the framebuffer is written to⁶. What to render and how to render it depends on the renderer's parameters. If the framebuffer supports accumulation (i.e., it was created with `OSP_FB_ACCUM`) then successive calls to `ospRenderFrame` will progressively refine the rendered image. If additionally the framebuffer has an `OSP_FB_VARIANCE` channel then `ospRenderFrame` returns an estimate of the current variance of the rendered image, otherwise `inf` is returned. The estimated variance can be used by the application as a quality indicator and thus to decide whether to stop or to continue progressive rendering.

⁶ This is currently not implemented, i.e., all channels of the framebuffer are always updated.

Progress and Cancel

To be informed about the progress of rendering the current frame the application can register a callback function of type

```
typedef int (*OSPProgressFunc)(void* userPtr, const float progress);
```

via

```
void ospSetProgressFunc(OSPProgressFunc, void* userPtr);
```

The provided user pointer `userPtr` is passed as first argument to the callback function⁷ and the reported progress is in (0–1]. If the callback function returns zero than the application requests to cancel rendering, i.e., the current `ospRenderFrame` will return at the first opportunity and the content of the framebuffer will be undefined. Therefore, better clear the framebuffer with `ospFrameBufferClear` then before a subsequent call of `ospRenderFrame`.

Passing `NULL` as `OSPProgressFunc` function pointer disables the progress callback.

⁷ That way applications can also register a member function of a C++ class together with the `this` pointer as `userPtr`.

Chapter 5

Parallel Rendering with MPI

OSPRay has the ability to scale to multiple nodes in a cluster via MPI. This enables applications to take advantage of larger compute and memory resources when available.

5.1 Prerequisites for MPI Mode

In addition to the standard build requirements of OSPRay, you must have the following items available in your environment in order to build & run OSPRay in MPI mode:

- An MPI enabled multi-node environment, such as an HPC cluster
- An MPI implementation you can build against (i.e., Intel MPI, MVAPICH2, etc...)

5.2 Enabling the MPI Module in your Build

To build the MPI module the CMake option `OSPRAY_MODULE_MPI` must be enabled, which can be done directly on the command line (with `-DOSPRAY_MODULE_MPI=ON`) or through a configuration dialog (`ccmake`, `cmake-gui`), see also [Compiling OSPRay](#).

This will trigger CMake to go look for an MPI implementation in your environment. You can then inspect the CMake value of `MPI_LIBRARY` to make sure that CMake found your MPI build environment correctly.

This will result in an OSPRay module being built. To enable using it, applications will need to either link `libospray_module_mpi`, or call

```
ospLoadModule("mpi");
```

before initializing OSPRay.

5.3 Modes of Using OSPRay's MPI Features

OSPRay provides two ways of using MPI to scale up rendering: offload and distributed.

5.3.1 Offload Rendering

The “offload” rendering mode is where a single (not-distributed) calling application treats the OSPRay API the same as with local rendering. However, OSPRay uses multiple MPI connected nodes to evenly distribute frame rendering work,

where each node contains a full copy of all scene data. This method is most effective for scenes which can fit into memory, but are expensive to render: for example, path tracing with many samples-per-pixel is compute heavy, making it a good situation to use the offload feature. This can be done with any application which already uses OSPRay for local rendering without the need for any code changes.

When doing MPI offload rendering, applications can optionally enable dynamic load balancing, which can be beneficial in certain contexts. This load balancing refers to the distribution of tile rendering work across nodes: thread-level load balancing on each node is still dynamic with the thread tasking system. The options for enabling/controlling the dynamic load balancing features on the `mpi_offload` device are found in the table below, which can be changed while the application is running. Please note that these options will likely only pay off for scenes which have heavy rendering load (e.g., path tracing a non-trivial scene) and have much variance in how expensive each tile is to render.

Type	Name	Default	Description
bool	<code>dynamicLoadBalancer</code>	false	whether to use dynamic load balancing

Table 5.1 – Parameters specific to the `mpi_offload` device.

5.3.2 Distributed Rendering

The “distributed” rendering mode is where a MPI distributed application (such as a scientific simulation) uses OSPRay collectively to render frames. In this case, the API expects all calls (both created objects and parameters) to be the same on every application rank, except each rank can specify arbitrary geometries and volumes. Each renderer will have its own limitations on the topology of the data (i.e., overlapping data regions, concave data, etc.), but the API calls will only differ for scene objects. Thus all other calls (i.e., setting camera, creating framebuffer, rendering frame, etc.) will all be assumed to be identical, but only rendering a frame and committing the model must be in lock-step. This mode targets using all available aggregate memory for huge scenes and for “in-situ” visualization where the data is already distributed by a simulation app.

5.4 Running an Application with the “offload” Device

As an example, our sample viewer can be run as a single application which offloads rendering work to multiple MPI processes running on multiple machines.

The example apps are setup to be launched in two different setups. In either setup, the application must initialize OSPRay with the offload device. This can be done by creating an “`mpi_offload`” device and setting it as the current device (via the `ospSetCurrentDevice()` function), or passing either “`--osp:mpi`” or “`--osp:mpi-offload`” as a command line parameter to `ospInit()`. Note that passing a command line parameter will automatically call `ospLoadModule("mpi")` to load the MPI module, while the application will have to load the module explicitly if using `ospNewDevice()`.

5.4.1 Single MPI Launch

OSPRay is initialized with the `ospInit()` function call which takes command line arguments in and configures OSPRay based on what it finds. In this setup,

the app is launched across all ranks, but workers will never return from `osInit()`, essentially turning the application into a worker process for OSPRay. Here's an example of running the `ospVolumeViewer` data-replicated, using `c1-c4` as compute nodes and `localhost` the process running the viewer itself:

```
mpirun -perhost 1 -hosts localhost,c1,c2,c3,c4 ./ospExampleViewer <scene file> --osp:mpi
```

5.4.2 Separate Application & Worker Launches

The second option is to explicitly launch the app on rank 0 and worker ranks on the other nodes. This is done by running `ospray_mpi_worker` on worker nodes and the application on the display node. Here's the same example above using this syntax:

```
mpirun -perhost 1 -hosts localhost ./ospExampleViewer <scene file> --osp:mpi \
: -hosts c1,c2,c3,c4 ./ospray_mpi_worker
```

This method of launching the application and OSPRay worker separately works best for applications which do not immediately call `osInit()` in their `main()` function, or for environments where application dependencies (such as GUI libraries) may not be available on compute nodes.

5.5 Running an Application with the “distributed” Device

Applications using the new distributed device should initialize OSPRay by creating (and setting current) an “`mpi_distributed`” device or pass “`--osp:mpi-distributed`” as a command line argument to `osInit()`. Note that due to the semantic differences the distributed device gives the OSPRay API, it is not expected for applications which can already use the offload device to correctly use the distributed device without changes to the application.

The following additional parameter can be set on the `mpi_distributed` device.

Table 5.2 – Parameters for the `mpi_distributed` device.

Type	Name	Description
<code>void*</code>	<code>worldCommunicator</code>	A pointer to the <code>MPI_Comm</code> which should be used as OSPRay’s world communicator. This will set how many ranks OSPRay should expect to participate in rendering. The default is <code>MPI_COMM_WORLD</code> where all ranks are expected to participate in rendering.

By setting the `worldCommunicator` parameter to a different communicator than `MPI_COMM_WORLD` the client application can tune how OSPRay is run within its processes. The default uses `MPI_COMM_WORLD` and thus expects all processes to also participate in rendering, thus if a subset of processes do not call collectives like `ospRenderFrame` the application would hang.

For example, an MPI parallel application may be run with one process per core, however OSPRay is multithreaded and will perform best when run with one process per-node. By splitting `MPI_COMM_WORLD` the application can create a communicator with one rank per-node to then run OSPRay on one process per-node. The remaining ranks on each node can then aggregate their data to the OSPRay process for rendering.

The model used by the distributed device takes three additional parameters, to allow users to express their data distribution to OSPRay. All models should

Table 5.3 – Parameters for the distributed OSPModel.

Type	Name	Description
int	id	An integer that uniquely identifies this piece of distributed data. For example, in a common case of one sub-brick per-rank, this would just be the region's MPI rank. Multiple ranks can specify models with the same ID, in which case the rendering work for the model will be shared among them.
vec3f	region.lower	Override the original model geometry + volume bounds with a custom lower bound position. This can be used to clip geometry in the case the objects cross over to another region owned by a different node. For example, rendering a set of spheres with radius.
vec3f	region.upper	Override the original model geometry + volume bounds with a custom upper bound position.

be disjoint to ensure correct sort-last compositing. Geometries used in the distributed MPI renderer can make use of the [SciVis renderer's OBJ material](#).

The renderer supported when using the distributed device is the `mpi_raycast` renderer. This renderer is an experimental renderer and currently only supports ambient occlusion (on the local data only, with optional ghost data). To compute correct ambient occlusion across the distributed data the application is responsible for replicating ghost data and specifying the ghost models and models as described above. Note that shadows and ambient occlusion are computed on the local geometries, in the model and the corresponding `ghostModel` in the ghost model array, if any where set.

Table 5.4 – Parameters for the `mpi_raycast` renderer.

Type	Name	Default	Description
OSPModel/OSPModel[]	model	NULL	the <code>model</code> to render, can optionally be a <code>data</code> array of multiple models
OSPModel/OSPModel[]	ghostModel	NULL	the optional <code>model</code> containing the ghost geometry for ambient occlusion; when setting a <code>data</code> array for both <code>model</code> and <code>ghostModel</code> , each individual ghost model shadows only its corresponding model
OSPLight[]	lights		<code>data</code> array with handles of the [lights]
int	aoSamples	0	number of rays per sample to compute ambient occlusion
bool	aoTransparencyEnabled	false	whether object transparency is respected when computing ambient occlusion (slower)
bool	oneSidedLighting	true	if true, backfacing surfaces (wrt. light source) receive no illumination

See the distributed device examples in the MPI module for examples.

Chapter 6

Scene graph

WARNING: USE AT YOUR OWN RISK. The Scene graph is currently in Alpha mode and will change frequently. It is not yet recommended for critical production work.

The scene graph is the basis of our Example Viewer which consists of a superset of OSPRay objects represented in a graph hierarchy (currently a tree). This graph functions as a hierarchical specification for scene properties and a self-managed update graph. The scene graph infrastructure includes many convenience functions for templated traversals, queries of state and child state, automated updates, and timestamped modifications to underlying state.

The scene graph nodes closely follow the dependencies of existing OSPRay API internals, i.e., a `sg::Renderer` has a “model” child, which in turn has a “`TriangleMesh`”, which in turn has a child named “`vertex`” similar to how you may set the “`vertex`” parameter on the `osp::TriangleMesh` which in turn is added to an `OSPMModel` object which is set as the model on the `OSPRenderer`. The scene graph is a superset of OSPRay functionality so there is no direct 1:1 mapping between the scene graph hierarchy in all cases, however it is kept as close as possible. This makes the scene graph viewer in the [Example Viewer] a great way to understand OSPRay state.

6.1 Hierarchy Structure

The root of the scene graph is based on `sg::Renderer`. The scene graph can be created by

```
auto renderer = sg::createNode("renderer", "Renderer");
```

which automatically creates child nodes for necessary OSPRay state. To update and commit all state and render a single function is provided which can be called with:

```
renderer.renderFrame(renderer["frameBuffer"].nodeAs<sg::FrameBuffer>());
```

Values can be set using:

```
renderer["spp"] = 16;
```

To explore the full set of nodes, simply launch the Example Viewer and traverse through the GUI representation of all scene graph nodes.

6.2 Traversals

The scene graph contains a set of builtin traversals as well as modular visitor functors for implementing custom passes over the scene graph. The required traversals are handled for you by default within the `renderFrame` function on the renderer. For any given node there are two phases to a traversal operation, pre and post traversal of the nodes children. `preTraversal` initializes node state and objects and sets the current traversal context with appropriate state. For instance, `sg::Model` will create a new `OSPModel` object, set its value to that object, and set `sg::RenderContext.currentOSPModel` to its own value. After `preTraversal` is finished, the children of `sg::Model` are processed in a similar fashion and now use the modified context. In `postTraversal`, `sg::Model` will commit the changes that its children have potentially set and it will pop its modifications from the current context. This behavior is replicated for every scene graph node and enables children to act on parent state without specific implementations from the parent node. An example of this are the `sg::NodeParam` nodes which are containers for values to be set on `OSPRObjects`, such as a float value. This is put on the scene graph with a call to:

```
renderer["lights"]["sun"].createChild("intensity", "float", 0.3f);
```

This call accesses the child named “lights” on the renderer, and in turn the child named “sun”. This child then gets its own child of a newly created node with the name “intensity” of type `float` with a value of `0.3f`. When committed, this node will call `ospSet1f` with the node value on the current `OSPRObject` on the context which is set by the parent. If you were to create a custom light called “MyLight” and had a float parameter called “`flickerFreq`”, a similar line would be used without requiring any additional changes in the scene graph internals beyond registering the new light class. Known parameters such as floats will also show up in the Example Viewer GUI without requiring any additional code beyond adding them to the scene graph and the internal implementation in OSPRay.

The base passes required to use the scene graph include `verification`, `commit`, and `render` traversals. Every node in the scene graph has a valid state which needs to be set before operating on the node. Nodes may have custom qualifications for validity, but by default they are set through `valid_` flags on the scene graph `Node` for things like whitelists and range checks. Once verified, `commit` traverses the scene graph and commits scene graph state to OSPRay. Commits are timestamped, so re-committing will only have any affect if a dependent child has been modified requiring a new commit. Because of this, each node does not have to track if it is valid or if anything in the scene has been modified, as `commit` will only be called on that node if those are already true. By default invalid nodes will throw exceptions, however this can be turned off which enables the program to keep running. In the Example Viewer GUI, invalid nodes will be marked in red but the previously committed state will keep rendering until the invalid state is corrected.

For examples of implementing custom traversals, see the `sg/visitors` folder. Here is an example of a visitor that collects all nodes with a given name:

```
struct GatherNodesByName : public Visitor
{
    GatherNodesByName(const std::string &_name);

    bool operator()(Node &node, TraversalContext &ctx) override;

    std::vector<std::shared_ptr<Node>> results();
};
```

```
private:
    std::string name;
    std::vector<std::shared_ptr<Node>> nodes;
};

// Inlined definitions ///////////////////////////////// ///////////////////////////////// /////////////////////////////////

inline GatherNodesByName::GatherNodesByName(const std::string &_name)
: name(_name)
{
}

inline bool GatherNodesByName::operator()(Node &node, TraversalContext &
{
    if (utility::longestBeginningMatch(node.name(), this->name) == this->name) {
        auto itr = std::find_if(
            nodes.begin(),
            nodes.end(),
            [&](const std::shared_ptr<Node> &nodeInList) {
                return nodeInList.get() == &node;
            }
        );
        if (itr == nodes.end())
            nodes.push_back(node.shared_from_this());
    }

    return true;
}

inline std::vector<std::shared_ptr<Node>> GatherNodesByName::results()
{
    return nodes;// TODO: should this be a move (i.e., reader 'consumes')?
}
```

6.3 Thread Safety

The scene graph is only thread safe for accessing and setting values on nodes. More advanced operations like adding or removing nodes are not thread safe. At some point we hope to add transactions to handle these, but for now the scene graph nodes must be added/removed on the same thread that is committing and rendering.

Chapter 7

Examples

7.1 Simple Tutorial

A minimal working example demonstrating how to use OSPRay can be found at `apps/ospTutorial.c`¹. On Linux build it in the build directory with

```
gcc -std=c99 ..\apps\ospTutorial.c -I ..\ospray\include -I .. \
./libospray.so -Wl,-rpath,. -o ospTutorial
```

On Windows build it in the “build_directory\\$Configuration” with

```
cl ..\..\apps\ospTutorial.c -I ..\..\ospray\include -I ..\.. ospray.lib
```

Running `ospTutorial` will create two images of two triangles, rendered with the Scientific Visualization renderer with full Ambient Occlusion. The first image `firstFrame.ppm` shows the result after one call to `ospRenderFrame` – jagged edges and noise in the shadow can be seen. Calling `ospRenderFrame` multiple times enables progressive refinement, resulting in antialiased edges and converged shadows, shown after ten frames in the second image `accumulated-Frames.ppm`.

¹ A C++ version that uses the C++ convenience wrappers of OSPRay’s C99 API via `include/ospray/ospray_cpp.h` is available at `apps/ospTutorial.cpp`.




Figure 7.1 – First frame.



Figure 7.2 – After accumulating ten frames.

7.2 Mini-App Tutorials

OSPRay also ships various mini-apps to showcase OSPRay features. These apps are all prefixed with `ospTutorial` and can be found in the `tutorials/` directory of the OSPRay source tree. Current tutorials include:

- structured volumes
- unstructured volumes
- spheres
- animated spheres
- subdivision surfaces

More apps will be created in future releases to further demonstrate other interesting OSPRay features.

7.3 Example Viewer

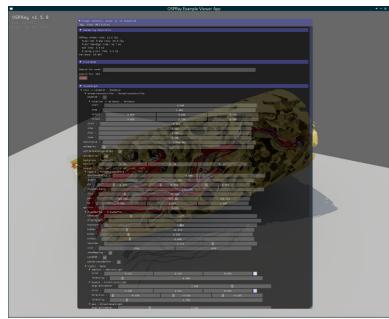


Figure 7.3 – Screenshot of using `ospExampleViewer` with a scene graph.

OSPRay includes an exemplary viewer application `ospExampleViewer`, showcasing most features of OSPRay which can be run as `./ospExampleViewer [options] <filename>`. The Example Viewer uses the `ImGui` library for user interface controls and is based on a prototype OSPRay [scene graph](#) interface where nodes can be viewed and edited interactively. Updates to scene graph nodes update OSPRay state automatically through the scene graph viewer which is enabled by pressing 'g'.

7.3.1 Exploring the Scene

The GUI shows the entire state of the program under the root scene graph node. Expanding nodes down to explore and edit the scene is possible, for example a material parameter may be found under `renderer→world→mesh→material→Kd`. Updates to values will be automatically propagated to the next render. Individual nodes can be easily found using the “Find Node” section, which will find nodes with a given name based on the input string. Scene objects can also be selected with the mouse by shift-left clicking in the viewer.

Click on nodes to expand their children, whose values can be set by dragging or double clicking and typing in values. You can also add new nodes where appropriate: for example, when “lights” is expanded right clicking on “lights” and selecting create new node and typing in a light type, such as “PointLight”, will add it to the scene. Similarly, right clicking on “world” and creating an “Importer” node will add a new scene importer from a file. Changing the filename to an appropriate file will load the scene and propagate the resulting state. Exporting and importing the scene graph is only partially supported at the moment through “`ospsg`” files. Currently, any nodes with Data members will break this functionality, however right clicking and selecting export on the camera or lights

nodes for instance will save out their respective state which can be imported on the command line. The Example Viewer also functions as an OSPRay state debugger – invalid values will be shown in red up the hierarchy and will not change the viewer until corrected.

7.3.2 Volume Rendering

Volumes are loaded into the viewer just as a mesh is. Volume appearance is modified according to the transfer function, which will show up in a popup window on the GUI after pressing ‘g’. Click and drag across the transfer function to set opacity values, and selecting near the bottom of the editable transfer function widget sets the opacity to zero. The colors themselves can only be modified by selecting from the dropdown menu ‘ColorMap’ or importing and exporting json colors. The range that the transfer function operates on can be modified on the scene graph viewer.

7.3.3 Controls

- ‘g’ - toggle scene graph display
- ‘q’ - quit
- Left click and drag to rotate
- Right click and drag or mouse wheel to zoom in and out.
- Mouse-Wheel click will pan the camera.
- Control-Left clicking on an object will select a model and all of its children which will be displayed in the
- Shift-Left click on an object will zoom into that part of the scene and set the focal distance.

7.3.4 Command Line Options

- Running `./ospExampleViewer -help` will bring up a list of command line options. These options allow you to load files, run animations, modify any scene graph state, and many other functions. See the [demos](#) page for examples.
- Supported file importers currently include: `obj`, `ply`, `x3d`, `vtu`, `osp`, `ospsg`, `xml` (`rivl`), `points`, `xyz`.

7.3.5 Denoiser

When the example viewer is built with OpenImageDenoise, the denoiser is automatically enabled when running the application. It can be toggled on/off at runtime via the `useDenoiser` GUI parameter found under the framebuffer in the scene graph.

7.4 Distributed Viewer

The application `ospDistribViewerDemo` demonstrates how to write a distributed SciVis style interactive renderer using the distributed MPI device. Note that because OSPRay uses sort-last compositing it is up to the user to ensure that the data distribution across the nodes is suitable. Specifically, each nodes’ data must be convex and disjoint. This renderer supports multiple volumes and geometries per node. To ensure they are composited correctly you specify a list of bounding regions to the model, within these regions can be arbitrary volumes/geometries and each rank can have as many regions as needed. As long as the regions are

disjoint/convex the data will be rendered correctly. In this demo we either generate a volume, or load a RAW volume file if one is passed on the command line.

7.4.1 Loading a RAW Volume

To load a RAW volume you must specify the filename (-f <file>), the data type (-dtype <dtype>), the dimensions (-dims <x> <y> <z>) and the value range for the transfer function (-range <min> <max>). For example, to run on the [CSAFE data set from the demos page](#) you would pass the following arguments:

```
mpirun -np <n> ./ospDistribViewerDemo \
-f <path to csafe>/csafe-heptane-302-volume.raw \
-dtype uchar -dims 302 302 302 -range 0 255
```

The volume file will then be chunked up into an $x \times y \times z$ grid such that $n = xyz$. See `loadVolume` in [gensv/generateSciVis.cpp](#) for an example of how to properly load a volume distributed across ranks with correct specification of brick positions and ghost voxels for interpolation at boundaries. If no volume file data is passed a volume will be generated instead, in that case see `makeVolume`.

7.4.2 Geometry

The viewer can also display some randomly generated sphere geometry if you pass `-spheres <n>` where n is the number of spheres to generate per-node. These spheres will be generated inside the bounding box of the region's volume data.

In the case that you have geometry crossing the boundary of nodes and are replicating it on both nodes to render (ghost zones, etc.) the region will be used by the renderer to clip rays against allowing to split the object between the two nodes, with each rendering half. This will keep the regions rendered by each rank disjoint and thus avoid any artifacts. For example, if a sphere center is on the border between two nodes, each would render half the sphere and the halves would be composited to produce the final complete sphere in the image.

7.4.3 App-initialized MPI

Passing the `-appMPI` flag will have the application initialize MPI instead of letting OSPRay do it internally when creating the MPI distributed device. In this case OSPRay will not finalize MPI when cleaning up the device, allowing the application to use OSPRay for some work, shut it down and recreate everything later if needed for additional computation, without accidentally shutting down its MPI communication.

7.4.4 Interactive Viewer

Rank 0 will open an interactive window with GLFW and display the rendered image. When the application state needs to update (e.g., camera or transfer function changes), this information is broadcasted out to the other nodes to update their scene data.

7.5 Demos

Several ready-to-run demos, models and data sets for OSPRay can be found at the [OSPRay Demos and Examples](#) page.

© 2013–2019 Intel Corporation

Intel, the Intel logo, Xeon, Intel Xeon Phi, and Intel Core are trademarks of Intel Corporation in the U.S. and/or other countries.
*Other names and brands may be claimed as the property of others.

Optimization Notice: Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice Revision #20110804